

Introduction to Machine Learning

CMPT 726

Mo Chen

SFU Computing Science

Sept. 9, 2020

Bishop PRML Ch. 1

About Me


- Undergraduate from UBC
- PhD from UC Berkeley
- Postdoc at Stanford
- Assistant Professor at SFU CS since 2018
 - Multi-Agent Robotic Systems Lab (<https://sfumars.com>)



[Administrivia](#)




[Machine Learning](#)



[Curve Fitting](#)



[Coin Tossing](#)



Outline

[Administrivia](#)

[Machine Learning](#)

[Curve Fitting](#)

[Coin Tossing](#)

About This Course

Lectures: Mon 16:30-18:20, Wed 16:30-17:20, Canvas Bb Collab Ultra

- Website: <https://canvas.sfu.ca/courses/56287>

Piazza online discussion and Q&A :

- Links to sign-up and home page on website

Office hours:

- Tentatively Fridays 14:00-15:00, Zoom link on website
- See website for TA office hours

Teaching Assistants

- Xubo Lyu (xlv@sfu.ca)
- Roshni Shaik (rshaik@sfu.ca)
- Ria Thomas (riat@sfu.ca)

Administrivia

- We will cover techniques in the standard ML toolkit
 - maximum likelihood, regularization, neural networks, Markov random fields (MRF), graphical models, expectation-maximization (EM), mixture models, hidden Markov models (HMM), particle filters, ...
- There will be 3 assignments and 1 project
 - Assignment 0: 5%
 - Assignment 1-3: 15% each
 - Project: 50%
- No exams!

Administrivia

- Recommend doing associated readings from Bishop, *Pattern Recognition and Machine Learning* (PRML) after each lecture
 - Other reference books
 - *The Elements of Statistical Learning*, Trevor Hastie, Robert Tibshirani, and Jerome Friedman
 - *Machine Learning*, Tom Mitchell
 - *Pattern Classification (2nd ed.)*, Richard O. Duda, Peter E. Hart, and David G. Stork
 - *Information Theory, Inference, and Learning Algorithms*, David MacKay (available online)
 - *Deep Learning*, Ian Goodfellow, Yoshua Bengio and Aaron Courville (available online)
 - Online courses
 - Coursera, Udacity

Administrivia - Assignments

- Assignments will be monthly
- No late submissions other than exceptional circumstances (contact me directly)!
- Programming questions mainly use Python

Administrivia – Assignment Peer Grading

- Double blind peer grading:
 - Up to five peers will grade your assignment according to detailed solutions and rubrics
 - Assignment mark will be a combination of
 - Scores from peers (75%)
 - Quality of grading (25%)
 - Insightful comments (5% bonus)
 - Disputes to be handled directly by teaching staff
- Some benefits:
 - Frees up teaching staff's time for more help with assignments and projects
 - Ensures assignments are thoroughly reviewed
 - Ensures thorough review of assignment solutions

Administrivia – Assignments

- Assignment 0: introduce yourself to 5 peer reviewers
 - Becoming familiar with Canvas and peer review
 - Will not really be double blind this time
 - Assignment due Sept. 14 (next Monday)
 - Peer review due Sept. 18 (next Friday)
- Exact deadlines: always 23:59 on the specified date
- Assignment 1:
 - Assigned Sept. 18, due Oct. 9, reviews due Oct. 16
- Assignment 2:
 - Assigned Oct. 16, due Nov. 6, reviews due Nov. 13
- Assignment 3:
 - Assigned Nov. 13, due Dec. 4, reviews due Dec. 11

Administrivia - Project

- Project details
 - Practice doing research
 - Ideal project – take problem from your research/interests, use ML (properly)
 - We will provide a list of *pre-approved* projects proposed by TAs and SFU research labs
 - This gives some of you a chance to work with SFU research labs
 - Other projects will need to be approved: Check with a TA to make sure your project is reasonable/feasible

Administrivia - Project

- Project details
 - Work in groups (2-3 students)
 - Produce (short) research paper
 - Graded on proper research methodology, not just results
 - Choice of problem / algorithms
 - Relation to previous work
 - Comparative experiments
 - Quality of exposition
 - Students must list their contribution to the project within the group
 - Details on course webpage
 - Poster session/presentations TBA
 - Report due Dec. 16 at 23:59

Administrivia - Background

- Calculus:

$$E = mc^2 \Rightarrow \frac{\partial E}{\partial c} = 2mc$$

- Linear algebra:

$$Au_i = \lambda_i u_i; \frac{\partial}{\partial x}(x^\top a) = a$$

- See PRML Appendix C
- Probability:

$$p(X) = \sum_Y p(X, Y); \quad p(x) = \int p(x, y) dy; \quad \mathbb{E}_x[f] = \int p(x)f(x)dx$$

- See PRML Ch. 1.2



It will be possible to refresh, but if you've never seen these before this course will be **very** difficult.

What is Machine Learning (ML)?

- Algorithms that automatically improve performance through experience
- Often this means define a model by hand, and use data to fit its parameters

Why ML?

- The real world is complex – difficult to hand-craft solutions.
- ML is the preferred framework for applications in many fields:
 - Computer Vision
 - Natural Language Processing, Speech Recognition
 - Robotics
 - ...

Hand-written Digit Recognition



Belongie et al. PAMI 2002

- Difficult to hand-craft rules about digits

Hand-written Digit Recognition


 $x_i =$ 

$t_i = (0,0,0,1,0,0,0,0,0,0)$

- Represent input image as a vector $x_i \in \mathbb{R}^{784}$.
- Suppose we have a target vector t_i
- This is **supervised learning**
 - Discrete, finite label set: perhaps $t_i \in \{0,1\}^{10}$, a **classification** problem
- Given a **training set** $\{(x_1, t_1), \dots, (x_N, t_N)\}$, learning problem is to construct a “good” function $y(x)$ from these:


$$y: \mathbb{R}^{784} \rightarrow \mathbb{R}^{10}$$

Face Detection





Schneiderman and Kanade, IJCV 2004

- Classification problem
- $t_i \in \{0,1,2\}$, non-face, frontal face, profile face.




Spam Detection



- Classification problem
- $t_i \in \{0,1\}$, non-spam,spam
- x_i counts of words, e.g. Viagra, stock, outperform, multi-bagger

Stock Price Prediction



- Problems in which t_i is continuous are called **regression**
- E.g. t_i is stock price, x_i contains company profit, debt, cash flow, gross sales, number of spam emails sent, . . .

Clustering Images




Wang et al., CVPR 2006

- Only x_i is defined: **unsupervised learning**
- E.g. x_i describes image, find groups of similar images

Types of Learning Problems

- Supervised Learning
 - Classification
 - Regression
- Unsupervised Learning
- Reinforcement Learning

An Example - Polynomial Curve Fitting



- Suppose we are given training set of N observations (x_1, \dots, x_N) and (t_1, \dots, t_N) , $x_i, t_i \in \mathbb{R}$
- Regression problem, estimate $y(x)$ from these data

Polynomial Curve Fitting



- What form is $y(x)$?
 - Let's try polynomials of degree M :

$$y(x, w) = w_0 + w_1 x + w_2 x^2 + \cdots + w_M x^M$$

- This is the **hypothesis space**.
- How do we measure success?
- Sum of squared errors:

$$E(w) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, w) - t_n\}^2$$

- Among functions in the class, choose that which minimizes this error



Polynomial Curve Fitting

- Error function





$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

- Best coefficients

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} E(\mathbf{w})$$


- Found using pseudo-inverse (more later)

Which Degree of Polynomial?



- A **model selection** problem
- $M = 9 \rightarrow E(\mathbf{w}^*) = 0$: This is **over-fitting**

Generalization




- Generalization is the holy grail of ML
 - Want good performance for new data
- Measure generalization using a separate set
 - Use root-mean-squared (RMS) error: $E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$

Validation Set


- Split training data into **training set** and **validation set**
- Train different models (e.g. diff. order polynomials) on training set
- Choose model (e.g. order of polynomial) with minimum error on validation set

Cross-validation



- Data are often limited
- Cross-validation creates S groups of data, use $S - 1$ to train, other to validate
 - Extreme case leave-one-out cross-validation (LOO-CV): S is number of training data points
- Cross-validation is an effective method for model selection, but can be slow
 - Models with multiple complexity parameters: exponential number of runs

Controlling Over-fitting: Regularization





	$M = 0$	$M = 1$	$M = 3$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*				17.37
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

- As order of polynomial M increases, so do coefficient magnitudes
- Penalize large coefficients in error function:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$



Controlling Over-fitting: Regularization



	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
w_0^*	0.35	0.35	0.13
w_1^*	232.37	4.74	-0.05
w_2^*	-5321.83	-0.77	-0.06
w_3^*	48568.31	-31.97	-0.05
w_4^*	-231639.30	-3.89	-0.03
w_5^*	640042.26	55.28	-0.02
w_6^*	-1061800.52	41.32	-0.01
w_7^*	1042400.18	-45.95	-0.00
w_8^*	-557682.99	-91.53	0.00
w_9^*	125201.43	72.68	0.01

Fits for $M = 9$

Over-fitting: Dataset size



- With more data, more complex model ($M = 9$) can be fit
- Rule of thumb: 10 datapoints for each parameter

Summary

- Want models that **generalize** to new data
 - Train model on **training set**
 - Measure performance on held-out **test set**
 - Performance on test set is good estimate of performance on new data

Summary - Model Selection

- Which model to use? E.g. which degree polynomial?
 - Training set error is lower with more complex model
 - Can't just choose the model with lowest training error
 - Peeking at test error is unfair. E.g. picking polynomial with lowest test error
 - Performance on test set is no longer good estimate of performance on new data

Summary - Solutions I

- Use a validation set
 - Train models on **training set**. E.g. different degree polynomials
 - Measure performance on held-out **validation set**
 - Measure performance of that model on held-out **test set**
- Can use **cross-validation** on training set instead of a separate validation set if little data and lots of time
 - Choose model with lowest error over all cross-validation folds (e.g. polynomial degree)
 - Retrain that model using all training data (e.g. polynomial coefficients)

Summary - Solutions II

- Use regularization
 - Train complex model (e.g high order polynomial) but penalize being “too complex” (e.g. large weight magnitudes)
 - Need to balance error vs. regularization (λ)
 - Choose λ using cross-validation
- Get more data

Bayesianity

- Frequentist view – probabilities are frequencies of random, repeatable events
- Bayesian view – probability quantifies uncertain beliefs
- Important distinction for us: Bayesianity allows us to discuss probability distributions over parameters (such as w)
 - Include priors (e.g. $p(w)$) over model parameters
- Later, we will see Bayesian approaches to combatting over-fitting and model selection for curve fitting
- For now, an illustrative example . . .

Coin Tossing

- Let's say you're given a coin, and you want to find out $P(H)$, the probability that if you flip it it lands as "heads".
- Flip it a few times: $H\ H\ T$
- $P(H) = 2/3$, no need for CMPT 726
- Hmm... is this rigorous? Does this make sense?

Coin Tossing - Model

- Bernoulli distribution $P(H) = \mu, P(T) = 1 - \mu$
- Assume coin flips are independent and identically distributed (i.i.d.)
 - i.e. All are separate samples from the Bernoulli distribution
- Given data $\mathcal{D} = \{x_1, \dots, x_N\}$, heads: $x_i = 1$, tails: $x_i = 0$, the likelihood of the data is:

$$p(\mathcal{D}|\mu) = \prod_{n=1}^N p(x_n|\mu) = \prod_{n=1}^N \mu^{x_n} (1-\mu)^{(1-x_n)}$$

Maximum Likelihood Estimation

- Given D with h heads and t tails
- What should μ be?
- Maximum Likelihood Estimation (MLE): choose μ which maximizes the likelihood of the data

$$\mu_{ML} = \arg \max_{\mu} p(\mathcal{D}|\mu)$$

- Since $\ln(\cdot)$ is monotone increasing:

$$\mu_{ML} = \arg \max_{\mu} \ln p(\mathcal{D}|\mu)$$

Maximum Likelihood Estimation

- Likelihood:

$$p(\mathcal{D}|\mu) = \prod_{n=1}^N \mu^{x_n} (1-\mu)^{(1-x_n)}$$

- Log-likelihood:

$$\ln p(\mathcal{D}|\mu) = \sum_{n=1}^N [x_n \ln \mu + (1-x_n) \ln(1-\mu)]$$

- Take derivative, set to 0:

$$\begin{aligned} \frac{d}{d\mu} \ln p(\mathcal{D}|\mu) &= \sum_{n=1}^N \left[x_n \frac{1}{\mu} - (1-x_n) \frac{1}{1-\mu} \right] = \frac{1}{\mu} h - \frac{1}{1-\mu} t \\ \Rightarrow \mu &= \frac{h}{t+h} \end{aligned}$$

Maximum Likelihood Estimation

- Likelihood:

$$p(\mathcal{D}|\mu) = \prod_{n=1}^N \mu^{x_n} (1-\mu)^{(1-x_n)}$$

- Log-likelihood:

$$\ln p(\mathcal{D}|\mu) = \sum_{n=1}^N [x_n \ln \mu + (1-x_n) \ln(1-\mu)]$$

- Take derivative, set to 0:

$$\begin{aligned} \frac{d}{d\mu} \ln p(\mathcal{D}|\mu) &= \sum_{n=1}^N \left[x_n \frac{1}{\mu} - (1-x_n) \frac{1}{1-\mu} \right] = \frac{1}{\mu} h - \frac{1}{1-\mu} t \\ \Rightarrow \mu &= \frac{h}{t+h} \end{aligned}$$

Maximum Likelihood Estimation

- Likelihood:

$$p(\mathcal{D}|\mu) = \prod_{n=1}^N \mu^{x_n} (1-\mu)^{(1-x_n)}$$

- Log-likelihood:

$$\ln p(\mathcal{D}|\mu) = \sum_{n=1}^N [x_n \ln \mu + (1-x_n) \ln(1-\mu)]$$

- Take derivative, set to 0:

$$\begin{aligned} \frac{d}{d\mu} \ln p(\mathcal{D}|\mu) &= \sum_{n=1}^N \left[x_n \frac{1}{\mu} - (1-x_n) \frac{1}{1-\mu} \right] = \frac{1}{\mu} h - \frac{1}{1-\mu} t \\ \Rightarrow \mu &= \frac{h}{t+h} \end{aligned}$$

Maximum Likelihood Estimation

- Likelihood:

$$p(\mathcal{D}|\mu) = \prod_{n=1}^N \mu^{x_n} (1-\mu)^{(1-x_n)}$$

- Log-likelihood:

$$\ln p(\mathcal{D}|\mu) = \sum_{n=1}^N [x_n \ln \mu + (1-x_n) \ln(1-\mu)]$$

- Take derivative, set to 0:

$$\begin{aligned} \frac{d}{d\mu} \ln p(\mathcal{D}|\mu) &= \sum_{n=1}^N \left[x_n \frac{1}{\mu} - (1-x_n) \frac{1}{1-\mu} \right] = \frac{1}{\mu} h - \frac{1}{1-\mu} t \\ \Rightarrow \mu &= \frac{h}{t+h} \end{aligned}$$

Maximum Likelihood Estimation

- Likelihood:

$$p(\mathcal{D}|\mu) = \prod_{n=1}^N \mu^{x_n} (1-\mu)^{(1-x_n)}$$

- Log-likelihood:

$$\ln p(\mathcal{D}|\mu) = \sum_{n=1}^N [x_n \ln \mu + (1-x_n) \ln(1-\mu)]$$

- Take derivative, set to 0:

$$\begin{aligned} \frac{d}{d\mu} \ln p(\mathcal{D}|\mu) &= \sum_{n=1}^N \left[x_n \frac{1}{\mu} - (1-x_n) \frac{1}{1-\mu} \right] = \frac{1}{\mu} h - \frac{1}{1-\mu} t \\ \Rightarrow \mu &= \frac{h}{t+h} \end{aligned}$$

Bayesian Learning

- Wait, does this make sense? What if I flip 1 time, heads?
Do I believe $\mu = 1$?
- Learn μ the Bayesian way:

$$P(\mu|\mathcal{D}) = \frac{P(\mathcal{D}|\mu)P(\mu)}{P(\mathcal{D})}$$

$$\underbrace{P(\mu|\mathcal{D})}_{\text{posterior}} \propto \underbrace{P(\mathcal{D}|\mu)}_{\text{likelihood}} \underbrace{P(\mu)}_{\text{prior}}$$

- Prior encodes knowledge that most coins are 50-50
- Conjugate prior makes math simpler, easy interpretation
 - For Bernoulli, the beta distribution is its conjugate

Bayesian Learning

- Wait, does this make sense? What if I flip 1 time, heads?
Do I believe $\mu = 1$?
- Learn μ the Bayesian way:

$$P(\mu|\mathcal{D}) = \frac{P(\mathcal{D}|\mu)P(\mu)}{P(\mathcal{D})}$$

$$\underbrace{P(\mu|\mathcal{D})}_{\text{posterior}} \propto \underbrace{P(\mathcal{D}|\mu)}_{\text{likelihood}} \underbrace{P(\mu)}_{\text{prior}}$$

- Prior encodes knowledge that most coins are 50-50
- Conjugate prior makes math simpler, easy interpretation
 - For Bernoulli, the beta distribution is its conjugate

Bayesian Learning

- Wait, does this make sense? What if I flip 1 time, heads?
Do I believe $\mu = 1$?
- Learn μ the Bayesian way:

$$P(\mu|\mathcal{D}) = \frac{P(\mathcal{D}|\mu)P(\mu)}{P(\mathcal{D})}$$

$$\underbrace{P(\mu|\mathcal{D})}_{\text{posterior}} \propto \underbrace{P(\mathcal{D}|\mu)}_{\text{likelihood}} \underbrace{P(\mu)}_{\text{prior}}$$


- Prior encodes knowledge that most coins are 50-50
- **Conjugate prior** makes math simpler, easy interpretation
 - For Bernoulli, the **beta** distribution is its conjugate

Beta Distribution

- We will use the Beta distribution to express our prior knowledge about coins:

$$\text{Beta}(\mu|a, b) = \frac{\Gamma(a + b)}{\underbrace{\Gamma(a)\Gamma(b)}_{\text{normalization}}} \mu^{a-1} (1 - \mu)^{b-1}$$

- Parameters a and b control the shape of this distribution



Posterior

$$\begin{aligned}
 P(\mu|\mathcal{D}) &\propto P(\mathcal{D}|\mu)P(\mu) \\
 &\propto \underbrace{\prod_{n=1}^N \mu^{x_n}(1-\mu)^{1-x_n}}_{\text{likelihood}} \underbrace{\mu^{(a-1)}(1-\mu)^{b-1}}_{\text{prior}} \\
 &\propto \mu^h(1-\mu)^t \mu^{(a-1)}(1-\mu)^{b-1} \\
 &\propto \mu^{h+a-1}(1-\mu)^{t+b-1}
 \end{aligned}$$

- Simple form for posterior is due to use of conjugate prior
- Parameters a and b act as extra observations
- Note that as $N = h + t \rightarrow \infty$, prior is ignored

Posterior

$$\begin{aligned}
 P(\mu|\mathcal{D}) &\propto P(\mathcal{D}|\mu)P(\mu) \\
 &\propto \underbrace{\prod_{n=1}^N \mu^{x_n}(1-\mu)^{1-x_n}}_{\text{likelihood}} \underbrace{\mu^{(a-1)}(1-\mu)^{b-1}}_{\text{prior}} \\
 &\propto \mu^h(1-\mu)^t \mu^{(a-1)}(1-\mu)^{b-1} \\
 &\propto \mu^{h+a-1}(1-\mu)^{t+b-1}
 \end{aligned}$$

- Simple form for posterior is due to use of conjugate prior
- Parameters a and b act as extra observations
- Note that as $N = h + t \rightarrow \infty$, prior is ignored

Posterior

$$\begin{aligned}
 P(\mu | \mathcal{D}) &\propto P(\mathcal{D} | \mu)P(\mu) \\
 &\propto \underbrace{\prod_{n=1}^N \mu^{x_n} (1-\mu)^{1-x_n}}_{\text{likelihood}} \underbrace{\mu^{(a-1)}(1-\mu)^{b-1}}_{\text{prior}} \\
 &\propto \mu^h (1-\mu)^t \mu^{(a-1)}(1-\mu)^{b-1} \\
 &\propto \mu^{h+a-1} (1-\mu)^{t+b-1}
 \end{aligned}$$

- Simple form for posterior is due to use of conjugate prior
- Parameters a and b act as extra observations
- Note that as $N = h + t \rightarrow \infty$, prior is ignored

Maximum A Posteriori

- Given posterior $P(\mu|\mathcal{D})$ we could compute a single value, known as the Maximum a Posteriori (MAP) estimate for μ :

$$\mu_{MAP} = \arg \max_{\mu} P(\mu|\mathcal{D})$$

- Known as point estimation

Bayesian Learning

- However, correct Bayesian thing to do is to use the full distribution over μ
 - i.e. Compute

$$\mathbb{E}_\mu[f] = \int P(\mu|\mathcal{D})f(\mu)d\mu$$

- This integral is usually hard to compute

Bayesian Learning

- However, correct Bayesian thing to do is to use the full distribution over μ
 - i.e. Compute

$$\mathbb{E}_\mu[f] = \int P(\mu|\mathcal{D})f(\mu)d\mu$$

- This integral is usually hard to compute

Conclusion

- Readings: Ch. 1.1-1.3, 2.1
- Types of learning problems
 - Supervised: regression, classification
 - Unsupervised
- Learning as optimization
 - Squared error loss function
 - Maximum likelihood (ML)
 - Maximum a posteriori (MAP)
- Want generalization, avoid over-fitting
 - Cross-validation
 - Regularization
 - Bayesian prior on model parameters