

# Linear Models for Classification

CMPT 726

Mo Chen

SFU Computing Science

Sept. 23, 2020

Bishop PRML Ch. 4

# Classification: Hand-written Digit Recognition

 $x_i =$  $t_i = (0, 0, 0, 1, 0, 0, 0, 0, 0, 0)$ 

- Each input vector classified into one of  $K$  discrete classes
  - Denote classes by  $\mathcal{C}_k$
- Represent input image as a vector  $x_i \in \mathbb{R}^{784}$ .
- We have target vector  $t_i \in \{0, 1\}^{10}$
- Given a **training set**  $\{(x_1, t_1), \dots, (x_N, t_N)\}$ , learning problem is to construct a “good” function  $y(x)$  from these.
  - $y: \mathbb{R}^{784} \rightarrow \mathbb{R}^{10}$

# Generalized Linear Models

- Similar to previous chapter on linear models for regression, we will use a “linear” model for classification:

$$y(\boldsymbol{x}) = f(\boldsymbol{w}^\top \boldsymbol{x} + w_0)$$

- This is called a **generalized linear model**
- $f(\cdot)$  is a fixed non-linear function
  - e.g.

$$f(u) = \begin{cases} 1, & \text{if } u \geq 0 \\ 0, & \text{otherwise} \end{cases}$$

- Decision boundary** between classes will be linear function of  $\boldsymbol{x}$
- Can also apply non-linearity to  $\boldsymbol{x}$ , as in  $\phi_i(\boldsymbol{x})$  for regression

# Outline

Discriminant Functions

Generative Models

Discriminative Models


# Outline

Discriminant Functions

Generative Models

Discriminative Models

# Discriminant Functions with Two Classes




- Start with 2 class problem,  
 $t_i \in \{0,1\}$
- Simple linear discriminant

$$y(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + w_0$$


apply threshold function to get  
classification

# Discriminant Functions with Two Classes



- $y(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + w_0$ 
  - Gradient of  $y$  is  $\mathbf{w}$
  - Constant  $y$  values  $\Rightarrow$  parallel lines
- If  $y = 0$  (decision boundary),  
 $\mathbf{w}^\top \mathbf{x} = -w_0 \Rightarrow \left(\frac{\mathbf{w}}{\|\mathbf{w}\|}\right)^\top \mathbf{x} = -\frac{w_0}{\|\mathbf{w}\|}$
- In general,  $\frac{y}{\|\mathbf{w}\|} = \frac{\mathbf{w}^\top \mathbf{x}}{\|\mathbf{w}\|} + \frac{w_0}{\|\mathbf{w}\|}$ , or  
 $\left(\frac{\mathbf{w}}{\|\mathbf{w}\|}\right)^\top \mathbf{x} = \frac{y(\mathbf{x})}{\|\mathbf{w}\|} - \frac{w_0}{\|\mathbf{w}\|}$


# Multiple Classes



- A linear discriminant between two classes separates with a hyperplane
- How to use this for multiple classes?
- **One-versus-the-rest** method: build  $K - 1$  classifiers, between  $\mathcal{C}_k$  and all others
- **One-versus-one** method: build  $K(K - 1)/2$  classifiers, between all pairs

# Multiple Classes

Given  $y_k(\mathbf{x}_A) > y_j(\mathbf{x}_A), y_k(\mathbf{x}_B) > y_j(\mathbf{x}_B)$



- A solution is to build  $K$  linear functions:

$$y_k(\mathbf{x}) = \mathbf{w}_k^\top \mathbf{x} + w_{k0}$$

assign  $\mathbf{x}$  to class  $\arg \max_k y_k(\mathbf{x})$

- Gives connected, convex **decision regions**

$$\hat{\mathbf{x}} = \lambda \mathbf{x}_A + (1 - \lambda) \mathbf{x}_B, \lambda \in [0,1]$$

$$y_k(\hat{\mathbf{x}}) = \lambda y_k(\mathbf{x}_A) + (1 - \lambda) y_k(\mathbf{x}_B)$$

$$y_j(\hat{\mathbf{x}}) = \lambda y_j(\mathbf{x}_A) + (1 - \lambda) y_j(\mathbf{x}_B)$$

$$\Rightarrow y_k(\hat{\mathbf{x}}) > y_j(\hat{\mathbf{x}}), \quad \forall j \neq k$$



# Least Squares for Classification

- How do we learn the decision boundaries ( $\mathbf{w}_k, w_{k0}$ )?
- One approach is to use least squares, similar to regression
- Find  $\mathbf{W}$  to minimize squared error over all examples and all components of the label vector:

$$E(\mathbf{W}) = \frac{1}{2} \sum_{n=1}^N \sum_{k=1}^K (y_k(\mathbf{x}_n) - t_{nk})^2$$

- Some algebra, we get a solution using the pseudo-inverse as in regression


## Problems with Least Squares



- Looks okay... least squares decision boundary
  - Similar to logistic regression decision boundary (more later)

- Gets worse by adding easy points?!
- Why?
  - If target value is 1, points far from boundary will have high value, say 10; this is a large error so the boundary is moved

## More Least Squares Problems



- Easily separated by hyperplanes, but not found using least squares!
- We'll address these problems later with better models
- First, a look at a different criterion for linear discriminant


# Fisher's Linear Discriminant

- The two-class linear discriminant acts as a projection:

$$y = \mathbf{w}^T \mathbf{x} + w_0$$

- followed by a threshold
- In which direction  $\mathbf{w}$  should we project?
- One which separates classes “well”




# Fisher's Linear Discriminant



- A natural idea would be to project in the direction of the line connecting class means
- However, problematic if classes have variance in this direction
- Fisher criterion: maximize ratio of inter-class separation (between) to intra-class variance (inside)

## Pre-Project List Available

- In Canvas under Files -> Project
  - Have a read, and decide on top 1, 2, and 3 choices
  - You may also propose your own project
- Sign up forms will be released in the next couple of days
  - Please complete rankings *before* next Monday



## Math time - FLD

- Projection  $y_n = \mathbf{w}^\top \mathbf{x}_n$
- Inter-class separation is distance between class means (good):

$$m_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{w}^\top \mathbf{x}_n$$

- Intra-class variance (bad):

$$s_k^2 = \sum_{n \in \mathcal{C}_k} (y_n - m_k)^2$$

- Fisher criterion:

$$J(\mathbf{w}) = \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2}$$

maximize wrt  $\mathbf{w}$

## Math time - FLD

$$J(\mathbf{w}) = \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2} = \frac{\mathbf{w}^\top S_B \mathbf{w}}{\mathbf{w}^\top S_W \mathbf{w}}$$

Between-class covariance:

$$S_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^\top$$

Within-class covariance:

$$S_W = \sum_{n \in \mathcal{C}_1} (\mathbf{x}_n - \mathbf{m}_1)(\mathbf{x}_n - \mathbf{m}_1)^\top + \sum_{n \in \mathcal{C}_2} (\mathbf{x}_n - \mathbf{m}_2)(\mathbf{x}_n - \mathbf{m}_2)^\top$$

Lots of math:

$$\mathbf{w} \propto S_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$

If covariance  $S_W$  is isotropic, reduces to class mean difference vector

# FLD Summary

- FLD is a dimensionality reduction technique
- Criterion for choosing projection based on class labels
  - Still suffers from outliers (e.g. earlier least squares example)

# Perceptrons

- **Perceptrons** is used to refer to many neural network structures (more coming up)
- The classic type is a fixed non-linear transformation of input, one layer of adaptive weights, and a threshold:

$$y(\mathbf{x}) = f(\mathbf{w}^\top \phi(\mathbf{x}))$$

- Developed by Rosenblatt in the 50s
- The main difference compared to the methods we've seen so far is the learning algorithm

# Perceptron Learning

- Two class problem
- For ease of notation, we will use  $t = 1$  for class  $\mathcal{C}_1$  and  $t = -1$  for class  $\mathcal{C}_2$ ; we choose  $f$  such that  $f(a) = 1$  if  $a \geq 0$  and  $f(a) = -1$  otherwise
- We saw that squared error was problematic
- Instead, we'd like to minimize the number of misclassified examples
  - An example is mis-classified if  $\mathbf{w}^\top \phi(\mathbf{x}_n) t_n < 0$
  - Perceptron criterion:

$$E_P(\mathbf{w}) = - \sum_{n \in \mathcal{M}} E_{P,n}(\mathbf{w}) = - \sum_{n \in \mathcal{M}} \mathbf{w}^\top \phi(\mathbf{x}_n) t_n$$

sum over mis-classified examples only

$$y(x) = f(w^T \phi(x)), \quad f(a) = \begin{cases} 1, & \text{if } a \geq 0 \\ -1, & \text{if } a < 0 \end{cases}$$

$w^T \phi(x_n) \geq 0$       predict: 1  
 $w^T \phi(x_n) < 0$       predict: -1

mistake if:  
 ground truth is  $\underline{-1} = t_n$   
 $\underline{1} = t_n$

# Perceptron Learning Algorithm


- Minimize the error function using stochastic gradient descent (gradient descent per example):

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_{P,n}(\mathbf{w}) = \mathbf{w}^{(\tau)} + \underbrace{\eta \phi(\mathbf{x}_n) t_n}_{\text{if incorrect}}$$




- Iterate over all training examples, only change  $\mathbf{w}$  if the example is mis-classified
- Guaranteed to converge if data are **linearly separable**
- Will not converge if not
- May take many iterations
- Sensitive to initialization

# Perceptron Learning Illustration

$w$ : dir'n of red class




$$\begin{aligned} w^{(\tau+1)} \\ = w^{(\tau)} + \eta \phi(x_n) t_n \end{aligned}$$



- Note there are many hyperplanes with 0 error
  - Support vector machines have a nice way of choosing one

# Limitations of Perceptrons

- Perceptrons can only solve linearly separable problems in feature space
  - Same as the other models in this chapter
- Canonical example of non-separable problem is X-OR
  - Real datasets can look like this too



# Outline

Discriminant Functions

Generative Models

Discriminative Models

# Probabilistic Generative Models

- Up to now we've looked at learning classification by choosing parameters to minimize an error function
- We'll now develop a probabilistic approach
- With 2 classes,  $\mathcal{C}_1$  and  $\mathcal{C}_2$ :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x})} \quad \text{Bayes' Rule}$$

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}, \mathcal{C}_1) + p(\mathbf{x}, \mathcal{C}_2)} \quad \text{Sum rule}$$

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad \text{Product rule}$$

- In **generative models** we specify the distribution  $p(\mathbf{x}|\mathcal{C}_k)$  which generates the data for each class

# Probabilistic Generative Models - Example

- Let's say we observe  $x$  which is the current temperature
- Determine if we are in Vancouver ( $\mathcal{C}_1$ ) or Honolulu ( $\mathcal{C}_2$ )
- Generative model:

$$p(\mathcal{C}_1|x) = \frac{p(x|\mathcal{C}_1)p(\mathcal{C}_1)}{p(x|\mathcal{C}_1)p(\mathcal{C}_1) + p(x|\mathcal{C}_2)p(\mathcal{C}_2)}$$

- $p(x|\mathcal{C}_1)$  is distribution over typical temperatures in Vancouver
    - e.g.  $p(x|\mathcal{C}_1) = \mathcal{N}(x; 10, 5)$
  - $p(x|\mathcal{C}_2)$  is distribution over typical temperatures in Honolulu
    - e.g.  $p(x|\mathcal{C}_2) = \mathcal{N}(x; 25, 5)$
  - Class priors  $p(\mathcal{C}_1) = 0.1, p(\mathcal{C}_2) = 0.9$
- 
- $p(\mathcal{C}_1|x=15) = \frac{0.0484 \times 0.1}{0.0484 \times 0.1 + 0.0108 \times 0.9} \approx 0.33$


# Generalized Linear Models

- We can write the classifier in another form

$$\begin{aligned} p(\mathcal{C}_1|\mathbf{x}) &= \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \\ &= \frac{1}{1 + \frac{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}} \\ &= \frac{1}{1 + \exp(-a)} \equiv \sigma(a) \end{aligned}$$

where  $a = \log \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)}$

# Logistic Sigmoid



- The function  $\sigma(a) = \frac{1}{1 + \exp(-a)}$  is known as the logistic sigmoid
- It squashes the real axis down to  $[0, 1]$
- It is continuous and differentiable
- It avoids the problems encountered with the *too correct* least-squares error fitting

# Multi-class Extension


- There is a generalization of the logistic sigmoid to  $K > 2$  classes:

$$\begin{aligned} p(\mathcal{C}_k | \mathbf{x}) &= \frac{p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k)}{\sum_j p(\mathbf{x}|\mathcal{C}_j)p(\mathcal{C}_j)} \\ &= \frac{\exp(a_k)}{\sum_j \exp(a_j)} \end{aligned}$$

where  $a_k = \log p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k)$

- a.k.a. softmax function
  - If some  $a_k \gg a_j$ ,  $p(\mathcal{C}_k | \mathbf{x})$  goes to 1

# Gaussian Class-Conditional Densities



- Back to that  $a$  in the logistic sigmoid for 2 classes
- Let's assume the class-conditional densities  $p(x|\mathcal{C}_k)$  are Gaussians, and have the same covariance matrix  $\Sigma$ :

$$p(x|\mathcal{C}_k) = \frac{1}{(2\pi)^{D/2} |\Sigma|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu_k)^\top \Sigma^{-1} (x - \mu_k) \right\}$$

- $a$  takes a simple form:

$$a = \log \frac{p(x|\mathcal{C}_1)p(\mathcal{C}_1)}{p(x|\mathcal{C}_2)p(\mathcal{C}_2)} = w^\top x + w_0$$

- Note that quadratic terms  $x^\top \Sigma^{-1} x$  cancel

# Maximum Likelihood Learning

- We can fit the parameters to this model using **maximum likelihood**
  - Parameters are  $\mu_1, \mu_2, \Sigma^{-1}, p(\mathcal{C}_1) \equiv \pi, p(\mathcal{C}_2) \equiv 1 - \pi$
  - Refer to as  $\theta$
- For a datapoint  $x_n$  from class  $\mathcal{C}_1$  ( $t_n = 1$ ):

$$p(x_n, \mathcal{C}_1) = p(\mathcal{C}_1)p(x_n | \mathcal{C}_1) = \pi \mathcal{N}(x_n | \mu_1, \Sigma)$$

- For a datapoint  $x_n$  from class  $\mathcal{C}_2$  ( $t_n = 0$ ):

$$p(x_n, \mathcal{C}_2) = p(\mathcal{C}_2)p(x_n | \mathcal{C}_2) = (1 - \pi) \mathcal{N}(x_n | \mu_2, \Sigma)$$

# Maximum Likelihood Learning

- The likelihood of the training data is:

$$p(\mathbf{t}|\pi, \boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \boldsymbol{\Sigma}) = \prod_{n=1}^N [\pi \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_1, \boldsymbol{\Sigma})]^{t_n} [(1 - \pi) \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_2, \boldsymbol{\Sigma})]^{1-t_n}$$

- As usual,  $\log(\cdot)$  is our friend:

$$l(\mathbf{t}|\theta) = \sum_{n=1}^N \frac{(t_n \underbrace{\log \pi + (1 - t_n) \log(1 - \pi)}_{\pi} + t_n \underbrace{\log \mathcal{N}_1 + (1 - t_n) \ln \mathcal{N}_2}_{\boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \boldsymbol{\Sigma}})}{\pi}$$

- Maximize for each separately

$$l(\boldsymbol{t}|\theta) = \sum_{n=1}^N \left( t_n \underbrace{\log \pi + (1-t_n) \log(1-\pi)}_{\pi} + t_n \underbrace{\log \mathcal{N}_1 + (1-t_n) \ln \mathcal{N}_2}_{\mu_1, \mu_2, \Sigma} \right)$$

$$\frac{\partial}{\partial \pi} l(t|\theta) = \sum_{n=1}^N \left[ \frac{t_n}{\pi} - \frac{1-t_n}{1-\pi} \right]$$

$$O = \sum_{n=1}^N \left( \frac{b_n}{\pi} \right) - \sum_{n=1}^N \left( \frac{1-b_n}{1-\pi} \right)$$

$$O = \frac{N_1}{\pi} - \frac{N_2}{1-\pi}$$

$$N_1(1-\pi) = N_2\pi$$

$$N_1 = N_2 \pi + N_1 \pi \Rightarrow$$

$$\pi = \frac{N_1}{N_1 + N_2}$$



# Maximum Likelihood Learning - Class Priors

- Maximization with respect to the class priors parameter  $\pi$  is straightforward:

$$\frac{\partial}{\partial \pi} l(\mathbf{t}|\theta) = \frac{\partial}{\partial \pi} \sum_{n=1}^N (t_n \log \pi + (1 - t_n) \ln(1 - \pi) + t_n \log \mathcal{N}_1 + (1 - t_n) \log \mathcal{N}_2)$$

$$0 = \sum_{n=1}^N \left( \frac{t_n}{\pi} - \frac{1 - t_n}{1 - \pi} \right)$$

$$0 = \frac{N_1}{\pi} - \frac{N_2}{1 - \pi} \quad \Rightarrow N_1 - N_1 \pi = N_2 \pi$$

$$\Rightarrow \pi = \frac{N_1}{N_1 + N_2}$$

# Maximum Likelihood Learning - Class Priors

- Maximization with respect to the class priors parameter  $\pi$  is straightforward:

$$\frac{\partial}{\partial \pi} l(t|\theta) = \sum_{n=1}^N \left( \frac{t_n}{\pi} - \frac{1-t_n}{1-\pi} \right)$$

$$\Rightarrow \pi = \frac{N_1}{N_1 + N_2}$$

- $N_1$  and  $N_2$  are the number of training points in each class
- Prior is simply the fraction of points in each class

# Maximum Likelihood Learning - Gaussian Parameters

- The other parameters can also be found in the same fashion
- Class means:

$$\mu_1 = \frac{1}{N_1} \sum_{n=1}^N t_n \mathbf{x}_n$$

$$\mu_2 = \frac{1}{N_2} \sum_{n=1}^N (1 - t_n) \mathbf{x}_n$$

- Means of training examples from each class
- Shared covariance matrix:

$$\Sigma = \frac{N_1}{N} \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} (\mathbf{x}_n - \boldsymbol{\mu}_1)(\mathbf{x}_n - \boldsymbol{\mu}_1)^T + \frac{N_2}{N} \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} (\mathbf{x}_n - \boldsymbol{\mu}_2)(\mathbf{x}_n - \boldsymbol{\mu}_2)^T$$

- Weighted average of class covariances

# Probabilistic Generative Models Summary

- Fitting Gaussian using ML criterion is sensitive to outliers
- Simple linear form for  $\alpha$  in logistic sigmoid occurs for more than just Gaussian distributions
  - Arises for any distribution in the **exponential family**, a large class of distributions

# Outline

Discriminant Functions

Generative Models

Discriminative Models

# Probabilistic Discriminative Models

- Generative model made assumptions about form of class-conditional distributions (e.g. Gaussian)
  - Resulted in logistic sigmoid of linear function of  $x$
- Discriminative model - explicitly use functional form

$$p(\mathcal{C}_1|x) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x} + w_0)}$$

and find  $\mathbf{w}$  directly

- For the generative model we had  $\underbrace{2M + M(M + 1)/2 + 1}_{\text{Means, variance, prior}}$  parameters
  - $M$  is dimensionality of  $x$
- Discriminative model will have  $M + 1$  parameters

# Generative vs. Discriminative

- Generative models
  - Can generate synthetic example data
  - Perhaps accurate classification is equivalent to accurate synthesis
    - e.g. vision and graphics
  - Tend to have more parameters
  - Require good model of class distributions
- Discriminative models
  - Only usable for classification
  - Don't solve a harder problem than you need to
  - Tend to have fewer parameters
  - Require good model of decision boundary

# Maximum Likelihood Learning - Discriminative Model

- As usual we can use the maximum likelihood criterion for learning

$$p(\mathbf{t}|\mathbf{w}) = \prod_{n=1}^N y_n^{t_n} \{1 - y_n\}^{1-t_n}, \text{ where } y_n = p(\mathcal{C}_1|\mathbf{x}_n)$$

- Taking  $\log(\cdot)$  and derivative gives:

$$\nabla l(\mathbf{w}) = \sum_{n=1}^N (t_n - y_n) \mathbf{x}_n$$

- This time no closed-form solution since  $y_n = \frac{1}{1+\exp(-\mathbf{w}^\top \mathbf{x}_n + w_0)}$
- Could use (stochastic) gradient descent
  - But there's a better iterative technique


# Iterative Reweighted Least Squares

- Iterative reweighted least squares (IRLS) is a descent method
  - As in [gradient descent](#), start with an initial guess, improve it
  - Gradient descent - take a step (how large?) in the gradient direction
- IRLS is a special case of a [Newton-Raphson](#) method
  - Approximate function using second-order Taylor expansion:

$$\hat{f}(\mathbf{w} + \mathbf{v}) = f(\mathbf{w}) + \nabla f(\mathbf{w})^\top (\mathbf{v} - \mathbf{w}) + \frac{1}{2} (\mathbf{v} - \mathbf{w})^\top H f(\mathbf{w}) (\mathbf{v} - \mathbf{w})$$

- Closed-form solution to minimize this is straight-forward: quadratic, derivatives linear
- In IRLS this second-order Taylor expansion ends up being a weighted least-squares problem, as in the regression case from last week
  - Hence the name IRLS

## Newton-Raphson



- Figure from Boyd and Vandenberghe, *Convex Optimization*
  - Excellent reference, free for download online  
<http://www.stanford.edu/~boyd/cvxbook/>

# Conclusion

- Readings: Ch. 4.1.1-4.1.4, 4.1.7, 4.2.1-4.2.2, 4.3.1-4.3.3
- Generalized linear models  $y(x) = f(\mathbf{w}^T \mathbf{x} + w_0)$
- Threshold/max function for  $f(\cdot)$ 
  - Minimize with least squares
  - Fisher criterion - class separation
  - Perceptron criterion - mis-classified examples
- Probabilistic models: logistic sigmoid / softmax for  $f(\cdot)$ 
  - Generative model - assume class conditional densities in exponential family; obtain sigmoid
  - Discriminative model - directly model posterior using sigmoid (a. k. a. **logistic regression**, though classification)
  - Can learn either using maximum likelihood
- All of these models are limited to linear decision boundaries in feature space