

# Neural Networks

CMPT 419/726

Mo Chen

SFU Computing Science

Jan. 29, 2020

Bishop PRML Ch. 5

# Neural Networks

- Neural networks arise from attempts to model human/animal brains
  - Many models, many claims of biological plausibility
- We will focus on **multi-layer perceptrons**
  - Mathematical properties rather than plausibility



# Applications of Neural Networks

- Many success stories for neural networks, old and new
  - Credit card fraud detection
  - Hand-written digit recognition
  - Face detection
  - Autonomous driving (CMU ALVINN)
  - Object recognition
  - Speech recognition

# Outline

Feed-forward Networks

Network Training

Error Backpropagation

Deep Learning

# Outline

Feed-forward Networks

Network Training

Error Backpropagation

Deep Learning

# Feed-forward Networks

- We have looked at generalized linear models of the form:

$$y(\mathbf{x}, \mathbf{w}) = f\left(\sum_{j=1}^M w_j \phi_j(\mathbf{x})\right)$$

for fixed non-linear basis functions  $\phi(\cdot)$

- We now extend this model by allowing adaptive basis functions, and learning their parameters
- In feed-forward networks (a.k.a. **multi-layer perceptrons**) we let each basis function be another non-linear function of linear combination of the inputs:

$$\phi_j(\mathbf{x}) = f\left(\sum_{j=1}^M \dots\right)$$

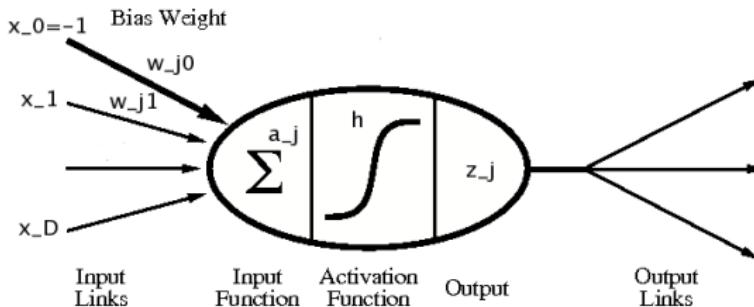
# Feed-forward Networks

- Starting with input  $x = (x_1, \dots, x_D)$ , construct linear combinations:

$$a_j = \sum_{i=1}^D \left( w_{ji}^{(1)} x_i + x_{j0}^{(1)} \right)$$

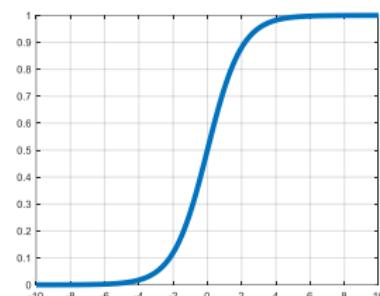
These  $a_j$  are known as activations

- Pass through an activation function  $h(\cdot)$  to get output  
$$z_j = h(a_j)$$
    - Model of an individual neuron



# Activation Functions

- Can use a variety of activation functions
  - Sigmoidal (S-shaped)
    - Logistic sigmoid  $1/(1 + \exp(-a))$  (useful for binary classification)
    - Hyperbolic tangent  $\tanh(\cdot)$
  - Radial basis function  $z_j = \sum_i (x_i - w_{ji})^2$
  - Softmax
    - Useful for multi-class classification
  - Identity
    - Useful for regression
  - Threshold
  - ...
- Needs to be differentiable for gradient-based learning (later)
- Can use different activation functions in each unit



# Activation Functions

Common choices of activation functions

Softplus:

$$\log(1 + e^x)$$

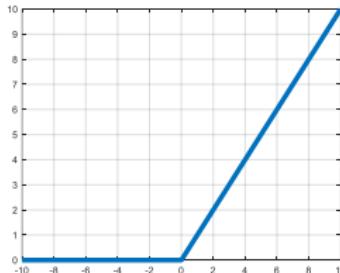
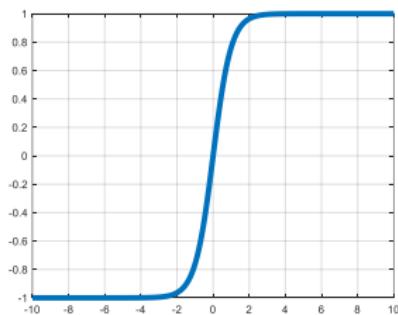
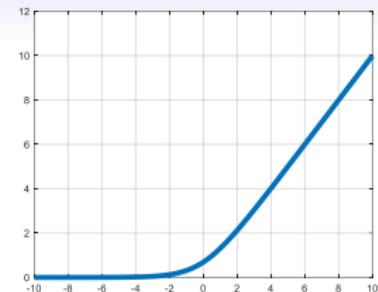
Hyperbolic tangent:

$$\tanh x$$

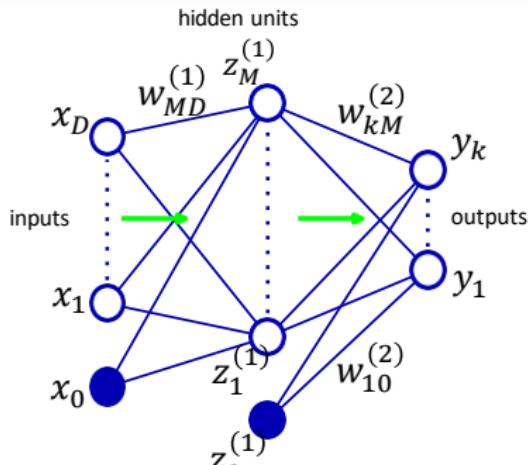
Rectified linear unit (ReLU):

$$\max(0, x)$$

Key feature: easy to differentiate



# Feed-forward Networks



- Connect together a number of these units into a feed-forward network (DAG)
- Above shows a network with one layer of **hidden units**
- Implements function

$$y_k(x, w) = h^{(2)} \left( \sum_{j=1}^M w_{kj}^{(2)} h^{(1)} \left( \sum_{i=1}^D w_{ij}^{(1)} x_i + w_{j0}^{(1)} \right) + w_{k0}^{(2)} \right)$$

# Outline

Feed-forward Networks

Network Training

Error Backpropagation

Deep Learning

## Network Training

- Given a specified network structure, how do we set its parameters (weights)?
  - As usual, we define a criterion to measure how well our network performs, optimize against it
- For regression, training data are  $(x_n, t_n), t_n \in \mathbb{R}$ 
  - Squared error naturally arises:

$$E(w) = \sum_{n=1}^N \{y(x_n, w) - t_n\}^2$$

- For binary classification, this is another discriminative model, ML:

$$p(t|w) = \prod_{n=1}^N y_n^{t_n} \{1 - y_n\}^{1-t_n}$$

$$E(w) = - \sum_{n=1}^N \{t_n \ln(y_n) + (1 - t_n) \ln(1 - y_n)\}$$

# Descent Methods

- Error function:

$$E(w) = \frac{1}{2} \sum_{n=1}^N \sum_k (y_{(n),k} - t_{(n),k})^2, \quad E_n(w) = \frac{1}{2} \sum_k (y_{(n),k} - t_{(n),k})^2$$

- $y(x, w)$  is a neural network, very complex
- Cannot solve  $\arg \min_w E(w)$  explicitly (like in linear regression)
- Gradient Descent:

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta^{(\tau)} \nabla E(w^{(\tau)})$$

- Stochastic Gradient Descent:

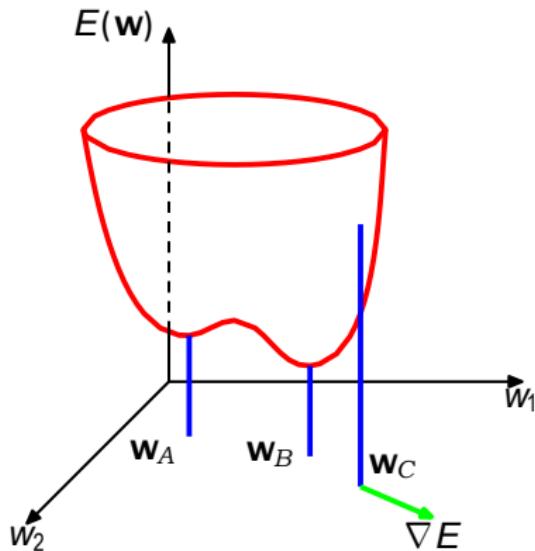
- $n$  chosen randomly

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta^{(\tau)} \nabla E_n(w^{(\tau)})$$

- A batch  $\mathcal{N}$  chosen randomly

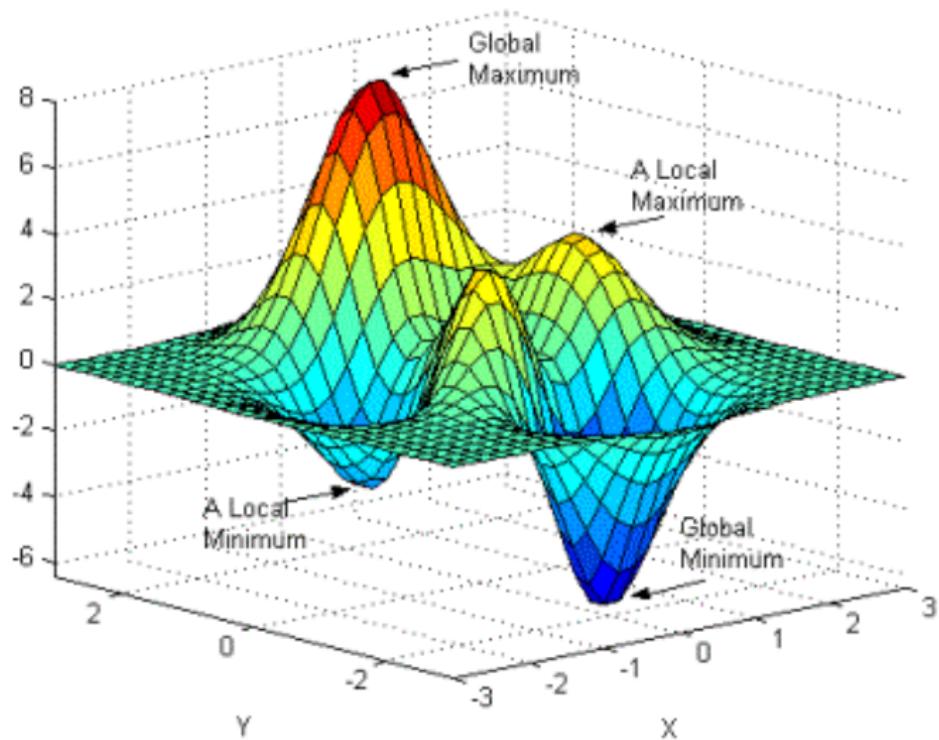
$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta^{(\tau)} \sum_{n \in \mathcal{N}} \nabla E_n(w^{(\tau)})$$

# Parameter Optimization



- For either of these problems, the error function  $E(\mathbf{w})$  is nasty
  - Nasty = non-convex
  - Non-convex = has **local minima**

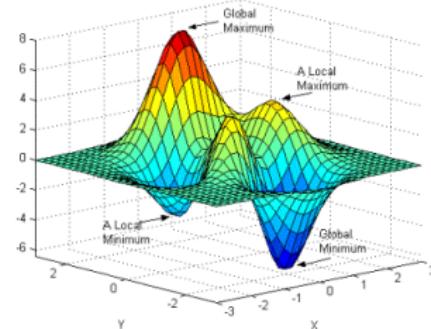
# A Non-Convex function



# Optimization Program

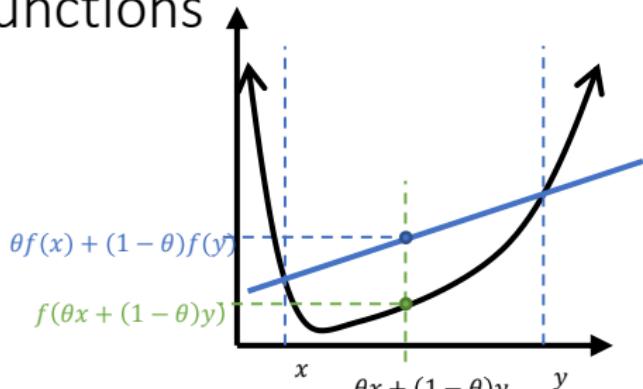
minimize  $f(x)$

subject to  $g_i(x) \leq 0, i = 1, \dots, n$   
 $h_j(x) = 0, j = 1, \dots, m$



- Very difficult to solve in general
  - Trade-offs to consider: computation time, solution optimality
  
- Easy cases:
  - Find global optimum for **linear program**:  $f, g_i, h_j$  are linear
  - Find global optimum for **convex program**:  $f, g_i$  are convex,  $h_j$  is linear
  - Find local optimum for **nonlinear program**:  $f, g_i, h_j$  are differentiable
  
- Neural Networks: Nonlinear and unconstrained

# Convex Functions



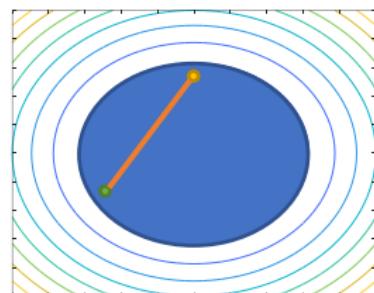
- **Convex function**

$f(\theta x + (1 - \theta)y) \leq \theta f(x) + (1 - \theta)f(y)$  for all  $x, y \in \mathbb{R}^n$ , for all  $\theta \in [0,1]$

- Sublevel sets of convex functions,  $\{x: f(x) \leq C\}$ , are convex

- **Convex shape  $\mathcal{C}$ :**

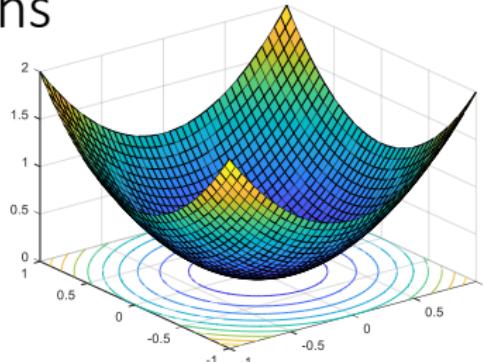
$$x_1, x_2 \in \mathcal{C}, \theta \in [0,1] \Rightarrow \theta x_1 + (1 - \theta)x_2 \in \mathcal{C}$$



# Convex Functions

- **Convex function**

$$f(\theta x + (1 - \theta)y) \leq \theta f(x) + (1 - \theta)f(y) \text{ for all } x, y \in \mathbb{R}^n, \text{ for all } \theta \in [0,1]$$

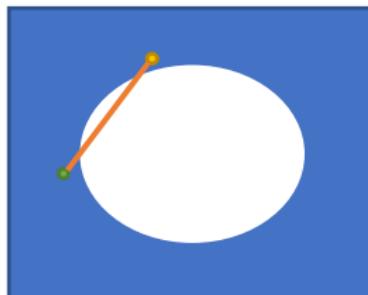


- Sublevel sets of convex functions,  $\{x: f(x) \leq C\}$ , are convex

- **Convex shape  $\mathcal{C}$ :**

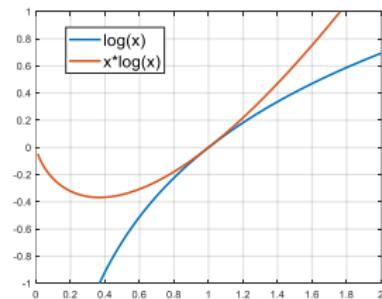
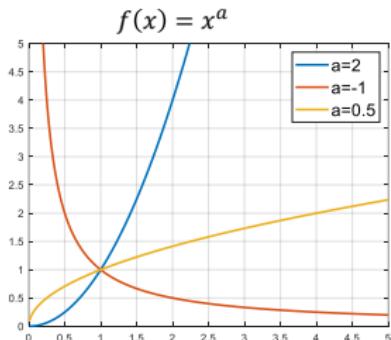
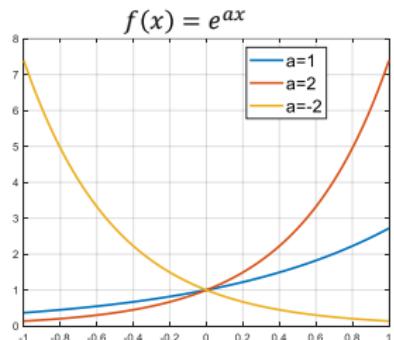
$$x_1, x_2 \in \mathcal{C}, \theta \in [0,1] \Rightarrow \theta x_1 + (1 - \theta)x_2 \in \mathcal{C}$$

- *Superlevel sets of convex functions are not convex!*



# Common Convex Functions on $\mathbb{R}$

- $f(x) = e^{ax}$  is convex for all  $x, a \in \mathbb{R}$
- $f(x) = x^a$  is convex on  $x > 0$  if  $a \geq 1$  or  $a \leq 0$ ; concave if  $0 < a < 1$
- $f(x) = \log x$  is concave
- $f(x) = x \log x$  is convex for  $x > 0$  (or  $x \geq 0$  if defined to be 0 when  $x = 0$ )



# Common Convex Functions on $\mathbb{R}^n$

- $f(x) = Ax + b$  is convex for any  $A, b$

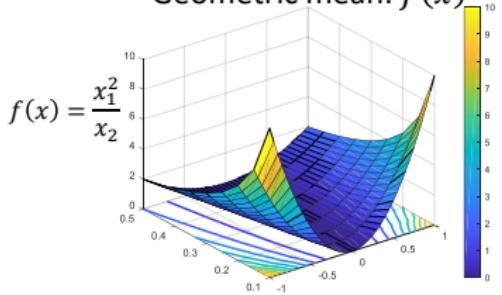
- Every norm on  $\mathbb{R}^n$  is convex

- $f(x) = \max(x_1, x_2, \dots, x_n)$  is convex

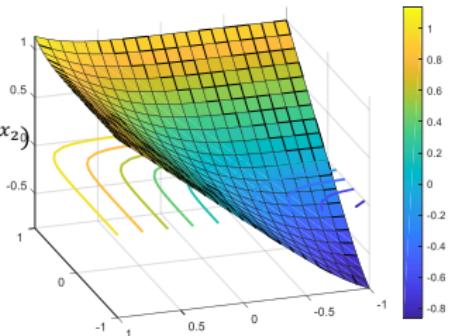
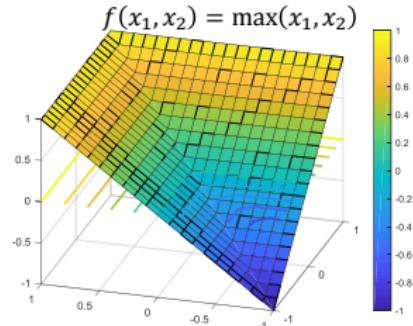
- $f(x) = \frac{x_1^2}{x_2}$  (for  $x_2 > 0$ )

- Log-sum-exp softmax:  $f(x) = \frac{1}{k} \log(e^{kx_1} + e^{kx_2} + \dots + e^{kx_n})$

- Geometric mean:  $f(x) = (\prod_{i=1}^n x_i)^{\frac{1}{n}}$ ,  $x_i > 0$



$$f(x) = \frac{1}{5} \log(e^{5x_1} + e^{5x_2})$$



# Descent Methods

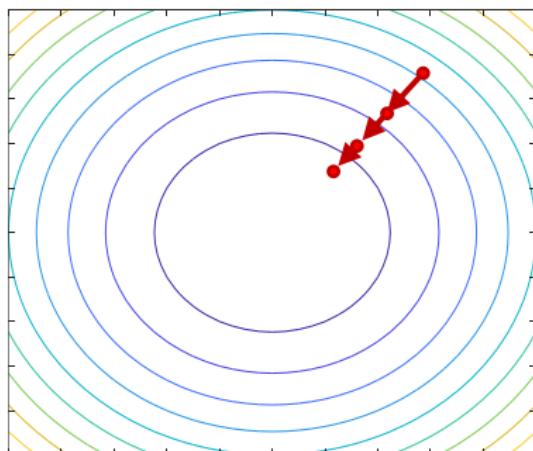
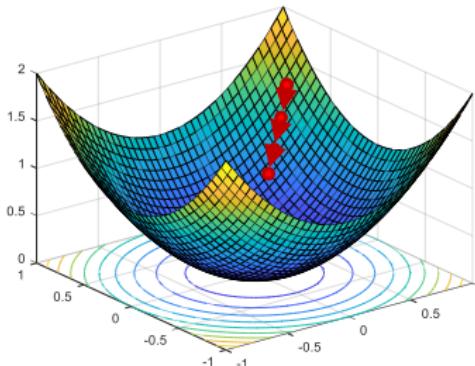
- The typical strategy for optimization problems of this sort is a descent method:

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \Delta\mathbf{w}^{(\tau)}$$

- As we've seen before, these come in many flavours
  - Gradient descent  $\nabla E(\mathbf{w}^{(\tau)})$
  - Stochastic gradient descent  $\nabla E_n(\mathbf{w}^{(\tau)})$
  - Newton-Raphson (second order)
- All of these can be used here, stochastic gradient descent is particularly effective
  - Redundancy in training data, escaping local minima

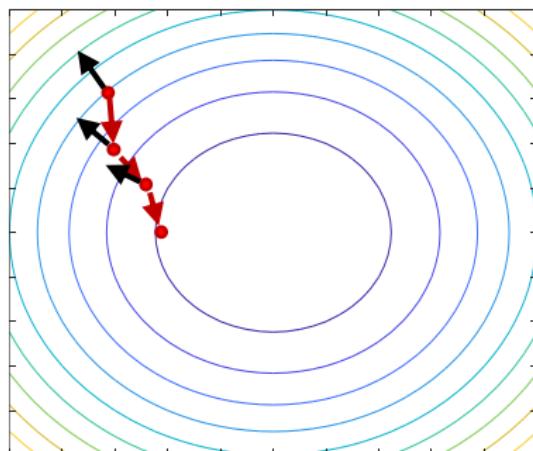
# Numerical Solution: Gradient Methods

- Start from  $x^0$  and construct a sequence  $x^k$  such that  $x^k \rightarrow x^*$ 
  - Calculate  $x^{k+1}$  from  $x^k$  by “going down the gradient”
  - Unconstrained case:  $x^{k+1} = x^k - \alpha^k \nabla f(x)$ ,  $\alpha^k > 0$



# Numerical Solution: Gradient Methods

- Start from  $x^0$  and construct a sequence  $x^k$  such that  $x^k \rightarrow x^*$ 
  - Calculate  $x^{k+1}$  from  $x^k$  by “going down the gradient”
  - Unconstrained case:  $x^{k+1} = x^k - \alpha^k \nabla f(x)$ ,  $\alpha^k > 0$
- More generally,  $x^{k+1} = x^k + \alpha^k d^k$  for some  $d$  such that
$$\nabla f(x^k)^\top d^k < 0$$
- Tuning parameters: descent direction  $d^k$ , and step size  $\alpha^k$



# Descent Direction

- Steepest descent:  $d^k = -\nabla f(x^k)$ 
  - $x^{k+1} = x^k - \alpha^k \nabla f(x)$
  - Simple but sometimes leads to slow convergence

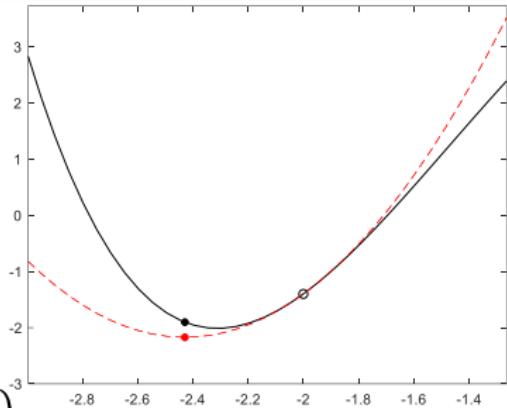
- Newton's method:  $d^k = (\nabla^2 f(x^k))^{-1} \nabla f(x^k)$ 
  - Minimize the quadratic approximation:

$$f^k(x) = f(x^k) + \nabla f(x^k)^T (x - x^k) + \frac{1}{2} (x - x^k)^T H f(x^k) (x - x^k)$$

- Set gradient to zero to obtain next iterate
 
$$\nabla f^k(x) = \nabla f(x^k) + H f(x^k) (x - x^k) = 0$$

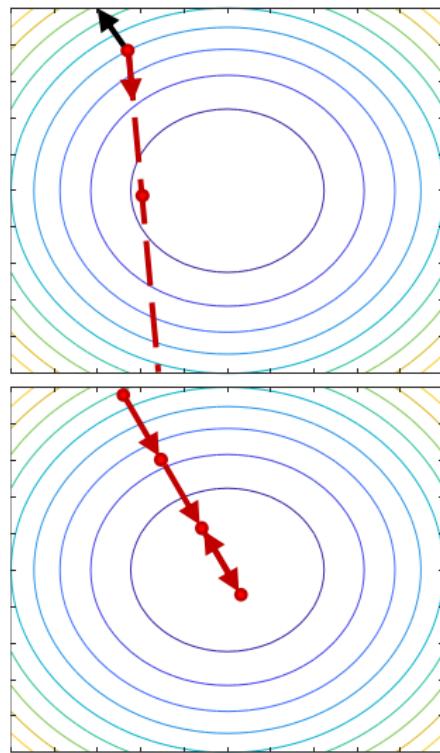
$$\Rightarrow x^{k+1} = x^k - (H f(x^k))^{-1} \nabla f(x^k)$$

- Fast convergence, but matrix inverse required
- Alternatively, use an algorithm to minimize a quadratic function



# Step Size (Learning rate)

- Recall  $x^{k+1} = x^k + \alpha^k d^k$ , with  $\nabla f(x^k)^\top d^k < 0$
- Line search: choose  $\alpha^k = \min_{\alpha \geq 0} f(x^k + \alpha^k d^k)$ 
  - Requires minimization
- Constant step size:  $\alpha^k = \alpha$ 
  - May not converge
- Diminishing step size:  $\alpha^k \rightarrow 0$ 
  - Still need to explore all regions  $\sum \alpha^k = \infty$
  - For example:  $\alpha^k = \frac{\alpha^0}{k}$



# Numerical Solution: Second Order Methods

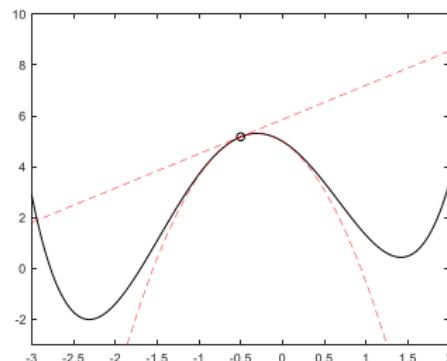
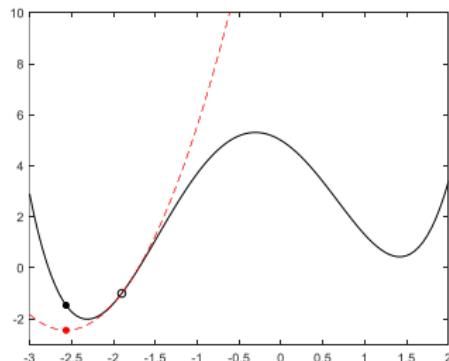
$$\text{minimize } f(x) \quad \longrightarrow \quad \text{minimize}_{d_x} (\mathbf{r}^k)^T d_x + \frac{1}{2} d_x^T \mathbf{B}_k d_x$$

where  $d_x := x - x^k$ ,

- Quadratize  $f(x)$ :

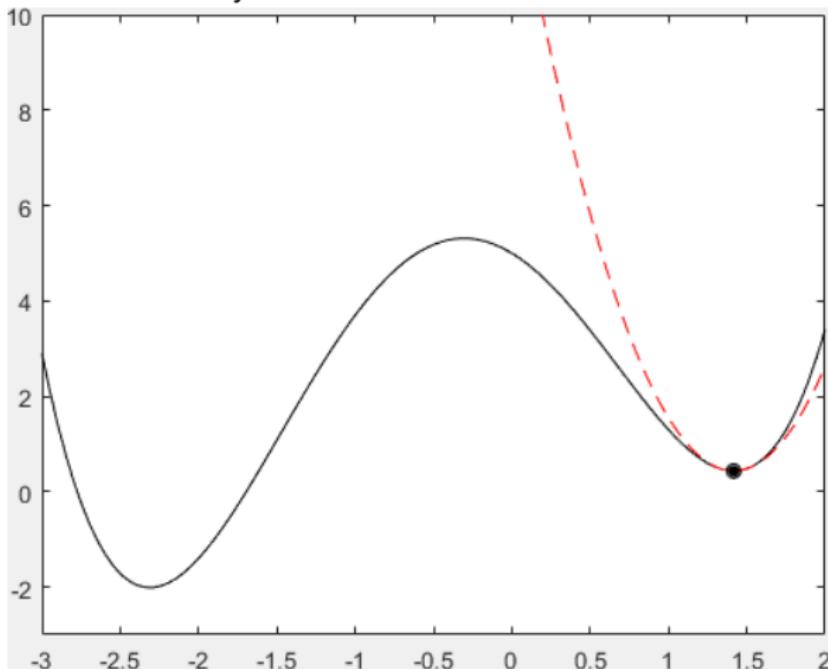
$$\begin{aligned} \mathbf{r}^k &= \nabla f(x_k) \\ \mathbf{B}_k &= \mathbf{H}f(x_k) \end{aligned}$$

- Convexify if needed, eg. by removing negative eigenvalues



# Example

minimize  $0.5x^4 + 0.8x^3 - 3x^2 - 2x + 5$   
subject to  $-3 \leq x \leq 2$



# Computing Gradients

- The function  $y(\mathbf{x}_n, \mathbf{w})$  implemented by a network is complicated
  - It isn't obvious how to compute error function derivatives with respect to weights
- Numerical method for calculating error derivatives, use finite differences:

$$\frac{\partial E_n}{\partial w_{ji}} \approx \frac{E_n(w_{ji} + \epsilon) - E_n(w_{ji} - \epsilon)}{2\epsilon}$$

- How much computation would this take with  $W$  weights in the network?
  - $O(|W|)$  per partial derivative (evaluation of  $E_n$ )
  - $O(|W|^2)$  total per gradient descent step (there are  $|W|$  partial derivatives)

# Outline

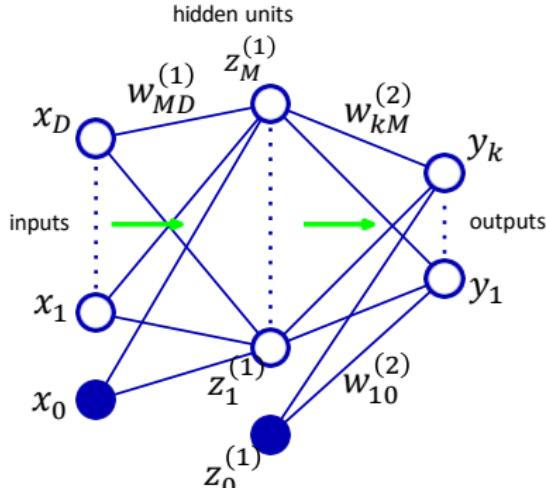
Feed-forward Networks

Network Training

Error Backpropagation

Deep Learning

# Feed-forward Networks

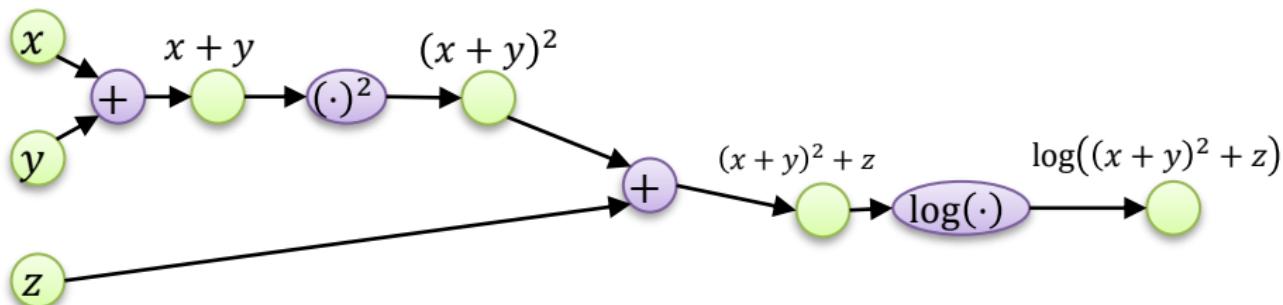


- Connect together a number of these units into a feed-forward network (DAG)
- Above shows a network with one layer of **hidden units**
- Implements function:

$$y_{(n),k}(x_n, w) = h^{(2)} \left( \sum_{j=1}^M w_{kj}^{(2)} h^{(1)} \left( \underbrace{\sum_{i=1}^D w_{ji}^{(1)} x_{(n),i} + w_{j0}^{(1)}}_{Z_{(n),j}} \right) + w_{k0}^{(2)} \right)$$

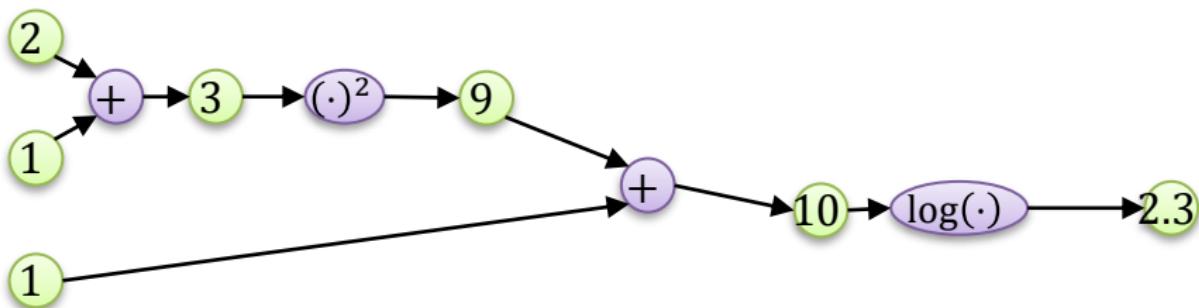
# Computation Graphs

- Consider the function  $f(x, y, z) = \log((x + y)^2 + z)$



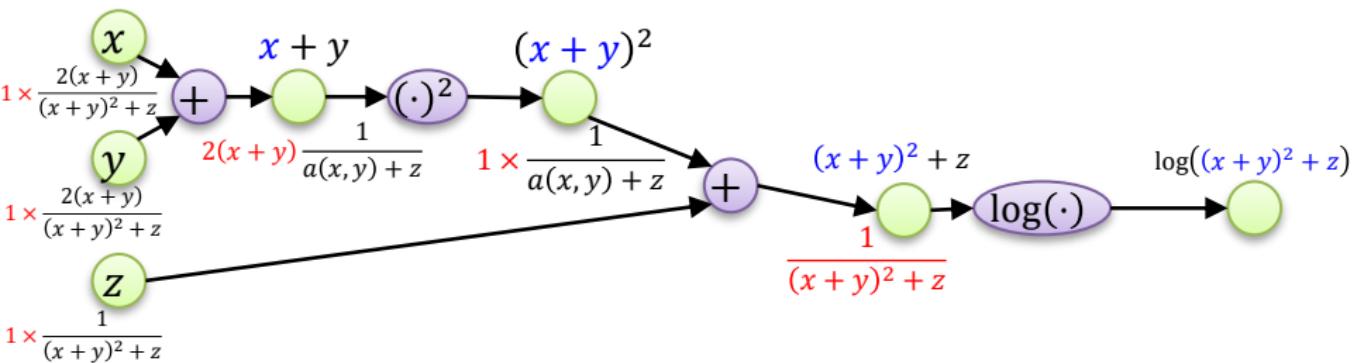
# Computation Graphs

- Consider the function  $f(x, y, z) = \log((x + y)^2 + z)$



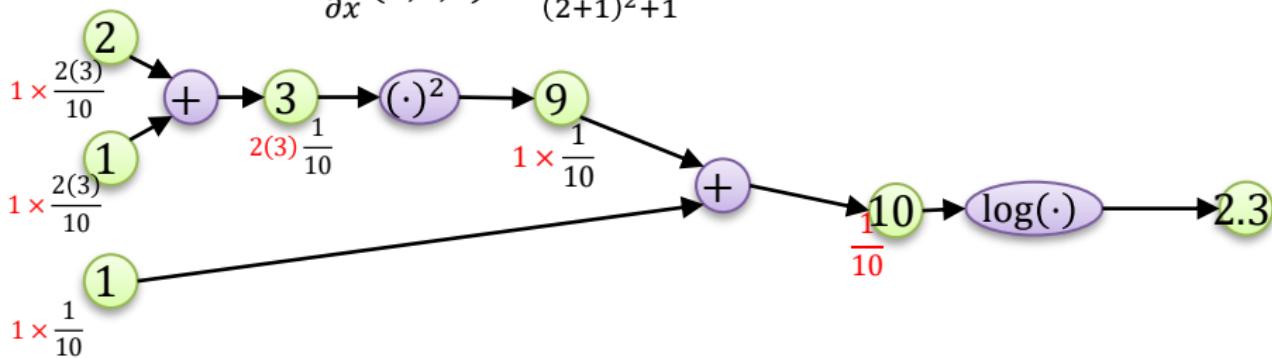
# Computation Graphs

- Consider the function  $f(x, y, z) = \log((x + y)^2 + z)$
- Gradients: let  $a(x, y) = (x + y)^2, b(s) = \log s$ 
  - $f(x, y, z) = b(a(x, y) + z)$
  - $\frac{\partial f}{\partial x}(x, y, z) = \frac{\partial b}{\partial x}(x, y, z) \frac{\partial a}{\partial x}(x, y, z) = \frac{1}{(x+y)^2+z} 2(x+y)$



# Computation Graphs

- Consider the function  $f(x, y, z) = \log((x + y)^2 + z)$
- Gradients: let  $a(x, y) = (x + y)^2, b(s) = \log s$ 
  - $f(x, y, z) = b(a(x, y) + z)$
  - $\frac{\partial f}{\partial x}(x, y, z) = \frac{\partial b}{\partial x}(x, y, z) \frac{\partial a}{\partial x}(x, y, z) = \frac{1}{(x+y)^2+z} 2(x+y)$
  - $\frac{\partial f}{\partial x}(2, 1, 1) = \frac{2(2+1)}{(2+1)^2+1} = 0.6$



# Error Backpropagation

- Backprop is an efficient method for computing error derivatives  $\frac{\partial E_n}{\partial w_{ji}^{(m)}}$ 
  - $O(W)$  to compute derivatives wrt all weights
- First, feed training example  $x_n$  forward through the network, storing all activations  $a_j$
- Calculating derivatives for weights connected to output nodes is easy
  - e.g. For linear output nodes  $y_k = \sum_i w_{ki}^{(L)} z_{(n),i}^{(L-1)}$ :
$$\frac{\partial E_n}{\partial w_{ki}^{(L)}} = \frac{\partial}{\partial w_{ki}^{(L)}} \frac{1}{2} (y_{(n),k} - t_{(n),k})^2 = (y_{(n),k} - t_{(n),k}) z_{(n),i}^{(L-1)}$$
- For hidden layers, propagate error backwards from the output nodes

# Chain Rule for Partial Derivatives

- A “reminder”
- For  $f(x, y)$ , with  $f$  differentiable wrt  $x$  and  $y$ , and  $x$  and  $y$  differentiable wrt  $u$ :

$$\frac{\partial f}{\partial u} = \frac{\partial f}{\partial x} \frac{\partial x}{\partial u} + \frac{\partial f}{\partial y} \frac{\partial y}{\partial u}$$

# Error Backpropagation

$y_{(n),k}, E_n$ :

- $n$ : data point
  - $k$ : component

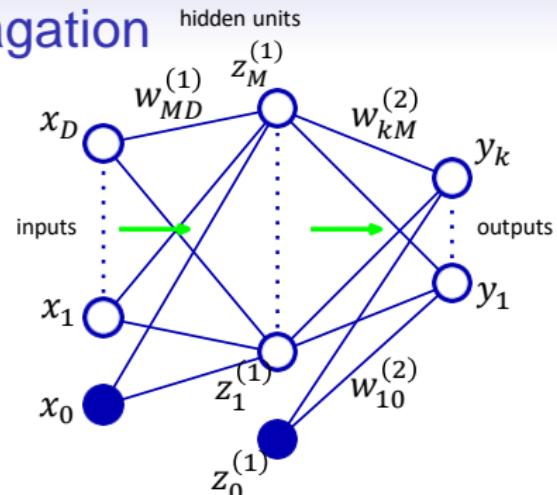
$$w_{jj}^{(m)}.$$

- $m$ : layer
  - $j$ : index matching output
  - $i$ : index matching input

$$E(w) = \frac{1}{2} \sum_{n=1}^N \sum_k (y_{(n),k} - t_{(n),k})^2, \quad y_{(n),k} = \sum_i w_{ki}^{(L)} z_{(n),i}^{(L-1)}$$

$$E_n(w) = \frac{1}{2} \sum_k (y_{(n),k} - t_{(n),k})^2$$

$$\frac{\partial E_n}{\partial w_{ki}^{(L)}} = \frac{\partial}{\partial w_{ki}^{(L)}} \frac{1}{2} \sum_{k'} (y_{(n),k'} - t_{(n),k'})^2 = (y_{(n),k} - t_{(n),k}) z_{(n),i}^{(L-1)} \quad (*)$$



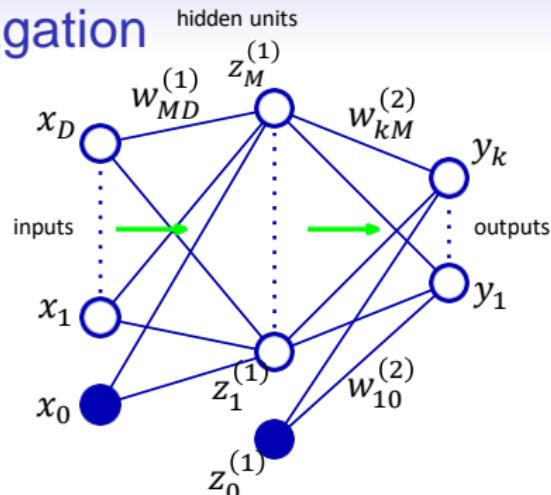
# Error Backpropagation

- We can write

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \frac{\partial}{\partial w_{ji}^{(m)}} E_n \left( a_{(n),1}^{(m)}, a_{(n),2}^{(m)}, \dots, a_{(n),D}^{(m)} \right)$$

- Using the chain rule:

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \frac{\partial E_n}{\partial a_{(n),j}^{(m)}} \frac{\partial a_{(n),j}^{(m)}}{\partial w_{ji}^{(m)}} + \sum_{k \neq j} \frac{\partial E_n}{\partial a_{(n),k}^{(m)}} \frac{\partial a_{(n),k}^{(m)}}{\partial w_{ji}^{(m)}}$$



where  $\sum_k (\dots)$  runs over all other nodes  $k$  in the same layer ( $m$ )

- Since  $a_{(n),k}^{(m)}$  does not depend on  $w_{ji}^{(m)}$ , all terms in the summation go to 0:

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \frac{\partial E_n}{\partial a_{(n),j}^{(m)}} \frac{\partial a_{(n),j}^{(m)}}{\partial w_{ji}^{(m)}}$$

## Error Backpropagation cont.

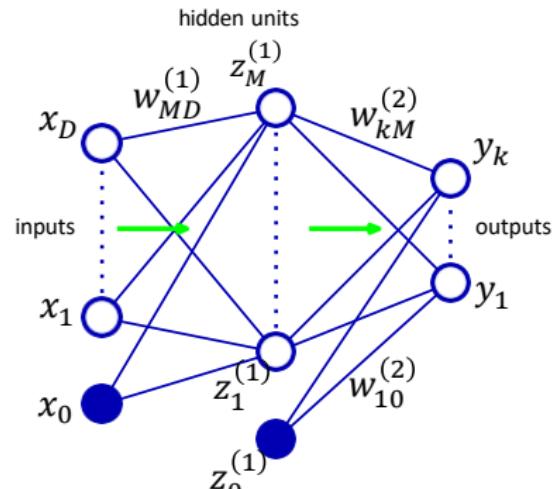
- Introduce error  $\delta_{(n),j}^{(m)} := \frac{\partial E_n}{\partial a_{(n),j}^{(m)}}$

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \delta_{(n),j}^{(m)} \frac{\partial a_{(n),j}^{(m)}}{\partial w_{ji}^{(m)}}$$

- Other factor is

$$\frac{\partial a_{(n),j}^{(m)}}{\partial w_{ji}^{(m)}} = \frac{\partial}{\partial w_{ji}^{(m)}} \sum_k w_{jk}^{(m)} z_k^{(m-1)} = z_0^{(m-1)}$$

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \delta_{(n),j}^{(m)} z_i^{(m-1)}$$



## Error Backpropagation cont.

- Error  $\delta_{(n),j}^{(m)}$  can also be computed using chain rule:

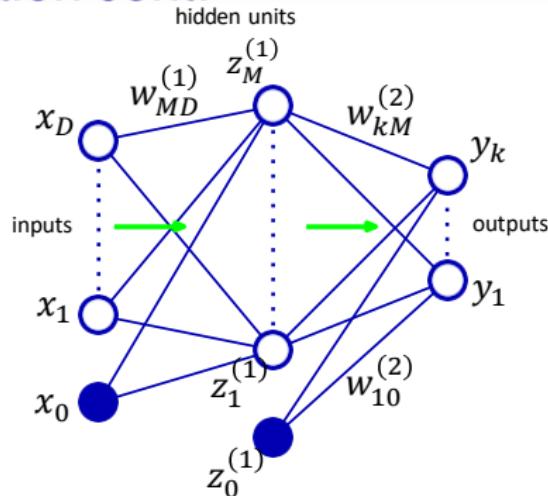
$$\delta_{(n),j}^{(m)} := \frac{\partial E_n}{\partial a_{(n),j}^{(m)}} = \sum_k \underbrace{\frac{\partial E_n}{\partial a_{(n),k}^{(m+1)}}}_{\delta_k^{(m+1)}} \frac{\partial a_{(n),k}^{(m+1)}}{\partial a_{(n),j}^{(m)}}$$

where  $\sum_k(\dots)$  runs over all nodes  $k$  in the layer after.

$$a_{(n),k}^{(m+1)} = \sum_i w_{ki}^{(m+1)} z_{(n),i}^{(m)} = \sum_i w_{ki}^{(m+1)} h^{(m)}\left(a_{(n),i}^{(m)}\right)$$

$$\frac{\partial a_{(n),k}^{(m+1)}}{\partial a_{(n),j}^{(m)}} = w_{kj}^{(m+1)} \left( h^{(m)} \right)' \left( a_{(n),j}^{(m)} \right)$$

$$\delta_{(n),j}^{(m)} = \sum_k \delta_{(n),k}^{(m+1)} w_{kj}^{(m+1)} (h^{(m)})' \left( a_{(n),j}^{(m)} \right) = (h^{(m)})' \left( a_{(n),j}^{(m)} \right) \sum_k \delta_{(n),k}^{(m+1)} w_{kj}^{(m+1)}$$



## Error Backpropagation cont.

- Error  $\delta_{(n),j}^{(m)}$  can also be computed using chain rule:

$$\delta_{(n),j}^{(m)} := \frac{\partial E_n}{\partial a_{(n),j}^{(m)}} = \sum_k \underbrace{\frac{\partial E_n}{\partial a_{(n),k}^{(m+1)}}}_{\delta_k} \frac{\partial a_{(n),k}^{(m+1)}}{\partial a_{(n),j}^{(m)}}$$

where  $\sum_k(\dots)$  runs over all nodes  $k$  in the layer **after**.

- Eventually:

$$\delta_{(n),j}^{(m)} = (h^{(m)})' \left( a_{(n),j}^{(m)} \right) \sum_k \delta_{(n),k}^{(m+1)} w_{kj}^{(m+1)}$$

- A weighted sum of the later error “caused” by this weight

## Error Backpropagation cont.

- Eventually:

$$\delta_{(n),j}^{(m)} = \left( h^{(m)} \right)' \left( a_{(n),j}^{(m)} \right) \sum_k \delta_{(n),k}^{(m+1)} w_{jk}^{(m+1)}$$

where  $\sum_k (\dots)$  runs over  $k$  all nodes  $k$  in the layer **after**.

- Above recursion relation needs last set of errors:  $\delta_j^{(L)}$

$$\frac{\partial E_n}{\partial w_{ji}^{(m)}} = \delta_{(n),j}^{(m)} z_i^{(m-1)} \quad (\text{by definition})$$

$$\frac{\partial E_n}{\partial w_{ji}^{(L)}} = \delta_{(n),j}^{(L)} z_{(n),i}^{(L-1)} = (y_{(n),j} - t_{(n),j}) z_{(n),i}^{(L-1)} \quad (\text{from before } *)$$

$$\delta_{(n),j}^{(L)} = y_{(n),j} - t_{(n),j} \quad (\text{by comparison})$$

# Summary

$O(W)$

## Output Definition / forward propagation

$$y_{(n),k}(x_n, w) = h^{(m+1)} \left( \sum_{j=1}^M w_{jk}^{(m+1)} h^{(m)} \left( \sum_{i=1}^D w_{ij}^{(m)} z_{(n),i}^{(m-1)} + w_{0j}^{(m)} \right) + w_{k0}^{(m+1)} \right)$$

- Save  $\mathbf{z}$ ,  $\mathbf{a}$

## Gradient computation / backpropagation

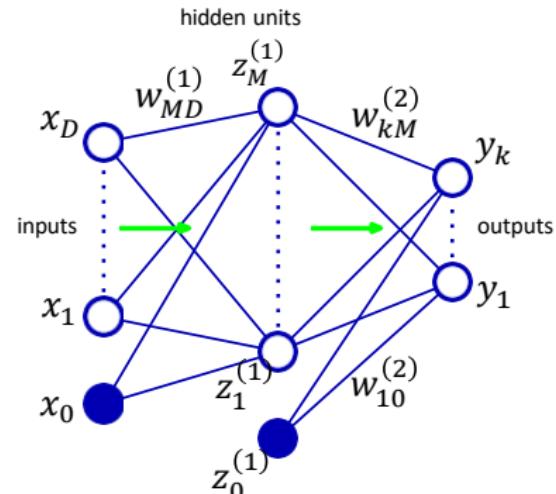
- Last layer:  $\frac{\partial E_n}{\partial w_{ik}^{(L)}} = (y_{(n),k} - t_{(n),k}) z_{(n),i}^{(L-1)}$
- Previous layers: Define  $\delta_{(n),j}^{(m)} := \frac{\partial E_n}{\partial a_{(n),j}^{(m)}}$

Starting from last layer,

$$\delta_{(n),j}^{(L)} = y_{(n),j} - t_{(n),j}$$

Recursion:  $\frac{\partial E_n}{\partial w_{ij}^{(m)}} = \delta_{(n),j}^{(m)} z_{(n),i}^{(m-1)},$

where  $\delta_{(n),j}^{(m)} = (h^{(m)})' (a_{(n),j}^{(m)}) \sum_k \delta_k^{(m+1)} w_{jk}^{(m+1)}$



# Summary

## Output Definition / forward propagation

$$y_{(n),k}(x_n, w) = h^{(m+1)} \left( \sum_{j=1}^M w_{jk}^{(m+1)} h^{(m)} \left( \sum_{i=1}^D w_{ij}^{(m)} z_{(n),i}^{(m-1)} + w_{0j}^{(m)} \right) + w_{k0}^{(m+1)} \right)$$

$O(W)$

- Save  $\mathbf{z}$ ,  $\mathbf{a}$

## Gradient computation / backpropagation

- Last layer:  $\frac{\partial E_n}{\partial w_{ik}^{(L)}} = (y_{(n),k} - t_{(n),k}) z_{(n),i}^{(L-1)}$
  - Previous layers: Define  $\delta_{(n),j}^{(m)} := \frac{\partial E_n}{\partial a_{(n),j}^{(m)}}$
- Starting from last layer,

$$\delta_{(n),j}^{(L)} = y_{(n),j} - t_{(n),j}$$

Recursion:  $\frac{\partial E_n}{\partial w_{ij}^{(m)}} = \delta_{(n),j}^{(m)} z_{(n),i}^{(m-1)},$       Goes through one layer of weights

where  $\delta_{(n),j}^{(m)} = (h^{(m)})' (a_{(n),j}^{(m)}) \sum_k \delta_k^{(m+1)} w_{jk}^{(m+1)}$

# Tensorflow Playground

- <https://playground.tensorflow.org>

# Outline

Feed-forward Networks

Network Training

Error Backpropagation

Deep Learning

# Deep Learning

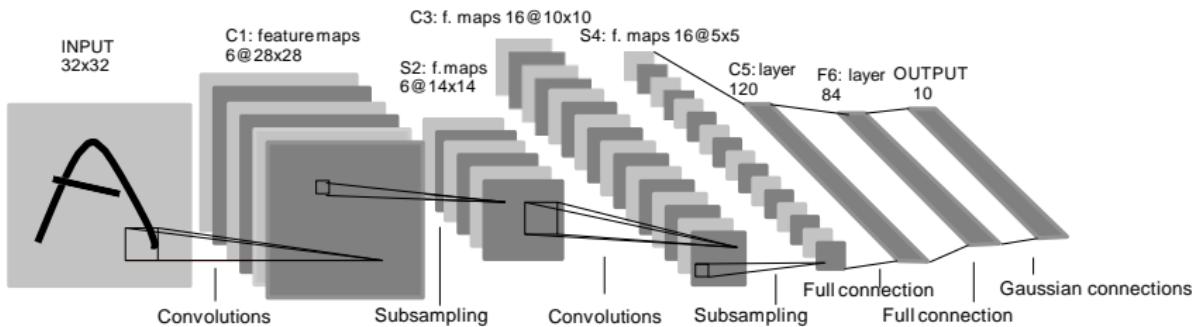
- Collection of important techniques to improve performance:
  - Multi-layer networks
  - Convolutional networks, parameter tying
  - Hinge activation functions (ReLU) for steeper gradients
  - Momentum
  - Drop-out regularization
  - Sparsity
  - Auto-encoders for unsupervised feature learning
  - ...
- **Scalability** is key, can use lots of data since stochastic gradient descent is memory-efficient, can be parallelized

# Hand-written Digit Recognition

3 6 8 1 7 9 6 6 9 1  
6 7 5 7 8 6 3 4 8 5  
2 1 7 9 7 1 2 8 4 6  
4 8 1 9 0 1 8 8 9 4  
7 6 1 8 6 4 1 5 6 0  
7 5 9 2 6 5 8 1 9 7  
2 2 2 2 2 3 4 4 8 0  
0 2 3 8 0 7 3 8 5 7  
0 1 4 6 4 6 0 2 4 3  
7 1 2 8 7 6 9 8 6 1

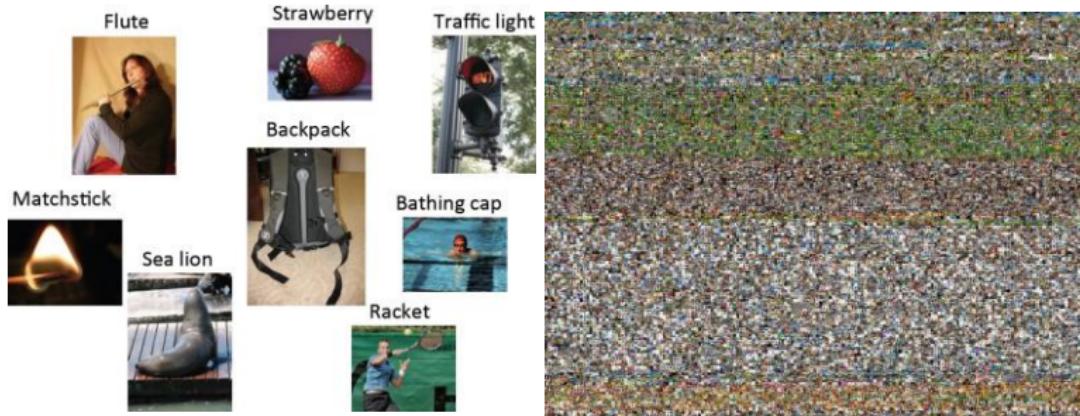
- MNIST - standard dataset for hand-written digit recognition
  - 60000 training, 10000 test images

# LeNet-5, circa 1998



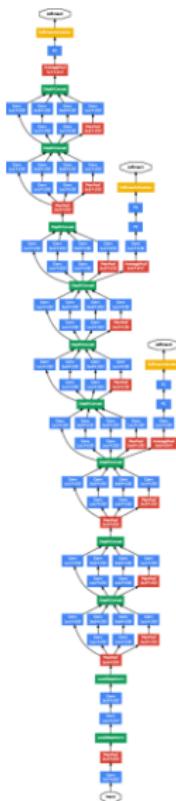
- LeNet developed by Yann LeCun et al.
  - Convolutional neural network
    - Local receptive fields (5x5 connectivity)
    - Subsampling (2x2)
    - Shared weights (reuse same 5x5 “filter”)
    - Breaking symmetry

# ImageNet



- ImageNet - standard dataset for object recognition in images (Russakovsky et al.)
  - 1000 image categories,  $\approx 1.2$  million training images (ILSVRC 2013)

# GoogLeNet, circa 2014

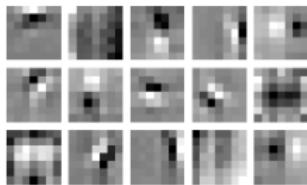


- GoogLeNet developed by Szegedy et al., CVPR 2015
- Modern deep network
- ImageNet top-5 error rate of 6.67% (later versions even better)
- Comparable to human performance (especially for fine-grained categories)

# ResNet, circa 2015

- 
- ResNet developed by He et al., ICCV 2015
  - 152 layers
  - ImageNet top-5 error rate of 3.57%
  - Better than human performance (especially for fine-grained categories)

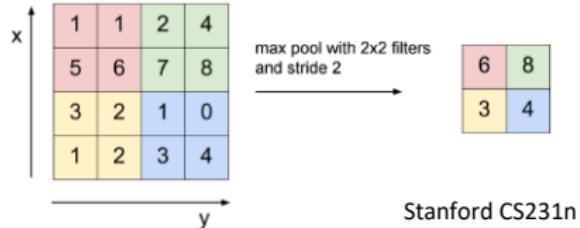
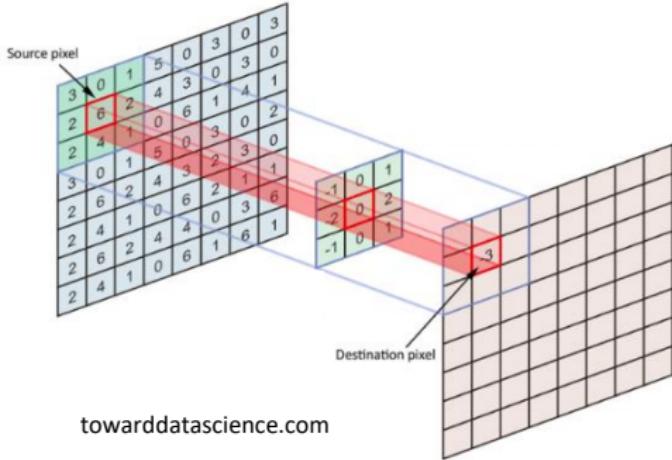
# Key Component 1: Convolutional Filters



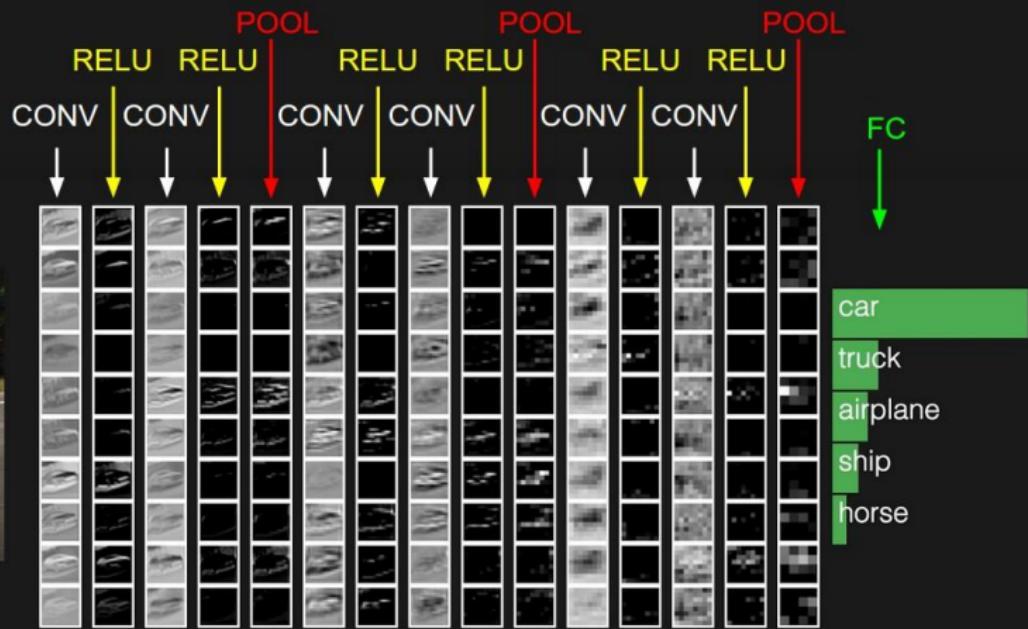
- Share parameters across network
- Reduce total number of parameters
- Provide **translation invariance**, useful for visual recognition

# Common Operations

- Fully connected (dot product)
- Convolution
  - Translationally invariant
  - Controls overfitting
- Pooling (fixed function)
  - Down-sampling
  - Controls overfitting
- Nonlinearity layer (fixed function)
  - Activation functions, e.g. ReLU

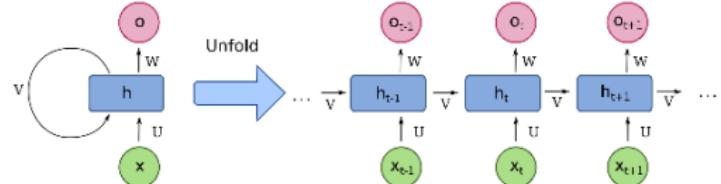
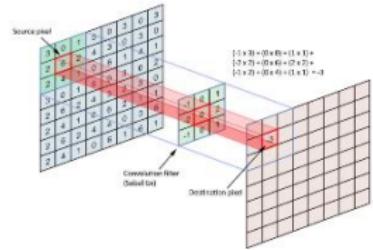


# Example: Small VGG Net From Stanford CS231n



# Neural Network Architectures

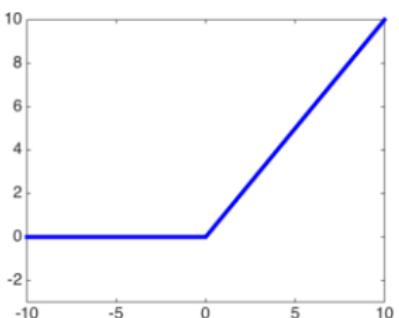
- Convolutional neural network (CNN)
  - Has translational invariance properties from convolution
  - Commonly used for computer vision
- Recurrent neural network RNN
  - Has feedback loops to capture temporal or sequential information
  - Useful for handwriting recognition, speech recognition, reinforcement learning
  - Long short-term memory (LSTM): special type of RNN with advantages in numerical properties
- Others
  - General feedforward networks, variational autoencoders (VAEs), conditional VAEs, generative adversarial networks



# Training Neural Networks

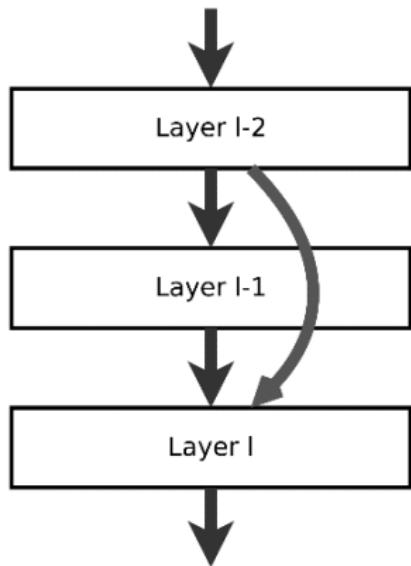
- Data preprocessing
  - Removing bad data
  - Transform input data (e.g. rotating, stretching, adding noise)
- Training process (optimization algorithm)
  - Choice of loss function (eg. L1 and L2 regularization)
  - Dropout: randomly set neurons to zero in each training iteration
  - **Learning rate** (step size) and other hyperparameter tuning
- Software packages: efficient gradient computation
  - Caffe, Torch, Theano, TensorFlow

## Key Component 2: Rectified Linear Units (ReLUs)



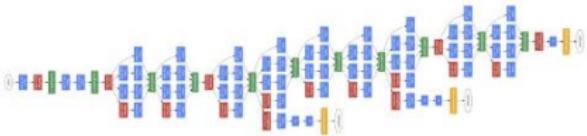
- Vanishing gradient problem
  - If derivatives very small, no/little progress via stochastic gradient descent
  - Occurs with sigmoid function when activation is large in absolute value
- ReLU:  $h(a_j) = \max(0, a_j)$
- Non-saturating, linear gradients (as long as non-negative activation on some training data)
- Sparsity inducing

## Key Component 3: Many, Many Layers



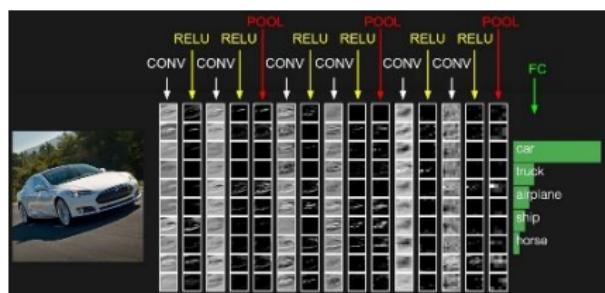
- **ResNet:**  $\approx 152$  layers (“shortcut connections”)
- GoogLeNet:  $\approx 27$  layers (“Inception” modules)
- VGG Net: 16-19 layers (Simonyan and Zisserman, 2014)
- AlexNet: 8 layers (Krizhevsky et al., 2012)

## Key Component 3: Many, Many Layers



- ResNet:  $\approx$ 152 layers (“shortcut connections”)
  - **GoogLeNet:**  $\approx$ 27 layers (“Inception” modules)
  - VGG Net: 16-19 layers (Simonyan and Zisserman, 2014)
  - AlexNet: 8 layers (Krizhevsky et al., 2012)

## Key Component 3: Many, Many Layers



- ResNet:  $\approx 152$  layers (“shortcut connections”)
- GoogLeNet:  $\approx 27$  layers (“Inception” modules)
- **VGG Net:** 16-19 layers  
(Simonyan and Zisserman, 2014)
- AlexNet: 8 layers (Krizhevsky et al., 2012)

## Key Component 4: Momentum

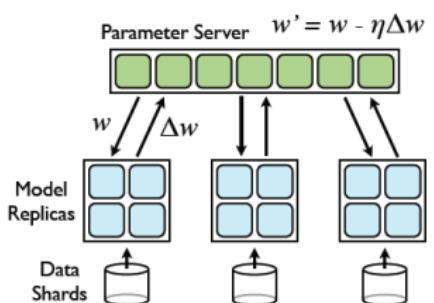
- Trick to escape plateaus / local minima
- Take exponential average of previous gradients

$$\frac{\overline{\partial E_n}^\tau}{\partial w_{ji}} = \frac{\overline{\partial E_n}^\tau}{\partial w_{ji}} + \alpha \frac{\overline{\partial E_n}^{\tau-1}}{\partial w_{ji}}$$

- Maintains progress in previous direction

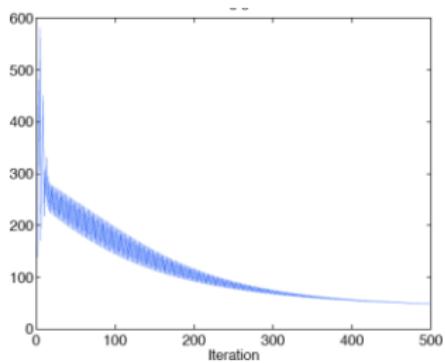
# Key Component 5: Asynchronous Stochastic Gradient Descent

- Big models won't fit in memory
- Want to use compute clusters (e.g. 1000s of machines) to run stochastic gradient descent
- How to parallelize computation?
- **Ignore synchronization across machines**
  - Just let each machine compute its own gradients and pass to a server storing current parameters
  - Ignore the fact that these updates are inconsistent
  - Seems to just work (e.g. Dean et al. NIPS 2012)



## Key Component 6: Learning Rate Schedule

- How to set learning rate  $\eta$ ?:



$$\mathbf{w}^\tau = \mathbf{w}^{\tau-1} + \eta \nabla \mathbf{w}$$

- **Option 1:** Run until validation error plateaus. Drop learning rate by x%
- **Option 2:** Adagrad, adaptive gradient. Per-element learning rate set based on local geometry (Duchi et al. 2010)

## Key Component 7: Data Augmentation



- Augment data with additional synthetic variants (10x amount of data)
- Or just use synthetic data, e.g. Sintel animated movie (Butler et al. 2012)

## Key Component 8: Data and Compute



- Get lots of data (e.g. ImageNet)
- Get lots of compute (e.g. CPU cluster, GPUs)
- Cross-validate like crazy, train models for 2-3 weeks on a GPU
- Researcher gradient descent (RGD) or Graduate student descent (GSD): get 100s of researchers to each do this, trying different network structures

# Challenges

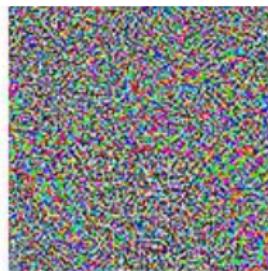
## Interpretability:



"panda"

57.7% confidence

$+$   $\epsilon$



=



"gibbon"

99.3% confidence

# Challenges

Data efficiency:

- ImageNet: 14 million images, 20000 categories
- AlphaStar: 200 years of gameplay



# Challenges

- Problem formulation (what are you trying to predict?)
- Choice of model and optimization algorithm
- Data collection, post-processing
- Feature selection
- ...

## More information

- <https://sites.google.com/site/deeplearningsummerschool>
- <http://tutorial.caffe.berkeleyvision.org/>
- [ufldl.stanford.edu/eccv10-tutorial](http://ufldl.stanford.edu/eccv10-tutorial)
- <http://www.image-net.org/challenges/LSVRC/2012/supervision.pdf>
- Courses: Deep Learning, Natural Language Processing, Computer Vision
- Project ideas
  - Long short-term memory (LSTM) models for temporal data
  - Learning embeddings (word2vec, FaceNet)
  - Structured output (multiple outputs from a network)
  - Zero-shot learning (learning to recognize new concepts without training data)
  - Transfer learning (use data from one domain/task, adapt to another)

# Conclusion

- Readings: Ch. 5.1, 5.2, 5.3
- Feed-forward networks can be used for regression or classification
  - Similar to linear models, except with **adaptive** non-linear basis functions
  - These allow us to do more than e.g. linear decision boundaries
- Different error functions
- Learning is more difficult, error function not convex
  - Use stochastic gradient descent, obtain (good?) local minimum
- Backpropagation for efficient gradient computation