

# Sequential Data

CMPT 726

Mo Chen

SFU Computing Science  
Oct. 28, 2020

Bishop PRML Ch. 13  
Russell and Norvig, AIMA

# Outline

Hidden Markov Models

Inference for HMMs

Learning for HMMs

# Outline

Hidden Markov Models

Inference for HMMs

Learning for HMMs

# Temporal Models

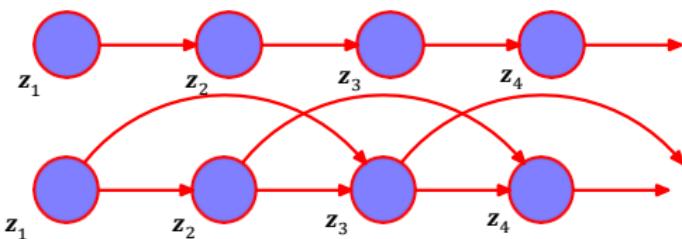
- The world changes over time
  - Explicitly model this change using Bayesian networks
  - Undirected models also exist (will not cover)
- Basic idea: copy state and evidence variables for each time step

e.g. Diabetes management

- $\mathbf{z}_t$  is set of **unobservable state variables** at time  $t$ 
  - $\text{bloodSugar}_t, \text{stomachContents}_t, \dots$
- $\mathbf{x}_t$  is set of **observable evidence variables** at time  $t$ 
  - $\text{measuredBloodSugar}_t, \text{foodEaten}_t, \dots$
- Assume **discrete time step, fixed**
- Notation:  $\mathbf{x}_{a:b} = \mathbf{x}_a, \mathbf{x}_{a+1}, \dots, \mathbf{x}_{b-1}, \mathbf{x}_b$

# Markov Chain

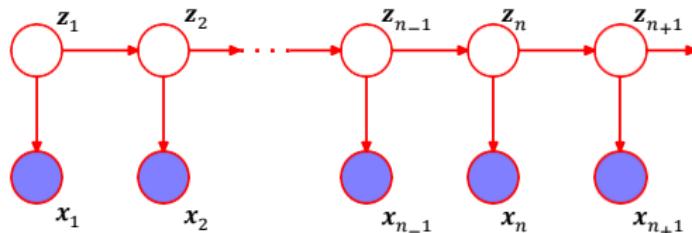
- Construct Bayesian network from these variables
  - parents? distributions? for state variables  $\mathbf{z}_t$ :
- Markov assumption:  $\mathbf{z}_t$  depends on **bounded** subset of  $\mathbf{z}_{1:t}$ 
  - First-order Markov process:  $p(\mathbf{z}_t | \mathbf{z}_{1:t-1}) = p(\mathbf{z}_t | \mathbf{z}_{t-1})$
  - Second-order Markov process:  $p(\mathbf{z}_t | \mathbf{z}_{1:t-1}) = p(\mathbf{z}_t | \mathbf{z}_{t-2}, \mathbf{z}_{t-1})$



- Stationary process:  $p(\mathbf{z}_t | \mathbf{z}_{t-1})$  fixed for all  $t$

# Hidden Markov Model (HMM)

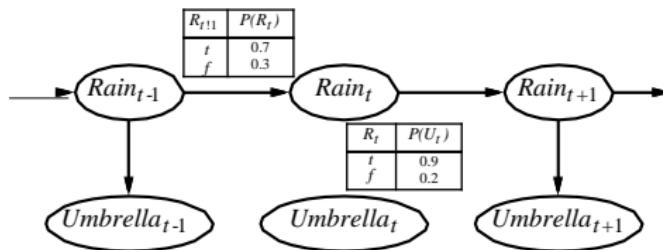
- **Sensor Markov assumption:**  $p(x_t|z_{1:t}, x_{1:t-1}) = p(x_t|z_t)$
- **Stationary process:** transition model  $p(z_t|z_{t-1})$  and sensor model  $p(x_t|z_t)$  fixed for all  $t$  (separate  $p(z_1)$ )
- HMM special type of Bayesian network,  $z_t$  is a **single discrete random variable**:



- **Joint distribution:**

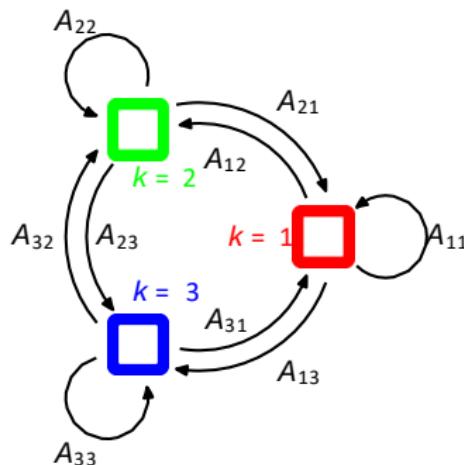
$$p(z_{1:t}, x_{1:t}) = p(z_1) \prod_{i=2:t} p(z_i|z_{i-1}) \prod_{i=1:t} p(x_i|z_i)$$

# HMM Example



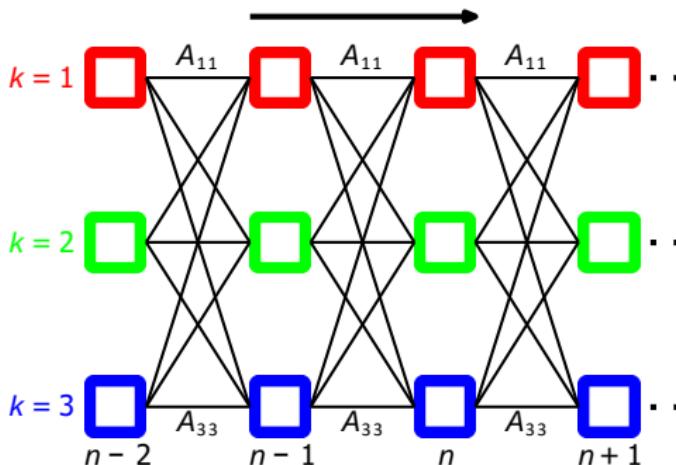
- First-order Markov assumption not true in real world
- Possible fixes:
  - Increase order of Markov process
  - Augment state, add  $temp_t, pressure_t$

# Transition Diagram



- $z_n$  takes one of 3 values
- Using one-of- $K$  coding scheme,  $z_{nk} = 1$  if in state  $k$  at time  $n$
- **Transition matrix  $A$**  where  $p(z_{nk} = 1 | z_{n-1,j} = 1) = A_{jk}$

# Lattice / Trellis Representation



- The lattice or trellis representation shows possible paths through the latent state variables  $z_n$

# Outline

Hidden Markov Models

Inference for HMMs

Learning for HMMs

# Inference Tasks

- **Filtering:**  $p(z_t | x_{1:t})$ 
  - Estimate current unobservable state given all observations to date
- **Prediction:**  $p(z_n | x_{1:t})$  for  $n > t$ 
  - Similar to filtering, without evidence
- **Smoothing:**  $p(z_n | x_{1:t})$  for  $n < t$ 
  - Better estimate of past states
- **Most likely explanation:**  $\arg \max_{z_{1:t}} p(z_{1:t} | x_{1:t})$ 
  - e.g. speech recognition, decoding noisy input sequence

# Filtering

- Aim: devise a **recursive** state estimation algorithm:

$$p(z_{t+1}|x_{1:t+1}) = f(x_{t+1}, p(z_t|x_{1:t}))$$

$$\begin{aligned} p(z_{t+1}|x_{1:t+1}) &= p(z_{t+1}|x_{1:t}, x_{t+1}) \\ &= \alpha p(x_{t+1}|x_{1:t}, z_{t+1}) p(z_{t+1}|x_{1:t}) && \text{(Bayes rule)} \\ &= \alpha p(x_{t+1}|z_{t+1}) p(z_{t+1}|x_{1:t}) && \text{(Markov assumption)} \end{aligned}$$

- i.e. **measurement + prediction**. Prediction by summing out  $z_t$ :

$$\begin{aligned} p(z_{t+1}|x_{1:t+1}) &= \alpha p(x_{t+1}|z_{t+1}) \sum_{z_t} p(z_{t+1}, z_t | x_{1:t}) && \text{(Marginalize)} \\ &= \alpha p(x_{t+1}|z_{t+1}) \sum_{z_t} p(z_{t+1}|z_t, x_{1:t}) p(z_t|x_{1:t}) && \text{(Product rule)} \\ &= \alpha p(x_{t+1}|z_{t+1}) \sum_{z_t} p(z_{t+1}|z_t) p(z_t|x_{1:t}) && \text{(Markov assumption)} \end{aligned}$$

# Filtering Example

Prior:  $p(\text{rain}_1 = \text{true}) = 0.5$

**Prediction:**  $\sum_{R_1} p(R_2|R_1)p(R_1|U_1 = T)$

$$0.7 \times 0.818 + 0.3 \times 0.182 (R_2 = T)$$

$$0.3 \times 0.818 + 0.7 \times 0.182 (R_2 = F)$$

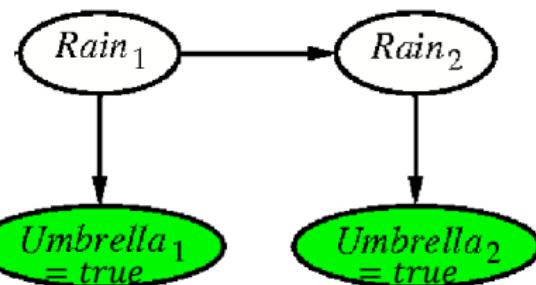
**Measurement:**  $p(R_1|U_1 = T)$

$$\begin{aligned} \text{Normalize } & 0.5 \times 0.9 (R_1 = T) \\ & 0.5 \times 0.2 (R_1 = F) \end{aligned}$$

$$\begin{array}{c} 0.500 \\ 0.500 \\ \downarrow \\ \left[ \begin{array}{c} 0.818 \\ 0.182 \end{array} \right] \end{array} \quad \begin{array}{c} 0.627 \\ 0.373 \\ \downarrow \\ \left[ \begin{array}{c} 0.883 \\ 0.117 \end{array} \right] \end{array}$$

**Measurement:**  $p(R_2|U_2 = T)$

$$\begin{aligned} \text{Normalize } & 0.627 \times 0.9 (R_2 = T) \\ & 0.117 \times 0.2 (R_2 = F) \end{aligned}$$

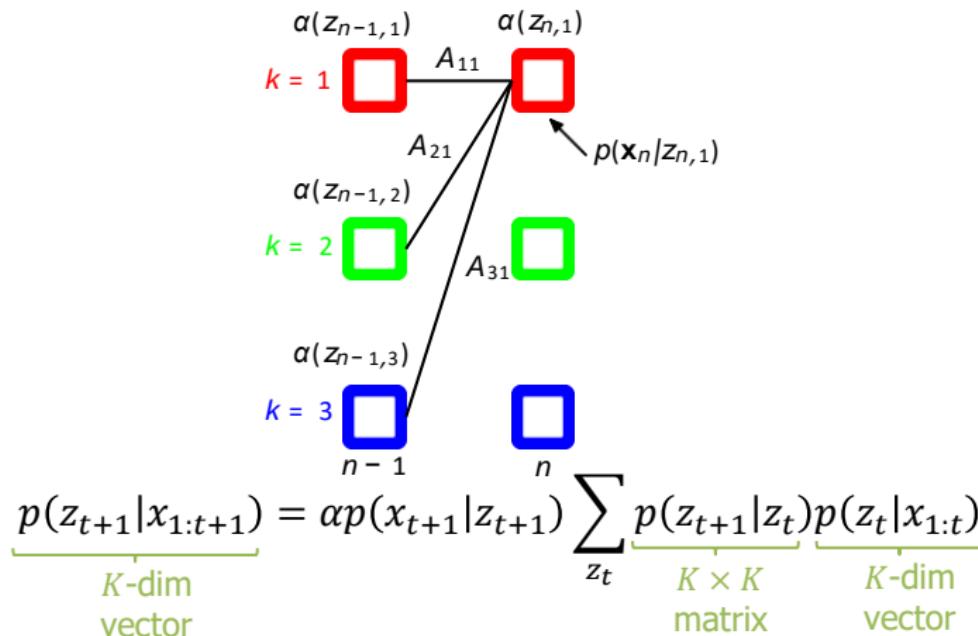


$$p(R_1|U_1) = \frac{p(U_1|R_1)p(R_1)}{P(U_1)}$$

$R_{t-1}$	$P(R_t)$	$R_t$	$P(U_t)$
t	0.7	t	0.9
f	0.3	f	0.2

$$p(z_{t+1}|x_{1:t+1}) = \alpha p(x_{t+1}|z_{t+1}) \sum_{z_t} p(z_{t+1}|z_t)p(z_t|x_{1:t})$$

# Filtering - Lattice

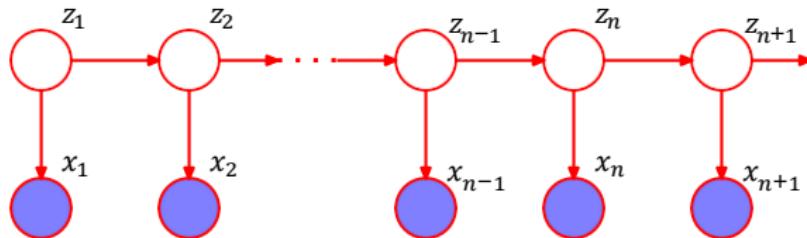


$\Rightarrow O(K^2)$  for each time step,  $O(NK^2)$  for  $N$  time steps

**Forward message passing:**  $\alpha(z_{t+1}) = p(x_{t+1} | z_{t+1}) \sum_{z_t} p(z_{t+1} | z_t) \alpha(z_t)$

- $\alpha(z_t) = p(x_{1:t}, z_t)$ ; previous normalization constant can be dropped
- Initial condition:  $\alpha(z_1) = p(x_1, z_1) = p(x_1 | z_1) p(z_1)$

# Smoothing



- Divide evidence  $x_{1:t}$  into  $x_{1:n}, x_{n+1:t}$

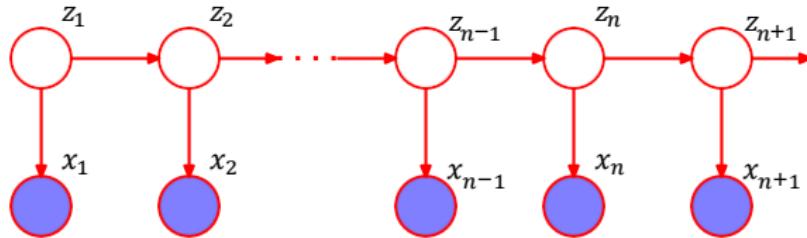
$$p(z_n|x_{1:t}) = \frac{p(x_{1:t}|z_n)p(z_n)}{p(x_{1:t})} \quad (\text{Bayes rule})$$

$$= \frac{p(x_{1:n}|z_n)p(x_{n+1:t}|z_n)p(z_n)}{p(x_{1:t})} \quad (\text{Cond. indep.})$$

$$= \frac{p(x_{1:n}, z_n)p(x_{n+1:t}|z_n)}{p(x_{1:t})} \quad (\text{Product rule})$$

$$= \frac{\alpha(z_n)\beta(z_n)}{p(x_{1:t})}$$

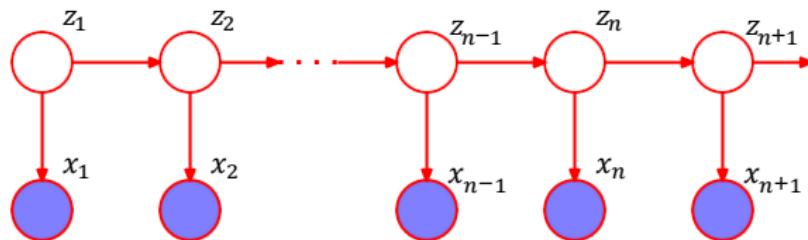
# Smoothing



- Divide evidence  $x_{1:t}$  into  $x_{1:n}, x_{n+1:t}$ ,  $p(z_n|x_{1:t}) = \eta \alpha(z_n) \beta(z_n)$
- Backwards message another recursion:

$$\begin{aligned}
 \frac{p(x_{n+1:t}|z_n)}{\beta(z_n)} &= \sum_{z_{n+1}} p(x_{n+1:t}, z_{n+1}|z_n) && \text{(Marginalize)} \\
 &= \sum_{z_{n+1}} p(x_{n+1:t}|z_{n+1}, z_n) p(z_{n+1}|z_n) && \text{(Product rule)} \\
 &= \sum_{z_{n+1}} p(x_{n+1:t}|z_{n+1}) p(z_{n+1}|z_n) && \text{(Markov assumption)} \\
 &= \sum_{z_{n+1}} p(x_{n+1}|z_{n+1}) \underbrace{p(x_{n+2:t}|z_{n+1})}_{\beta(z_{n+1})} p(z_{n+1}|z_n) && \text{(Cond. indep.)}
 \end{aligned}$$

# Smoothing



- Final condition: go back 2 slides and set  $n = t$

$$p(z_t | x_{1:t}) = \frac{\alpha(z_t)\beta(z_t)}{p(x_{1:t})}$$

$$p(z_t | x_{1:t}) = \frac{p(x_{1:t}, z_t)\beta(z_t)}{p(x_{1:t})}$$

$$\Rightarrow \beta(z_t) = 1$$

$$\alpha(z_1) = p(x_1|z_1)p(z_1)$$

## Smoothing Example

From previous slide

$$p(z_1)$$

$$\frac{\alpha(z_1)}{p(x_{1:1})} = p(z_1|x_{1:1})$$

$$p(z_1|x_{1:2}) = \frac{\alpha(z_1)\beta(z_1)}{p(x_{1:2})}$$

$$\begin{matrix} 0.500 & & 0.627 \\ 0.500 & & 0.373 \\ \downarrow & & \downarrow \\ 0.818 & & 0.883 \\ 0.182 & & 0.117 \end{matrix}$$

forward

$$\frac{\alpha(z_2)}{p(x_{1:2})} = p(z_2|x_{1:2})$$

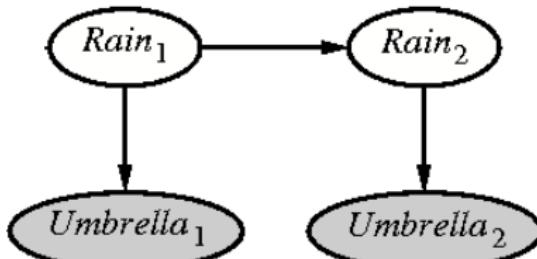
smoothed

$$p(z_2|x_{1:2}) = \frac{\alpha(z_2)\beta(z_2)}{p(x_{1:2})}$$

$$\begin{matrix} 0.883 & & 0.883 \\ 0.117 & & 0.117 \\ \uparrow & & \uparrow \\ 0.690 & & 1.000 \\ 0.410 & & 1.000 \end{matrix}$$

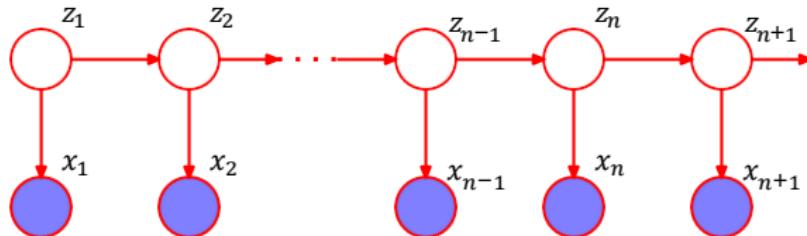
backward

$$\beta(z_2)$$



$$\beta(z_1) = \sum_{z_2} p(x_2|z_2)\beta(z_2)p(z_2|z_1)$$

# Smoothing

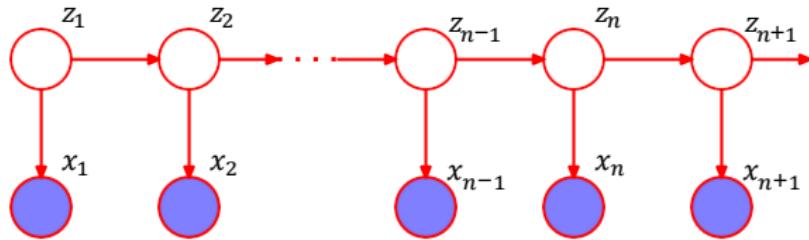


- Backwards message another recursion:

$$\underbrace{\beta(z_n)}_{\substack{K\text{-dim} \\ \text{vector}}} = \sum_{z_{n+1}} p(x_{n+1}|z_{n+1}) \underbrace{\beta(z_{n+1})}_{\substack{K\text{-dim} \\ \text{vector}}} \underbrace{p(z_{n+1}|z_n)}_{K \times K \text{ matrix}}$$

$\Rightarrow O(K^2)$  for each time step,  $O(NK^2)$  for  $N$  time steps

# Forward-Backward Algorithm



- Filter from time 1 to  $N$ , and cache forward messages  $\alpha(z_n)$
- Smooth from time  $N$  to 1, and cache backward messages  $\beta(z_n)$
- Can now compute  $p(z_n | x_1, x_2, \dots, x_t)$  for all  $n$
- Total complexity  $O(NK^2)$
- a.k.a Baum-Welch algorithm

# Outline

Hidden Markov Models

Inference for HMMs

Learning for HMMs

# HMM Parameters

- The **parameters** of an HMM:
  - Transition matrix  $A$  where  $p(z_{nk} = 1 | z_{n-1,j} = 1) = A_{jk}$
  - Sensor model  $\phi_k$  parameters to each  $p(x_n | z_{nk} = 1, \phi_k)$  (e.g.  $\phi_k$  could be mean and variance of Gaussian)
  - Prior for initial state  $z_1$ , model as multinomial  $p(z_{1k} = 1) = \pi_k$ , parameters  $\pi$
- Call these parameters  $\theta = (A, \pi, \phi)$
- **Learning problem:** given one sequence  $x$ , find best  $\theta$ 
  - Extension to multiple sequences straight-forward (assume independent, log of product is sum)

## Maximum Likelihood for HMMs

- We can use maximum likelihood to choose the best parameters:

$$\boldsymbol{\theta}_{ML} = \arg \max p(\mathbf{x}|\boldsymbol{\theta})$$

- Unfortunately this is hard to do: we can get  $p(\mathbf{x}|\boldsymbol{\theta})$  by summing out from the joint distribution:

$$\begin{aligned} p(\mathbf{x}|\boldsymbol{\theta}) &= \sum_{z_1} \sum_{z_2} \cdots \sum_{z_N} p(\mathbf{x}, z_1, z_2, \dots, z_N | \boldsymbol{\theta}) \\ &\equiv \sum_{\mathbf{z}} p(\mathbf{x}, \mathbf{z} | \boldsymbol{\theta}) \end{aligned}$$

- But this sum has  $K^N$  terms in it
- No simple closed-form solution
- Instead, use expectation-maximization (EM)

# EM for HMMs

- Start with initial guess for parameters  $\theta^{old} = (\mathbf{A}, \boldsymbol{\pi}, \boldsymbol{\phi})$
- E-step:** Calculate posterior on latent variables  $p(z|x, \theta^{old})$

Forward-backward algorithm

$$\mathbb{E}_{z \sim p(z|x, \theta^{old})} [\ln p(x, z|\theta)]$$

- M-step:** Maximize  $Q(\theta, \theta^{old}) = \sum_z p(z|x, \theta^{old}) \ln p(x, z|\theta)$  wrt  $\theta$
- Let's look at the M-step, and see how the HMM structure helps us

## HMM M-step

- **M-step:** Maximize  $Q(\boldsymbol{\theta}, \boldsymbol{\theta}^{old}) = \sum_z p(\mathbf{z}|x, \boldsymbol{\theta}^{old}) \ln p(x, \mathbf{z}|\boldsymbol{\theta})$  wrt  $\boldsymbol{\theta}$ :
- The complete data log-likelihood factors nicely:

$$\begin{aligned}\ln p(x, z|\boldsymbol{\theta}) &= \ln \left\{ p(z_1|\boldsymbol{\pi}) \prod_{i=2}^N p(z_i|z_{i-1}, \mathbf{A}) \prod_{i=1}^N p(x_i|z_i, \boldsymbol{\phi}) \right\} \\ &= \ln p(z_1|\boldsymbol{\pi}) + \sum_{i=2}^N \ln p(z_i|z_{i-1}, \mathbf{A}) + \sum_{i=1}^N \ln p(x_i|z_i, \boldsymbol{\phi})\end{aligned}$$

- To maximize  $Q$  we now have 3 separate problems, one for each parameter
  - Let's consider each in turn

# Prior $\pi$

- Maximize  $Q$  wrt prior on initial state  $\pi$ :

$$\begin{aligned} Q(\pi, \theta^{old}) &= \sum_z p(z|x, \theta^{old}) \ln p(z_1|\pi) \\ &= \sum_z p(z|x, \theta^{old}) \ln \prod_{k=1}^K \pi_k^{z_{1k}} \\ &= \sum_z p(z|x, \theta^{old}) \sum_{k=1}^K z_{1k} \ln \pi_k \\ &= \sum_{k=1}^K \ln \pi_k \sum_z p(z|x, \theta^{old}) z_{1k} \\ &= \sum_{k=1}^K p(z_{1k} = 1|x, \theta^{old}) \ln \pi_k \end{aligned}$$

- i.e. smoothed value for  $z_1$  being in state  $k$

$$Q(\pi, \theta^{old}) = \sum_{k=1}^K p(z_{1k} = 1 | \mathbf{x}, \boldsymbol{\theta}^{old}) \ln \pi_k$$

- Can solve for best  $\pi$
- Use Lagrange multiplier to enforce constraint  $\sum_k \pi_k = 1$

$$\pi_k = \frac{p(z_{1k} = 1 | \mathbf{x}, \boldsymbol{\theta}^{old})}{\sum_{j=1}^K p(z_{1j} = 1 | \mathbf{x}, \boldsymbol{\theta}^{old})}$$

- Intuitively sensible result: new  $\pi_k$  is smoothed probability of being in state  $k$  at time 1 using old parameters
- E-step needs to calculate smoothed  $p(z_{1k} = 1 | \mathbf{x}, \boldsymbol{\theta}^{old})$ ; this is fast  $O(NK^2)$

# Transition Matrix $A$

- Maximize  $Q$  wrt transition matrix  $A$ :

$$\begin{aligned}
 Q(A, \theta^{old}) &= \sum_z p(z|x, \theta^{old}) \sum_{i=2}^N \ln p(z_i|z_{i-1}, A) \\
 &= \sum_z p(z|x, \theta^{old}) \sum_{i=2}^N \ln \prod_{k=1}^N \prod_{j=1}^K A_{jk}^{z_{i-1,j} z_{i,k}} \\
 &= \sum_z p(z|x, \theta^{old}) \sum_{i=2}^N \sum_{k=1}^K \sum_{j=1}^K z_{i-1,j} z_{i,k} \ln A_{jk} \\
 &= \sum_{k=1}^K \sum_{j=1}^K \ln A_{jk} \sum_{i=2}^N \sum_z p(z|x, \theta^{old}) z_{i-1,j} z_{i,k} \\
 &= \sum_{k=1}^K \sum_{j=1}^K \ln A_{jk} \sum_{i=2}^N p(z_{i-1} = j, z_i = k | x, \theta^{old})
 \end{aligned}$$

- E-step needs to calculate  $p(z_{i-1} = j, z_i = k | x, \theta^{old})$ ; can be done quickly using forward and backward messages

$$Q(\mathbf{A}, \boldsymbol{\theta}^{old}) = \sum_{k=1}^K \sum_{j=1}^K \ln A_{jk} \sum_{i=2}^N p(z_{i-1} = j, z_i = k | \mathbf{x}, \boldsymbol{\theta}^{old})$$

- Can solve for best  $\mathbf{A}$
- Again use Lagrange multipliers to enforce constraint  $\sum_k A_{jk} = 1$

$$A_{jk} = \frac{\sum_{n=2}^N p(z_{n-1} = j, z_n = k | \mathbf{x}, \boldsymbol{\theta}^{old})}{\sum_{l=1}^K \sum_{n=2}^N p(z_{n-1} = j, z_n = l | \mathbf{x}, \boldsymbol{\theta}^{old})}$$

- Again sensible result:  $A_{jk}$  set to expected number of times we transition from state  $j$  to  $k$  using the smoothed results from old parameters

# Sensor Model

- Similar derivation for sensor model parameters  $\phi$
- Again end up with weighted parameter estimated based on expected values of states given smoothed estimates

# HMM EM Summary

- Start with initial guess for parameters  $\theta^{old} = (A, \pi, \phi)$
- Run forward-backward algorithm to get all messages  $\alpha(z_n), \beta(z_n)$  (E-step)
  - $O(NK^2)$  time complexity
  - Can use these to compute any smoothed posterior  $p(z_{nk} = 1 | x, \theta^{old})$
  - Also can compute any  $p(z_{nk} = 1, z_{n,k} = 1 | x, \theta^{old})$
  - Using these, update values for parameters (M-step)
  - $\pi_k$  is smoothed probability of being in state  $k$  at time 1
  - $A_{jk}$  is smoothed probability of transitioning from state  $j$  to  $k$  averaged over all time steps
  - $\phi$  is weighted sensor parameters using smoothed probabilities (e.g. similar to mixture of Gaussians)
- Repeat until convergence

# Inference Tasks

- **Filtering:**  $p(z_t | x_{1:t})$ 
  - Estimate current unobservable state given all observations to date
- **Prediction:**  $p(z_n | x_{1:t})$  for  $n > t$ 
  - Similar to filtering, without evidence
- **Smoothing:**  $p(z_n | x_{1:t})$  for  $n < t$ 
  - Better estimate of past states
- **Most likely explanation:**  $\arg \max_{z_{1:t}} p(z_{1:t} | x_{1:t})$ 
  - e.g. speech recognition, decoding noisy input sequence

# Sequence of Most Likely States

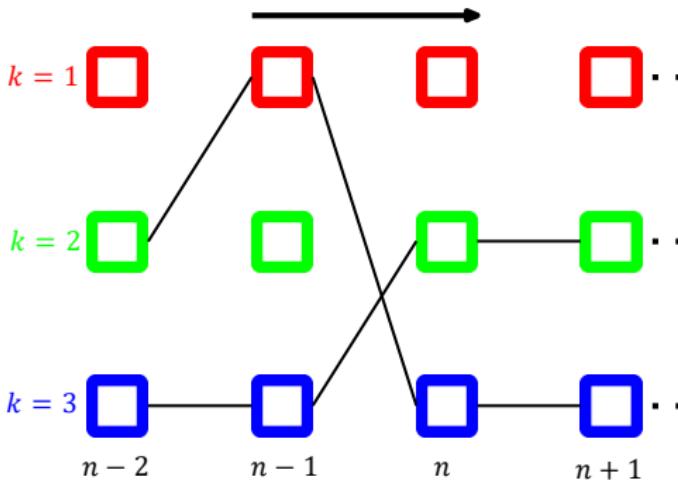
- Most likely sequence is not same as sequence of most likely states:

$$\arg \max_{z_{1:N}} p(z_{1:N} | x_{1:N})$$

versus

$$\left( \arg \max_{z_1} p(z_1 | x_{1:N}), \dots, \arg \max_{z_N} p(z_N | x_{1:N}) \right)$$

## Paths Through HMM

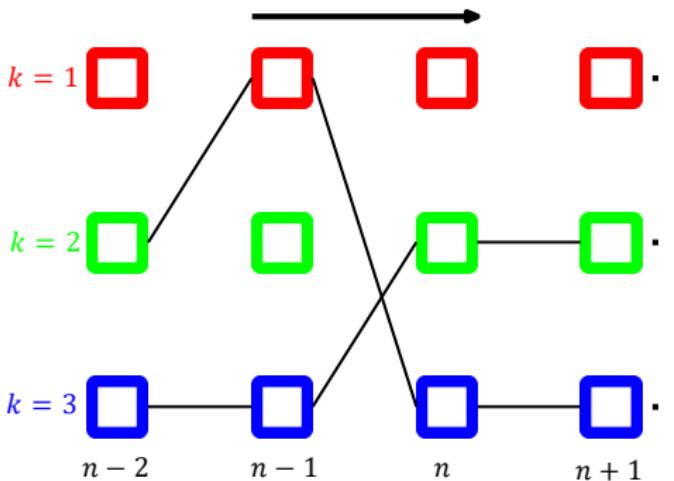


- There are  $K^N$  paths to consider through the HMM for computing

$$\arg \max_{z_{1:N}} p(z_{1:N} | x_{1:N})$$

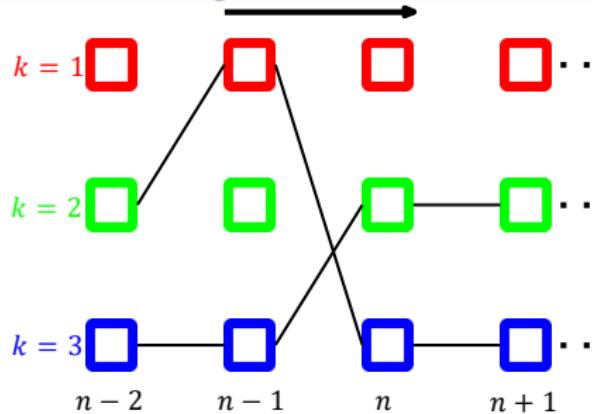
- Need a faster method

# Viterbi Algorithm



- Insight: for any value  $k$  for  $z_n$ , the best path  $(z_1, z_2, \dots, z_n = k)$  ending in  $z_n = k$  consists of the best path  $(z_1, z_2, \dots, z_{n-1} = j)$  **for some  $j$** , plus one more step
  - Don't need to consider exponentially many paths, just  $K$  at each time step
  - Dynamic programming algorithm – **Viterbi algorithm**

## Viterbi Algorithm - Math



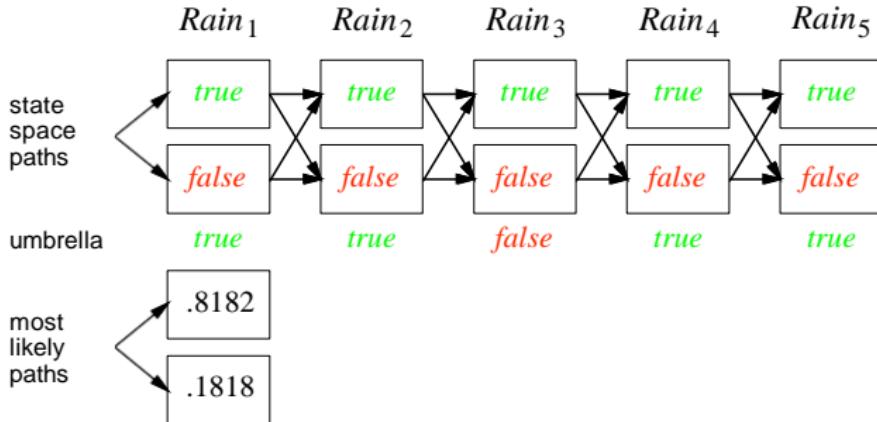
- Define message

$$w(n, k) = \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k)$$

- From factorization of joint distribution:

$$\begin{aligned} w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_{n-1}, z_1, \dots, z_{n-1}) p(x_n | z_n = k) p(z_n = k | z_{n-1}) \\ &= \max_{z_{n-1}} \max_{z_1, \dots, z_{n-2}} p(x_{1:n-1}, z_{1:n-1}) p(x_n | z_n = k) p(z_n = k | z_{n-1}) \\ &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j) \end{aligned}$$

# Viterbi Algorithm - Example



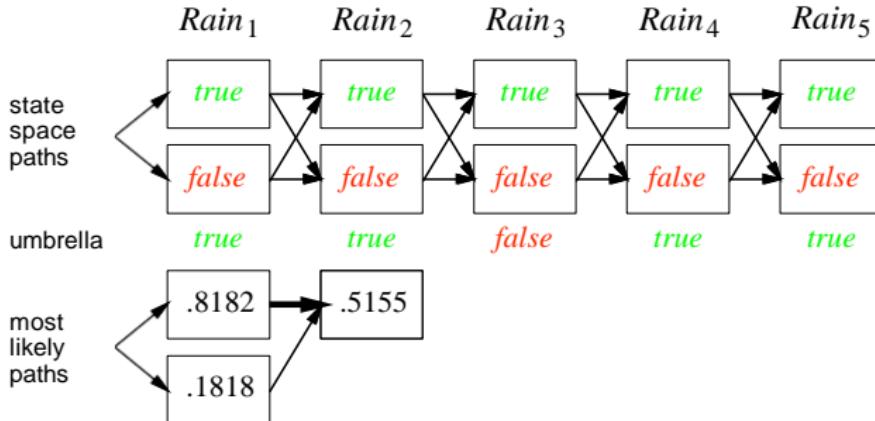
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(rain_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example

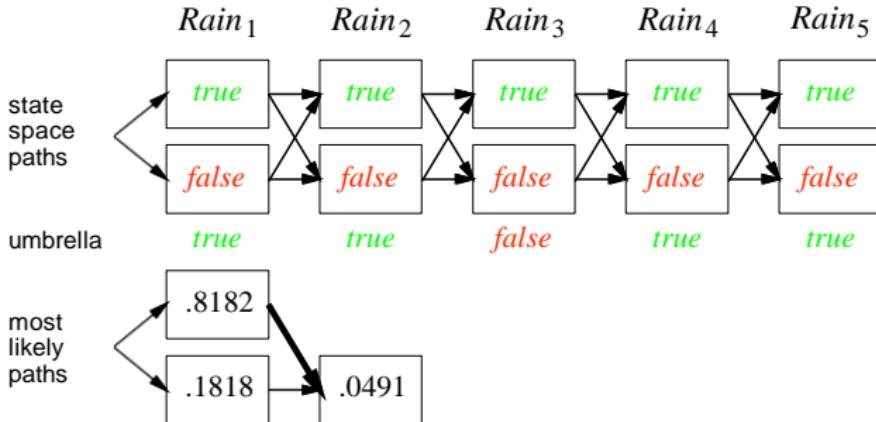


R <sub>t-1</sub>	P(R <sub>t</sub> )	R <sub>t</sub>	P(U <sub>t</sub> )
t	0.7	t	0.9
f	0.3	f	0.2

$$p(\text{rain}_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



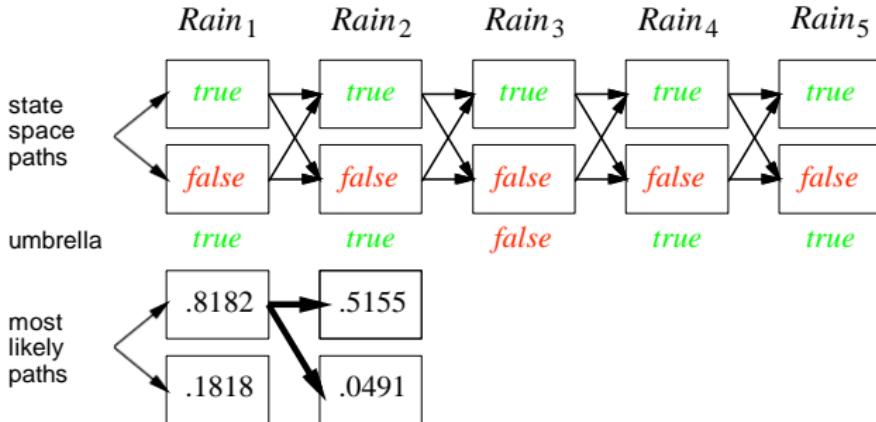
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(\text{rain}_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



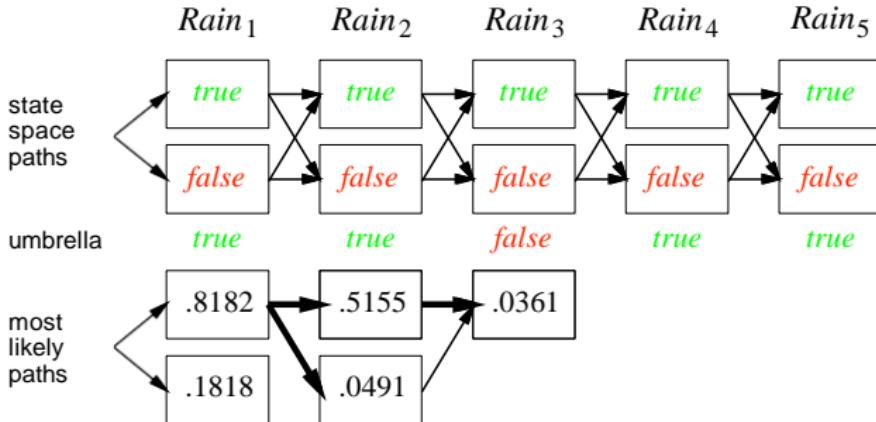
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(\text{rain}_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



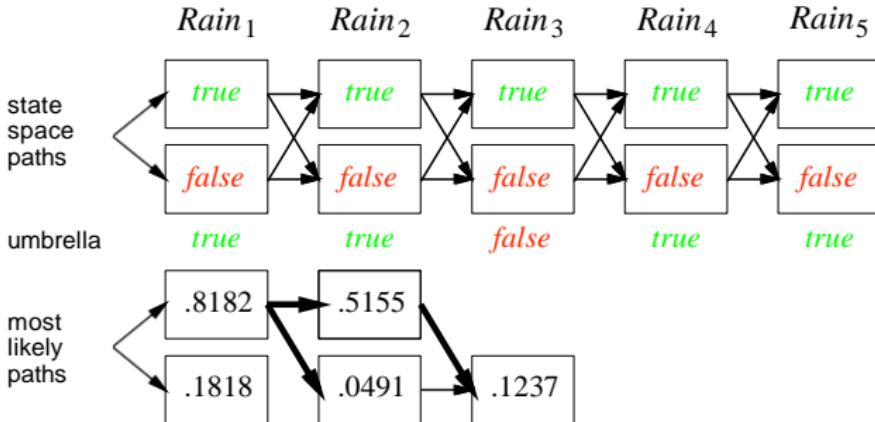
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(\text{rain}_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



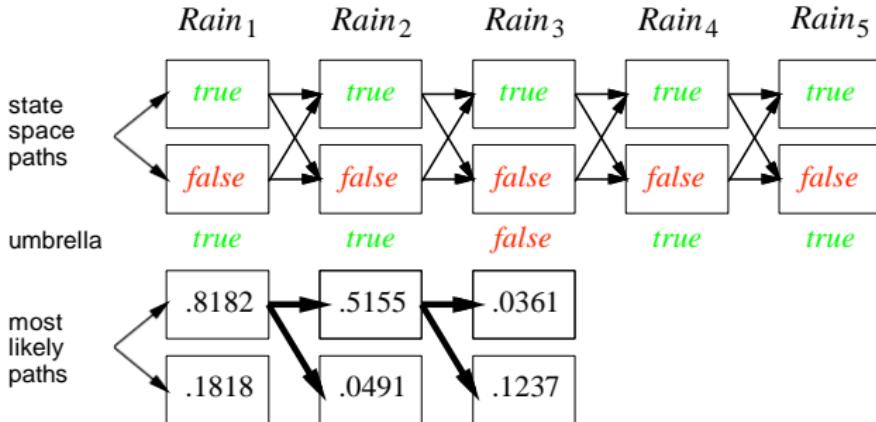
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(rain_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



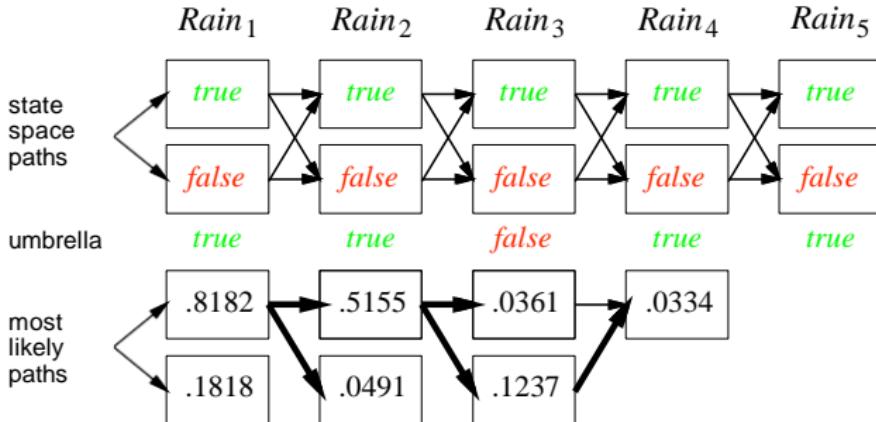
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(rain_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



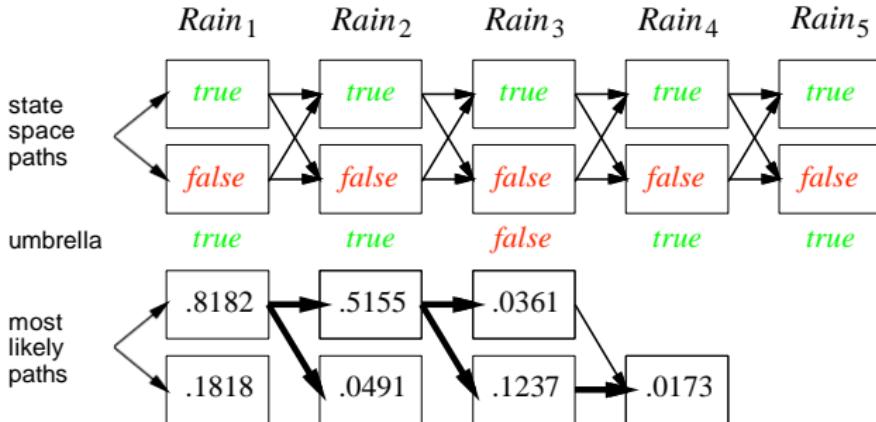
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(\text{rain}_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



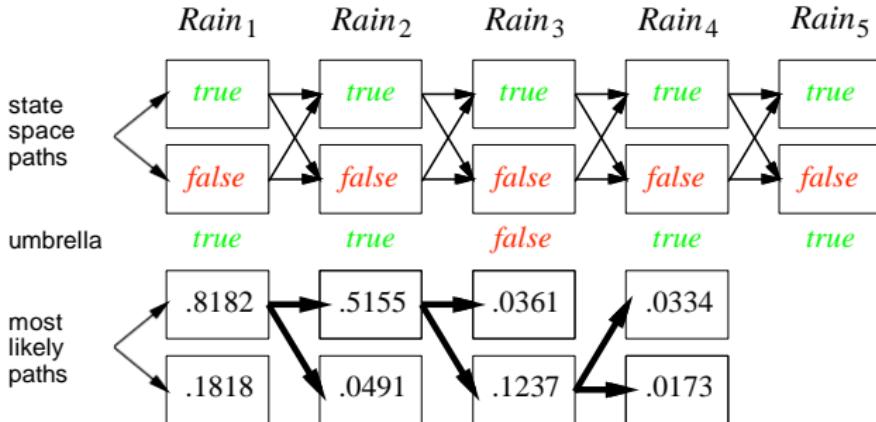
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(rain_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



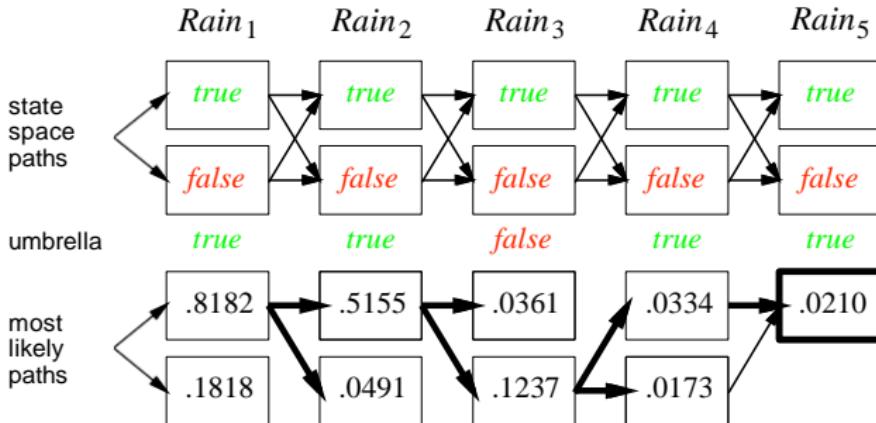
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(rain_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example

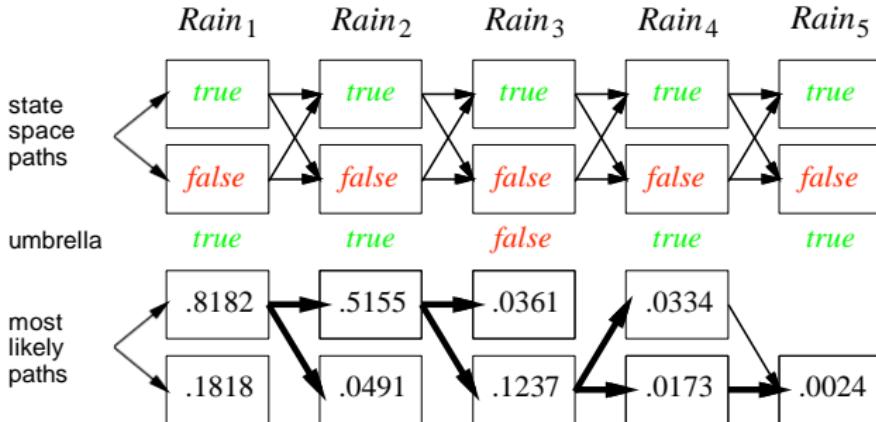


R <sub>t-1</sub>	P(R <sub>t</sub> )	R <sub>t</sub>	P(U <sub>t</sub> )
t	0.7	t	0.9
f	0.3	f	0.2

$$p(rain_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



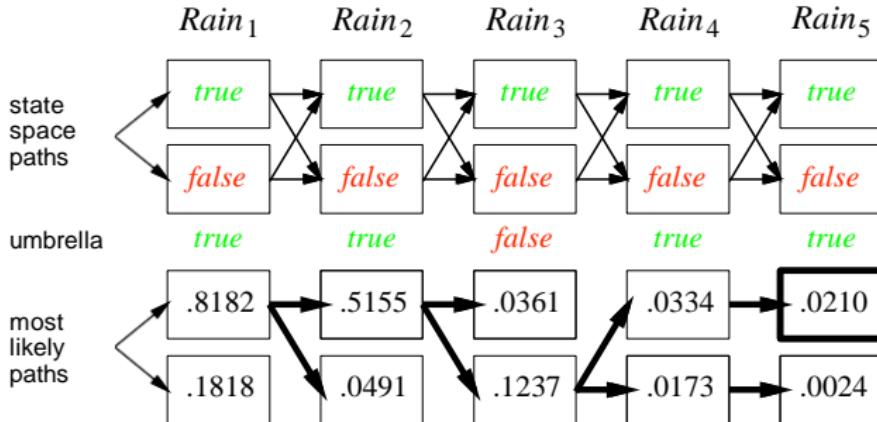
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(rain_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



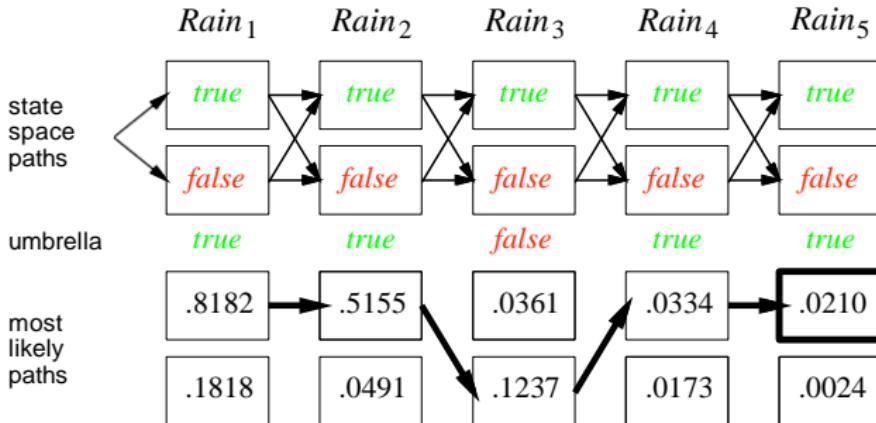
R <sub>t-1</sub>	P(R <sub>t</sub> )
t	0.7
f	0.3

R <sub>t</sub>	P(U <sub>t</sub> )
t	0.9
f	0.2

$$p(rain_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Example



R <sub>t-1</sub>	P(R <sub>t</sub> )	R <sub>t</sub>	P(U <sub>t</sub> )
t	0.7	t	0.9
f	0.3	f	0.2

$$p(\text{rain}_1 = \text{true}) = 0.5$$

$$\begin{aligned}
 w(n, k) &= \max_{z_1, \dots, z_{n-1}} p(x_1, \dots, x_n, z_1, \dots, z_n = k) \\
 &= \max_j w(n-1, j) p(x_n | z_n = k) p(z_n = k | z_{n-1} = j)
 \end{aligned}$$

# Viterbi Algorithm - Complexity

- Each step of the algorithm takes  $O(K^2)$  work
- With  $N$  time steps,  $O(NK^2)$  complexity to find most likely sequence
- Much better than naive algorithm evaluating all  $K^N$  possible paths

# Conclusion

- Readings: Ch. 13.2, 13.2.1, 13.2.2, 13.2.5
- HMM - Probabilistic model of temporal data
  - Discrete hidden (unobserved, latent) state variable at each time
  - Observation (can be discrete / continuous) at each time
  - Conditional independence assumptions (Markov)
  - Assumptions on distributions (stationary)
- Inference
  - Filtering
  - Smoothing
  - Most likely sequence
- Maximum likelihood learning
  - EM – efficient computation  $O(NK^2)$  time using forward-backward smoothing
- Most likely sequence in HMM
  - Viterbi algorithm –  $O(NK^2)$  time, dynamic programming algorithm