

Reinforcement Learning

CMPT 726

Mo Chen

SFU Computing Science

4/11/2020

Outline


- Reinforcement learning problem setup
- Imitation learning
- Basic ideas in RL
- Model-free value-based RL
- Policy-based and actor-critic RL

Outline

- **Reinforcement learning problem setup**
- Imitation learning
- Basic ideas in RL
- Model-free value-based RL
- Policy-based and actor-critic RL

Markov Decision Process

- Probabilistic model of robots and other systems
- State: $s \in \mathcal{S}$, discrete or continuous
- Action (control): $a \in \mathcal{A}$, discrete or continuous
- Transition operator (dynamics): \mathcal{T}
 - $\mathcal{T}_{ijk} = p(s_{t+1} = i | s_t = j, a_t = k) \leftarrow$ a tensor (multidimensional array)



State in MDPs and Reinforcement Learning

- State includes the internal states of an agent, but often also include
 - State of other agents
 - State of the environment
 - Sensor measurements
- Distinction between state and observation can be blurred
- In general, the state contains all variables other than actions that determine the next state through the transition probability $p(s_{t+1}|s_t, a_t)$



Policy and Reward

- Control policy (feedback control): $\pi(a|s)$

- Parametrized by θ

$$\theta: \pi_\theta(a|s) := p(a|s)$$

- Can be stochastic: probability of applying action a at state s



Policy and Reward

- Control policy (feedback control): $\pi(a|s)$


- Parametrized by θ

$$\theta: \pi_\theta(a|s) := p(a|s)$$

- Can be stochastic: probability of applying action a at state s


- Reward function: $r(s_t, a_t)$

- Reward received for being at state s_t and applying action a_t



Extensions of Problem Setup

- Partially observability
 - Partially Observable Markov Decision Process (POMDP)
 - State not fully known; instead, act based on observations



- Policy: $\pi_\theta(a|o)$
- In this class, state s will be synonymous with observation o .

Reinforcement Learning Objective

- Given: an MDP with state space \mathcal{S} , action space \mathcal{A} , transition probabilities \mathcal{T} , and reward function $r(s, a)$

- Objective: Maximize discounted sum of rewards (“return”)

$$\underset{\pi_\theta}{\text{maximize}} \mathbb{E} \sum_t \gamma^t r(s_t, a_t)$$

- $\gamma \in (0, 1]$: discount factor – larger roughly means “far-sighted”
 - Prioritizes immediate rewards
 - $\gamma < 1$ avoids infinite rewards; $\gamma = 1$ is possible if all sequences are finite

- Constraints: often implicit, and part of the objective
 - Subject to transition matrix \mathcal{T} (system dynamics)

Outline

- Reinforcement learning problem setup
- **Imitation learning**
- Basic ideas in RL
- Model-free value-based RL
- Policy-based and actor-critic RL

Imitation Learning

- Collect data through expert demonstration – sequence of states and actions, $\{s_0, a_0, s_1, a_1, \dots, s_{N-1}, a_{N-1}, s_N\}$
 - Note: Expert may not be solving $\underset{\pi}{\text{maximize}} \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)]$



- Learn $\pi_{\theta}(a_t | s_t)$ from data via regression

Imitation Learning

- Collect data through expert demonstration – sequence of states and actions, $\{s_0, a_0, s_1, a_1, \dots, s_{N-1}, a_{N-1}, s_N\}$
 - Note: Expert may not be solving $\underset{\pi}{\text{maximize}} \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)]$



- Learn $\pi_{\theta}(a_t | s_t)$ from data via regression


Imitation Learning

- Collect data through expert demonstration – sequence of states and actions, $\{s_0, a_0, s_1, a_1, \dots, s_{N-1}, a_{N-1}, s_N\}$
 - Note: Expert may not be solving $\underset{\pi}{\text{maximize}} \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)]$
- Learn $\pi_{\theta}(a_t | s_t)$ from data via regression
- Usually doesn't work due to "drift": small mistakes add up, and takes the system far from trained states
 - Sometimes, there can be "tricks" to make imitation learning work!

Autonomous Driving Through Imitation



Training:




Testing:



Dataset Aggregation

- Imitation learning drawback:
 - Distribution of observations in training is different from distribution of observations during test
 - Some states have never been seen during demonstration



- How to make the distributions equal?
 - Train perfect policy
 - Change data set → DAgger (Dataset Aggregation)

Dataset Aggregation (DAgger) Algorithm

1. Train policy from some initial data, $\mathcal{D}_i = \{s_0, a_0, s_1, a_1, \dots, s_{N-1}, a_{N-1}, s_N\}$
2. Run policy to obtain new observations $\{s_{N+1}, s_{N+2}, \dots, s_{N+M}\}$
 - Note: time indices and states here may not continue from initial data
3. Use humans to label data by providing actions for new observations, $\{a_{N+1}, \dots, a_{N+M-1}\}$
 - This creates another data set, $\bar{\mathcal{D}}_i = \{s_{N+1}, a_{N+1}, s_{N+2}, a_{N+2} \dots, a_{N+M-1}, s_{N+M}\}$
4. Combine two datasets, $\mathcal{D}_i \leftarrow \mathcal{D}_i \cup \bar{\mathcal{D}}_i$
 - Go back to first step

Challenges

- Non-Markovian behaviour
 - Perhaps augment state/observation space to include some history
 - Use neural networks that implicitly capture time series data: RNNs/LSTMs
- Unnatural data collection
 - Humans are probably not very good at collecting correction data in this manner
- Inconsistencies in human action

Addressing Drift

- Main goal: Teach system to correct errors
- Explicitly demonstrate corrections (DAgger)
- During demonstration, add noise to “force” mistakes, and see how humans correct them
- Ask humans to intentionally make mistakes
- Prior knowledge and heuristics
 - Example: Learn from stabilizing controller




Imitation Learning Tricks

- Common neural network architectures
 - LSTM – since we have time-series data
 - CNN – usually in combination with LSTM, if the observations are images

- Simplify action space:
 - Driving example: action space simplified to {left, centre, right}

- Clever data collection
 - Driving example: side cameras

- Inverse reinforcement learning
 - Learn goal, instead of policy, from data
 - Use reinforcement learning to learn to achieve the same goal



Imitation Learning Drawbacks

- Very small amount of data – challenging for training deep neural networks
- Humans are not very good at providing some kinds of actions
 - Quadrotor motor speed
 - Non-humanoid machines
- Hard to perform better at tasks humans are not very good at

Outline

- Reinforcement learning problem setup
- Imitation learning
- **Basic ideas in RL**
- Model-free value-based RL
- Policy-based and actor-critic RL

Reinforcement Learning

- Humans can learn without imitation
 - Given goal/task
 - Try an initial strategy
 - See how well the task is performed
 - Adjust strategy next time
- Reinforcement learning agent
 - Given goal/task in the form of reward function $r(s, a)$
 - Start with initial policy $\pi_\theta(a|s)$; execute policy
 - Obtain sum of rewards, $\sum_t r(s_t, a_t)$
 - Improve policy by updating θ , based on rewards



Reinforcement Learning Objective

- Given: an MDP with state space \mathcal{S} , action space \mathcal{A} , transition probabilities \mathcal{T} , and reward function $r(s, a)$

- Objective: Maximize expected discounted sum of rewards (“return”)

$$\underset{\pi_\theta}{\text{maximize}} \mathbb{E} \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)$$

- $\gamma \in (0, 1]$: discount factor – larger roughly means “far-sighted”
 - Prioritizes immediate rewards
 - $\gamma < 1$ avoids infinite rewards; $\gamma = 1$ is possible if all sequences are finite

- Constraints: now incorporated into the reward function

- Only constraint (usually implicit): subject to transition matrix \mathcal{T} (system dynamics)

RL vs. Other ML Paradigms

- No supervisor
 - But we will often draw inspiration from supervised learning
- Sequential data in time
- Reward feedback is obtained after a long time
 - Many actions combined together will receive reward
 - Actions are dependent on each other
- In robotics/health sciences/etc.: lack of data

Reinforcement Learning Categories

- Model-based
 - Explicitly involves an MDP model
- Model-free
 - Does not explicitly involve an MDP model
- Value based
 - Learns value function, and derives policy from value function
- Policy based
 - Learns policy without value function
- Actor critic
 - Incorporates both value function and policy

Value Functions

- “**State-value function**”: $V_\pi(s)$ -- expected return starting from state s and following policy π
 - $V_\pi(s) = \mathbb{E}_{a_t \sim \pi} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$
 - Expectation is on the random sequence $\{s_0, a_0, s_1, a_1, \dots\}$
- “**Action-value function**”, or “ **Q function**”: $Q_\pi(s, a)$ -- expected return starting from state s , taking action a , and then following policy π
 - $Q_\pi(s, a) = \mathbb{E}_{a_t \sim \pi} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s, a_0 = a]$


Principal of Optimality

- Optimal discounted sum of rewards:

- $$V_{\pi^*}(s) = \max_{\pi} \mathbb{E} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \mid s_0 = s \right]$$

- Dynamic programming:

- $$V_{\pi^*}(s) = \max_{a_t} \mathbb{E}[r(s_t, a_t) + \gamma V_{\pi^*}(s_{t+1}) \mid s_t = s]$$



Principal of Optimality

- Optimal discounted sum of rewards:

- $V_{\pi^*}(s) = \max_{\pi} \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$

- Dynamic programming:


- $V_{\pi^*}(s) = \max_{a_t} \mathbb{E}[r(s_t, a_t) + \gamma V_{\pi^*}(s_{t+1}) | s_t = s]$

- $Q_{\pi^*}(s, a) = \mathbb{E}[r(s_t, a_t) + \gamma V_{\pi^*}(s_{t+1}) | s_t = s, a_t = a]$

- Actually, recurrence is true even without maximization

- $V_{\pi}(s) = \mathbb{E}[r(s_t, a_t) + \gamma V_{\pi}(s_{t+1}) | s_t = s]$

- $Q_{\pi}(s, a) = \mathbb{E}[r(s_t, a_t) + \gamma V_{\pi}(s_{t+1}) | s_t = s, a_t = a]$



Basic Properties of Value Functions

- $V_{\pi^*}(s) = \max_{\pi} V_{\pi}(s)$
- $Q_{\pi^*}(s, a) = \max_{\pi} Q_{\pi}(s, a)$
- $V_{\pi^*}(s) = \max_a Q_{\pi^*}(s, a)$
- For now, value functions are stored in multi-dimensional arrays
- DP leads to deterministic policies – we will come back to stochastic policies

Optimizing the RL Objective via DP

- State-value function

- $V_{\pi^*}(s) = \max_{a_t} \mathbb{E}[r(s_t, a_t) + \gamma V(s_{t+1}) | s_t = s]$

- $V_{\pi^*}(s) = \max_a \{r(s, a) + \gamma \mathbb{E}[V(s_{t+1}) | s_t = s]\}$


- $V_{\pi^*}(s) = \max_a \{r(s, a) + \gamma \sum_{s'} [p(s' | s, a) V_{\pi^*}(s')]\}$

- **"Bellman backup"**: $V(s) \leftarrow \max_a \{r(s, a) + \gamma \sum_{s'} [p(s' | s, a) V(s')]\}$

- This is done for all s
- Iterate until convergence

- Optimal policy: $a = \arg \max_{a'} \{r(s, a') + \gamma \sum_{s'} [p(s' | s, a') V(s')]\}$

- Deterministic



Optimizing the RL Objective via DP

- Action-value function

- $Q_{\pi^*}(s, a) = \mathbb{E} \left[r(s_t, a_t) + \gamma \max_{a_{t+1}} Q_{\pi^*}(s_{t+1}, a_{t+1}) \mid s_t = s, a_t = a \right]$

- $Q_{\pi^*}(s, a) = r(s, a) + \gamma \mathbb{E} \left[\max_{a_{t+1}} Q_{\pi^*}(s_{t+1}, a_{t+1}) \mid s_t = s, a_t = a \right]$

- $Q_{\pi^*}(s, a) = r(s, a) + \gamma \sum_{s'} [p(s' | s, a) V_{\pi^*}(s')]$

- “Bellman backup”:

- $V(s) \leftarrow \max_a Q(s, a)$


- $Q(s, a) \leftarrow r(s, a) + \gamma \sum_{s'} [p(s' | s, a) V(s')]$

- This is done for all s and all a

- Iterate until convergence

- Optimal policy: $a = \arg \max_{a'} Q(s, a')$

- Deterministic




Approximate Dynamic Programming

- Use a function approximator (eg. neural network) $\hat{V}(s; w)$, where w are weights, to approximate V
 - $V(s)$ is no longer stored at every state
 - Weights w are updated using Bellman backups
- Basic algorithm: (We will learn about other variants too)
 - Sample some states, $\{s_i\}$
 - For each s_i , generate $\tilde{V}(s_i) = \max_a \{r(s, a) + \gamma \sum_{s'} [p(s'|s_t, a)\hat{V}(s'; w)]\}$
 - Using $\{s_i, \tilde{V}(s_i)\}$, update weights w via regression (supervised learning)


Generalized Policy Evaluation and Policy Improvement

- Start with initial policy π and value function V or Q
- Use policy π to update V or Q : $a = \pi(s)$
 - $V(s) \leftarrow r(s, a) + \gamma \sum_{s'} [p(s'|s, a)V(s')]$
 - $Q(s, a) \leftarrow r(s, a) + \gamma \sum_{s'} [p(s'|s, a)V(s')]$
 - In general, any **policy evaluation algorithm**
- Use V or Q to update policy π :
 - Given $V(s)$, $\pi(s) = \arg \max_a \{r(s, a) + \gamma \sum_{s'} [p(s'|s, a)V(s')]\}$
 - Given $Q(s, a)$, $\pi(s) = \arg \max_a Q(s, a)$
 - In general, any **policy improvement algorithm**



Convergence

- At convergence, the following are simultaneously satisfied:
 - $V(s) = r(s, a) + \gamma \sum_{s'} [p(s'|s, a)V(s')]$
 - $\pi(s) = \arg \max_{a'} \{r(s, a') + \gamma \sum_{s'} [p(s'|s, a')V(s')]\}$
- This is the principle of optimality
- Therefore, the value function and policy are optimal



Key Definitions So Far

- “**State-value function**”: $V_\pi(s)$ -- expected return starting from state s and following policy π
 - $V_\pi(s) = \mathbb{E}_{a_t \sim \pi} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$
 - Expectation is on the random sequence $\{s_0, a_0, s_1, a_1, \dots\}$
- “**Action-value function**”, or “ **Q function**”: $Q_\pi(s, a)$ -- expected return starting from state s , taking action a , and then following policy π
 - $Q_\pi(s, a) = \mathbb{E}_{a_t \sim \pi} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s, a_0 = a]$

Key Definitions So Far


- “**Value iteration**”: The process of iteratively updating value function
 - With DP, we only need to keep track of value function V or Q , and the policy π is implicit – determined from value function
- “**Policy iteration**”: The process of iteratively updating policy
 - This is done implicitly with Bellman backups
- “**Greedy policy**”: the policy obtained from choosing the best action based on the current value function
 - If the value function is optimal, the greedy policy is optimal

Towards Model-Free Learning

- Policy evaluation
 - Monte-Carlo (MC) Sampling
 - Temporal-difference (TD)
- Policy improvement
 - ϵ -greedy policies

Actually Learning from Experience: Monte-Carlo Policy Evaluation

- Start with initial policy π and value function V or Q
- Use policy π to update V : $a = \pi(s)$
 - Apply π to obtain trajectory $\{s_0, a_0, s_1, a_1, \dots\}$
 - Compute return: $R := \sum \gamma^t r(s_t, a_t)$
 - Repeat for many episodes to obtain empirical mean
 - “**Episode**”: a single “try” that produces a single trajectory
- Use V or Q to update policy π




Actually Learning from Experience: Monte-Carlo Policy Evaluation

- To obtain empirical mean, we record $N(s)$, # of times s is visited for every state
 - Start at $N(s) = 0$ for all s
 - Note that this means storing N (and S below) at every state
- First-visit MC Policy Evaluation:
 - At the first time t that s is visited in an episode,
 - Increment $N(s) \leftarrow N(s) + 1$
 - Record return $R(s) \leftarrow R(s) + \sum \gamma^t r(s_t, a_t)$
 - Repeat for many episodes
 - Estimate value: $V(s) = \frac{R(s)}{N(s)}$

Actually Learning from Experience: Monte-Carlo Policy Evaluation

- To obtain empirical mean, we record $N(s)$, # of times s is visited for every state
 - Start at $N(s) = 0$ for all s
 - Note that this means storing N (and S below) at every state
- **Every**-visit MC Policy Evaluation:
 - **Every** time t that s is visited in an episode,
 - Increment $N(s) \leftarrow N(s) + 1$
 - Record return $R(s) \leftarrow R(s) + \sum \gamma^t r(s_t, a_t)$
 - Repeat for many episodes
 - Estimate value: $V(s) \approx \frac{R(s)}{N(s)}$


Incremental Updates

- Instead of estimating $V_\pi(s)$ after many episodes, we can update it incrementally after every episode after receiving return R
 - $N(s) \leftarrow N(s) + 1$
 - $V(s) \leftarrow V(s) + \frac{1}{N(s)}(R - V(s))$
- More generally, we can weight the second term differently
 - $V(s) \leftarrow V(s) + \alpha(R - V(s))$ 

“step size” “target”


Monte-Carlo Policy Evaluation

- Start with initial policy π and value function V or Q
- Use policy π to update V : $a = \pi(s)$
 - MC policy evaluation provides estimate of V_π
 - Many episodes are needed to obtain accurate estimate
 - Model-free with MC!
- Use V or Q to update policy π
 - Greedy policy?



Monte-Carlo Policy Evaluation

- Start with initial policy π and value function V or Q
- Use policy π to update V : $a = \pi(s)$
 - MC policy evaluation provides estimate of V_π
 - Many episodes are needed to obtain accurate estimate
 - Model-free with MC!
- Use V or Q to update policy π
 - ~~Greedy policy?~~
 - Greedy policy lacks exploration, so V_π is not estimated at many states
 - ϵ -greedy policy



ϵ -Greedy Policy

- Also known as ϵ -greedy exploration
- Choose random action with probability ϵ
 - Typically uniformly random
 - If a takes on discrete values, then all actions will be chosen eventually
- Choose action from greedy policy with probability $1 - \epsilon$
 - $a = \arg \max_{a'}\{r(s, a') + \gamma \sum_s [p(s|s_t, a')V(s)]\}$
 - Still requires model, $p(s|s_t, a)$...
 - Solution: Q function

Monte-Carlo Policy Evaluation for Q Functions


- To obtain empirical mean, we record $N(s, a)$, # of times s is visited for every state and action a is applied
 - Start at $N(s, a) = 0$ for all s and a
 - Note that this means N (and S, A below) must be stored for every s and a
- First-visit MC Policy Evaluation:
 - At the first time t that s is visited and action a is applied in an episode,
 - Increment $N(s, a) \leftarrow N(s, a) + 1$
 - Record return $R(s, a) \leftarrow R(s, a) + \sum \gamma^t r(s_t, a_t)$
 - Repeat for many episodes
 - Estimate action-value function: $Q(s, a) = \frac{R(s, a)}{N(s, a)}$

Incremental Updates

- Instead of estimating $Q(s, a)$ after many episodes, we can update it incrementally after every episode after receiving return R
 - $N(s, a) \leftarrow N(s, a) + 1$
 - $Q(s, a) \leftarrow Q(s, a) + \frac{1}{N(s, a)}(R - Q(s, a))$
- More generally, we can weight the second term differently
 - $Q(s, a) \leftarrow Q(s, a) + \alpha(R - Q(s, a))$

Basic Q Learning

- Start with initial policy π and value function V or Q
- Use policy π to update Q : $a = \pi(s)$
 - Repeat for many episodes:
 - $N(s, a) \leftarrow N(s, a) + 1$
 - $Q(s, a) \leftarrow Q(s, a) + \frac{1}{N(s, a)} (R(s, a) - Q(s, a))$
- Use Q to update policy π
 - ϵ -greedy policy
 - With probability ϵ , choose random control
 - With probability $1 - \epsilon$, choose $a = \arg \max_{a'} \{Q(s, a')\}$
 - Pick $\epsilon = \frac{1}{k}$, where k is the # of algorithm iterations
 - Explore less as value function becomes more accurate



Outline

- Reinforcement learning problem setup
- Imitation learning
- Basic ideas in RL
- **Model-free value-based RL**
- Policy-based and actor-critic RL


DP vs. MC Policy Evaluation

- Suppose the policy π is given

- Dynamic Programming

$$V(s) \leftarrow \max_a Q(s, a)$$

$$Q(s, a) \leftarrow r(s, a) + \gamma \sum_{s'} [p(s'|s, a)V(s')]$$




DP vs. MC Policy Evaluation

- Suppose the policy π is given

- Dynamic Programming

$$V(s) \leftarrow \max_a Q(s, a)$$

$$Q(s, a) \leftarrow r(s, a) + \gamma \sum_{s'} [p(s'|s, a)V(s')]$$




- Monte-Carlo

- Repeat for many episodes:

$$N(s, a) \leftarrow N(s, a) + 1$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R - Q(s, a))$$



Temporal-Difference (TD) Policy Evaluation

- Temporal-difference: a class of policy evaluation techniques $\text{TD}(\lambda)$
- Most basic version: $\text{TD}(0)$
 - From any state s , apply policy $a = \pi(s)$ for one time step, obtain reward $r(s, a)$
 - Get to next state s' , and estimate return from then on using Q function
 - Note: next action is also from the same policy, $a' = \pi(s')$
 - $$Q(s, a) \leftarrow Q(s, a) + \alpha(r(s, a) + \gamma Q(s', a') - Q(s, a))$$
 - Repeat for many episodes to obtain $Q(s, a)$ estimates at many states s and actions a

Temporal-Difference (TD) Policy Evaluation

- Most basic version: TD(0)


$$Q(s, a) \leftarrow Q(s, a) + \alpha(r(s, a) + \gamma Q(s', a') - Q(s, a))$$

- Advantages:

- Online algorithm: Q can be updated during an episode
- Does not require complete episodes

- Disadvantages:


- System may not be Markov
- Initial Q can be very bad and Q may never improve enough



n -step TD

- TD: Look ahead one step

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r(s, a) + \gamma Q(s', a') - Q(s, a))$$




- n -step TD: look ahead n steps

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left(r(s, a) + \gamma r(s_{+1}, a_{+1}) + \cdots + \gamma^{n-1} r(s_{+(n-1)}, a_{+(n-1)}) + \gamma^n Q(s_{+n}, a_{+n}) - Q(s, a) \right)$$

$\underbrace{\quad\quad\quad}_{:= R_n}$

$$:= R_n$$




- MC: Look ahead until the end of the episode

TD(λ)

- n -step return estimate:
 - $R_n = r(s, a) + \gamma r(s_{+1}, a_{+1}) + \cdots \gamma^{n-1} r(s_{+(n-1)}, a_{+(n-1)}) + \gamma^n Q(s_{+n}, a_{+n})$
- λ -return: weighted average of different n -step returns
 - Weights: $(1 - \lambda)\lambda^{n-1}$
 - Estimated return: $(1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} R_n$
 - Small $\lambda \rightarrow$ near-future rewards are more important
 - Large $\lambda \rightarrow$ far-future rewards are more important
- TD(λ) policy evaluation:
 - $Q(s, a) \leftarrow Q(s, a) + \alpha \left(\underbrace{(1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} R_n}_{\text{target}} - Q(s, a) \right)$

SARSA Algorithm

- Start with initial policy π and value function V or Q
- Use ϵ -greedy policy to update Q : $a, a' \sim \pi(s)$, π is ϵ -greedy
 - Repeat for many episodes:
 - $Q(s, a) \leftarrow Q(s, a) + \alpha(r(s, a) + \gamma Q(s', a') - Q(s, a))$
- New policy π is derived from new Q
 - ϵ -greedy policy
 - With probability ϵ , choose random control
 - With probability $1 - \epsilon$, choose $a = \arg \max_{a'}\{Q(s, a')\}$
- If $\epsilon, \alpha \propto \frac{1}{k}$, then $Q(s, a) \rightarrow Q_{\pi^*}(s, a)$



On-Policy and Off-Policy Learning

- From SARSA:
 - Use ϵ -greedy policy to update Q : $a, a' \sim \pi(s)$, π is ϵ -greedy
 - Repeat for many episodes: $Q(s, a) \leftarrow Q(s, a) + \alpha(r(s, a) + \gamma Q(s', a') - Q(s, a))$
- “**Behaviour policy**”: policy used to collect rewards -- $a \sim \pi_B(s)$
- “**Target policy**”: policy used to estimate future rewards -- $a' \sim \pi_T(s)$
- “**On-policy learning**”: $\pi_B = \pi_T$
 - SARSA is an on-policy learning algorithm
- “**Off-policy learning**”: $\pi_B \neq \pi_T$
$$Q(s, a) \leftarrow Q(s, a) + \alpha(r(s, a) + \gamma Q(s', a') - Q(s, a)), \text{ where } a \sim \pi_B(s), a' \sim \pi_T(s)$$

Off-Policy Learning

- Off-policy learning: Behaviour and target policies are different
$$Q(s, a) \leftarrow Q(s, a) + \alpha(r(s, a) + \gamma Q(s', a') - Q(s, a)), \text{ where } a \sim \pi_B(s), a' \sim \pi_T(s)$$
- Advantages:
 - Learn from observing another agent (eg. human) execute a different policy
 - Learn from experience generated from old policies
 - Improve two policies at once, while following one policy
- Example: Q-Learning algorithm
 - π_B is ϵ -greedy with respect to Q
 - π_T is greedy with respect to Q

Q-Learning Algorithm

- Start with initial policy π and value function V or Q
- Update Q :
 - Repeat for many episodes with ϵ -greedy policy $a \sim \pi_B(s)$:
 - $$Q(s, a) \leftarrow Q(s, a) + \alpha \left(r(s, a) + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$$
 - Both the ϵ -greedy π_B and the greedy π_T are derived from Q
- If $\epsilon, \alpha = \frac{1}{k}$, then $Q(s, a) \rightarrow Q_{\pi^*}(s, a)$

Function Approximation

- So far, $Q(s, a)$ is stored in a multi-dimensional array
 - Model-free, but cannot solve large problems
- Parametrize value functions with parameters (or weights) w
 - $\hat{Q}(s, a; w) \approx Q(s, a)$
 - Update parameters w using MC- or TD-based learning
 - Hopefully, Q is generalizable to different states s and actions a

Fitting to a Known Q_π

- Fit $\hat{Q}(s, a; w)$ to $Q_\pi(s, a)$

$$\underset{w}{\text{minimize}} \left\| Q_\pi(S, A) - \hat{Q}(S, A; w) \right\|_2^2$$

- Training data: $\{(s_i, a_i), Q_\pi(s_i, a_i)\}$
- The collection of states and actions in training data is denoted S and A
- Gradient with respect to w :
 - $\frac{\partial}{\partial w} \left\| \hat{Q}(S, A; w) - Q_\pi(S, A) \right\|_2^2 = 2 \left(Q_\pi(S, A) - \hat{Q}(S, A; w) \right) \frac{\partial \hat{Q}(S, A; w)}{\partial w}$
- Gradient descent:
 - $w \leftarrow w - \alpha \left(Q_\pi(S, A) - \hat{Q}(S, A; w) \right) \frac{\partial \hat{Q}(S, A; w)}{\partial w}$
 - In practice, use stochastic gradient descent

Monte-Carlo Incremental Weight Updates

- First-visit MC policy evaluation
 - At the first time t that s is visited in an episode,
 - Increment $N(s, a) \leftarrow N(s, a) + 1$
 - Record return $R(s, a) \leftarrow R(s, a) + \sum \gamma^t r(s_t, a_t)$
 - Repeat for many episodes
 - Before: estimate action-value function at visited s and a . $Q(s, a) \approx \frac{R(s, a)}{N(s, a)}$
- Now: Above procedure produces “training data” $\{S, A, R\}$
 - Storing a set of S, A, R , etc. is called “**experience replay**”
 - This is as opposed to updating w as data is being collected
- Update weights:
 - $w \leftarrow w - \alpha \left(R - \hat{Q}(S, A; w) \right) \frac{\partial \hat{Q}(S, A; w)}{\partial w}$
- Guaranteed to converge to local optimum

$$Q(s, a) \approx \frac{R(s, a)}{N(s, a)}$$

Temporal-Difference Incremental Weight Updates

- Most basic version: TD(0)
 - From any state s , apply policy $a = \pi(s)$ for one time step, obtain reward $r(s, a)$
 - Before: get to next state s' , and estimate return from then on using Q function
 - $Q(s, a) \leftarrow Q(s, a) + \alpha(r(s, a) + \gamma Q(s', a') - Q(s, a))$
 - Repeat for many episodes to obtain $Q(s, a)$ estimates at many states s and actions a
- Above procedure produces a collection of current and next states and actions, S, A, R, S', A'
 - reward*
- Update weights using TD target:
 - $w \leftarrow w - \alpha \left(R + \gamma \hat{Q}(S', A'; w) - \hat{Q}(S, A; w) \right) \frac{\partial \hat{Q}(S, A; w)}{\partial w}$
- Not always guaranteed to converge to local minimum

Q-Learning With Function Approximation

Goal: Given a set of weights w^- , find the next set of weights w in $\hat{Q}(s, a; w)$

1. From any state s , apply ϵ -greedy policy with respect to $\hat{Q}(s, a; w^-)$
 - This produces a collection S, A, R, S'
2. Sample from the above collection to obtain a smaller data set $\tilde{S}, \tilde{A}, \tilde{R}, \tilde{S}'$
3. Update weights using stochastic gradient descent

$$\underset{w}{\text{minimize}} \left\| \tilde{R} + \gamma \max_{a'} \hat{Q}(\tilde{S}', a'; w^-) - \hat{Q}(\tilde{S}, \tilde{A}; w) \right\|_2^2$$

- Use deep Q -network (DQN) for $\hat{Q}(\tilde{S}, \tilde{A}; w) \rightarrow$ deep Q-learning

Deep Q-Learning Example: Atari Games

- Mnih et al. “Playing Atari with Deep Reinforcement Learning,” 2013




- States: pixels from last few frames
- Actions: controls in the game
- Reward: game score
- Deep Q network: convolutional and fully connected layers


Starting out - 10 minutes of training

**The algorithm tries to hit the ball back, but
it is yet too clumsy to manage.**

Deep Q-Learning: Robotic Arms

- Gu et al. “Deep Reinforcement Learning for Robotic Manipulation with Asynchronous Off-Policy Updates,” 2017.
- States: joint angles, end-effector positions, and their time derivatives, target position
- Actions: joint velocities of arm, torque of fingers
- Task: open door, pick up object and place it elsewhere
- Deep Q network: two fully connected hidden layers, 100 units each
- Main challenge: use multiple robots to learn at the same time and share knowledge





Single Worker - 4 hours

1.5x

Outline

- Reinforcement learning problem setup
- Imitation learning
- Basic ideas in RL
- Model-free value-based RL
- **Policy-based and actor-critic RL**

Categories of RL


- Model-based
 - Explicitly involves an MDP model
- Model-free
 - Does not involve an MDP model
- Value based
 - Learns value function, and derives policy from value function
- **Policy based**
 - **Learns policy without value function**
- **Actor critic**
 - **Incorporates both value function and policy**

Policy Gradients

- If we executed a policy π_θ from state s_0 , we obtain a trajectory
 - $\tau := (s_0, a_0, s_1, a_1, \dots)$
 - Note: this is a random variable
- The return is given by $R(\tau) := \sum_{t \geq 0} \gamma^t r(s_t, a_t)$
 - Also a random variable
- Expected return given parameters θ : $J(\theta) := \mathbb{E}_{\tau \sim p(\tau; \theta)}[R(\tau)]$
- Parameters for the optimal policy:
 - $\theta^* = \arg \max_{\theta} \mathbb{E}_{\tau \sim p(\tau; \theta)}[R(\tau)]$


Policy Gradients

- Strategy: differentiate $J(\theta)$ w.r.t. θ and perform stochastic gradient ascent
 - Do this in a way that is model-free and computationally tractable




Policy Gradients

- Strategy: differentiate $J(\theta)$ w.r.t. θ and perform stochastic gradient ascent
 - Do this in a way that is model-free and computationally tractable




Policy Gradients

- Strategy: differentiate $J(\theta)$ w.r.t. θ and perform stochastic gradient ascent
 - Do this in a way that is model-free and computationally tractable




Policy Gradients

- Strategy: differentiate $J(\theta)$ w.r.t. θ and perform stochastic gradient ascent
 - Do this in a way that is model-free and computationally tractable



Policy Gradients

- Strategy: differentiate $J(\theta)$ w.r.t. θ and perform stochastic gradient ascent
 - Do this in a way that is model-free and computationally tractable
- To achieve this
 - Write out $J(\theta)$
 - Take gradient
 - Do a math trick
 - Obtain gradient expression that can be estimated easily



Write Out $J(\theta)$ and Take Gradient

- $J(\theta) := \mathbb{E}_{\tau \sim p(\tau; \theta)}[R(\tau)]$
- $J(\theta) = \int_{\tau} R(\tau)p(\tau; \theta)d\tau$
- $\nabla_{\theta}J(\theta) = \int_{\tau} R(\tau)\nabla_{\theta}p(\tau; \theta)d\tau$
 - Hard...

Log Gradient Trick

- $\nabla_{\theta} J(\theta) = \int_{\tau} R(\tau) \nabla_{\theta} p(\tau; \theta) d\tau$
- Trick:
 - $\nabla_{\theta} p(\tau; \theta) = p(\tau; \theta) \frac{\nabla_{\theta} p(\tau; \theta)}{p(\tau; \theta)} = p(\tau; \theta) \nabla_{\theta} \log p(\tau; \theta)$
 - $\nabla_{\theta} J(\theta) = \int_{\tau} R(\tau) p(\tau; \theta) \nabla_{\theta} \log p(\tau; \theta) d\tau$
 - $\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)} [R(\tau) \nabla_{\theta} \log p(\tau; \theta)]$
 - Gradient is an expectation – can estimate this using techniques we learned before!

Model-Free Estimate of Gradient

- $\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)} [R(\tau) \nabla_{\theta} \log p(\tau; \theta)]$
- $p(\tau; \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$
- $\log p(\tau; \theta) = \sum_{t \geq 0} [\log p(s_{t+1} | s_t, a_t) + \log \pi_{\theta}(a_t | s_t)]$
- $\nabla_{\theta} \log p(\tau; \theta) = \sum_{t \geq 0} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$
 - Amazingly, model-free
 - Markov property is not used
 - $\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$ is known: since the form of π_{θ} is known
 - Eg. Backprop if π_{θ} is a neural network

Monte-Carlo Gradient Estimate


- Results so far:
 - $\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)}[R(\tau) \nabla_{\theta} \log p(\tau; \theta)]$
 - $\nabla_{\theta} \log p(\tau; \theta) = \sum_{t \geq 0} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$
- Some more algebra to write out gradient of $\nabla_{\theta} J(\theta)$
 - $\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)}[R(\tau) \sum_{t \geq 0} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]$
 - $\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)}[\sum_{t \geq 0} R(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]$
 - $\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N [\sum_{t \geq 0} R(\tau_i) \nabla_{\theta} \log \pi_{\theta}(a_{t,i} | s_{t,i})]$

REINFORCE Algorithm

- (Monte-Carlo Policy Gradient)
- Use policy $\pi_\theta(a|s)$ to obtain trajectories $\tau_i = \{s_{0,i}, a_{0,i}, \dots\}$
- Estimate the gradient of the reward
 - $\nabla_\theta J(\theta) \approx \sum_{i=1}^N [\sum_{t \geq 0} R(\tau_i) \nabla_\theta \log \pi_\theta(a_{t,i}|s_{t,i})]$
- Update policy parameters via (stochastic) gradient ascent
 - $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

Observation 1

- Gradient estimate:
 - $\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau; \theta)} [\sum_{t \geq 0} R(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]$
 - $\nabla_{\theta} J(\theta) \approx \sum_{i=1}^N [\sum_{t \geq 0} R(\tau_i) \nabla_{\theta} \log \pi_{\theta}(a_{t,i} | s_{t,i})]$
- Gradient estimate also works for POMDPs without modification



Observation 1

- Gradient estimate:
 - $\nabla_{\theta}J(\theta) \approx \mathbb{E}_{\tau \sim p(\tau; \theta)}[\sum_{t \geq 0} R(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]$
 - $\nabla_{\theta}J(\theta) \approx \sum_{i=1}^N [\sum_{t \geq 0} R(\tau_i) \nabla_{\theta} \log \pi_{\theta}(a_{t,i} | s_{t,i})]$
- Gradient estimate also works for POMDPs without modification
- Parameter updates: $\theta \leftarrow \theta + \alpha \nabla_{\theta}J(\theta)$
 - Trajectories have high reward will be made more likely
 - Trajectories with low reward will be made less likely
 - A high-reward trajectory has good actions... **on average**


Observation 2

- Gradient estimate:
 - $\nabla_{\theta} J(\theta) \approx \mathbb{E}_{\tau \sim p(\tau; \theta)} [\sum_{t \geq 0} R(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]$
- Causality?
 - $R(\tau)$ is the reward of the entire trajectory
 - $R(\tau)$ is multiplied in every term of the sum
 - τ includes times before t
 - So, according to the above, the weight of $\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$ depends on times prior to t ?
- Simple fix:
 - $\nabla_{\theta} J(\theta) \approx \mathbb{E}_{\tau \sim p(\tau; \theta)} [\sum_{t \geq 0} [(\sum_{t' \geq t} \gamma^{t'-t} r(s_t, a_t)) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]]$

Observation 3

- Gradient estimate:


- $\nabla_{\theta} J(\theta) \approx \mathbb{E}_{\tau \sim p(\tau; \theta)} [\sum_{t \geq 0} R(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]$



Observation 3

- Gradient estimate:


- $\nabla_{\theta} J(\theta) \approx \mathbb{E}_{\tau \sim p(\tau; \theta)} [\sum_{t \geq 0} R(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]$



Observation 3

- Gradient estimate:

- $\nabla_{\theta} J(\theta) \approx \mathbb{E}_{\tau \sim p(\tau; \theta)} [\sum_{t \geq 0} R(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]$



- Performance is measured by reward $R(\tau)$

- But what is considered “good”?
- Need a baseline of comparison!
- $\nabla_{\theta} J(\theta) \approx \mathbb{E}_{\tau \sim p(\tau; \theta)} [\sum_{t \geq 0} (R(\tau) - b) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]$
- Fact: expectation is unchanged as long as b does not depend on θ

Revised REINFORCE


- (Monte-Carlo Policy Gradient)
- Use policy $\pi_\theta(a|s)$ to obtain a trajectory $\tau = \{s_0, a_0, \dots\}$
- Estimate the gradient of the reward
 - $\nabla_\theta J(\theta) \approx \mathbb{E}_{\tau \sim p(\tau; \theta)} [\sum_{t \geq 0} [(\sum_{t' \geq t} \gamma^{t'-t} r(s_t, a_t) - b) \nabla_\theta \log \pi_\theta(a_t | s_t)]]$
- Update policy parameters via (stochastic) gradient ascent
 - $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

Picking a Baseline

- Many choices
- One intuitive choice
 - $b = V_\pi(s)$
 - Define $\mathcal{A}_\pi(s, a) := r(s, a) + \gamma V_\pi(s') - V_\pi(s)$
 - Good action: one that gives a **return** that is large relative to V
 - Bad action: one that gives a **return** that is small relative to V
 - $\mathcal{A}_\pi(s, a)$ -- “**advantage function**”
- But we don’t know V ...
 - Learn it!

Actor-Critic Methods

- Actor (policy π) decides which actions to take
- Critic (value function V) decides how good the action is



Actor-Critic Methods

- Basic algorithm, combining everything we've learned:
 1. Start with some initial policy π_θ and value function $\hat{V}(s; w)$
 - θ and w are parameters
 2. Collect data S, R, S' by executing policy
 3. Update V_ϕ : minimize
$$\min_w \left\| \tilde{R} + \gamma \hat{V}(\tilde{S}'; w^-) - V(\tilde{S}; w) \right\|_2^2$$
 - Many methods (eg. stochastic gradient descent)
 4. Estimate policy gradient: $\nabla_\theta J(\theta) \approx \mathbb{E}_{\tau \sim p(\tau; \theta)} \left[\sum_{t \geq 0} (\tilde{R} + \gamma V_\pi(\tilde{S}') - V_\pi(\tilde{S})) \nabla_\theta \log \pi_\theta(a_t | s_t) \right]$
 5. Improve policy via gradient ascent: $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$
 6. Repeat 2-5 many times

State-of-the-Art Policy Gradient Methods

- Trust region policy optimization (TRPO)
 - <https://arxiv.org/abs/1502.05477>
- Proximal policy optimization (PPO)
 - <https://arxiv.org/abs/1707.06347>