

Recurrent Neural Networks

CMPT 726

Mo Chen

SFU Computing Science

18 Nov. 2020

Goodfellow, Bengio, and Courville: Deep Learning textbook Ch. 10

Sequential Data with Neural Networks

- Sequential input / output
 - Many inputs, many outputs $x_{1:T} \rightarrow y_{1:S}$
 - e.g. object tracking, speech recognition with HMMs; on-line/batch processing
 - One input, many outputs $x \rightarrow y_{1:S}$
 - e.g. image captioning
 - Many inputs, one output $x_{1:T} \rightarrow y$
 - e.g. video classification

Outline

Recurrent Neural Networks

Long Short-Term Memory

Temporal Convolutional Networks

Examples

Outline


Recurrent Neural Networks

Long Short-Term Memory

Temporal Convolutional Networks

Examples

Hidden State




- Basic idea: maintain a state h_t
- State at time t depends on input x_t and previous state h_{t-1}
- It's a neural network, so relation is non-linear function of these inputs and some parameters W :

$$h_t = f_x(h_{t-1}, x_t; W) = \sigma(W_x x_t + W_h h_{t-1})$$

- Parameters W and function $f(\cdot)$ reused at all time steps

Outputs




- Output z_t also depends on the hidden state:

$$z_t = f_z(\mathbf{h}_t; \mathbf{W}_z) = \sigma(\mathbf{W}_z \mathbf{h}_t)$$


- Again, parameters/function reused across time

Recurrent neural network (RNN)

- Has feedback loops to capture temporal or sequential information
- Has the ability to learn tasks that require “memory” of events from many time steps ago
- Long short-term memory (LSTM): special type of RNN with advantages in numerical properties



Unfolding an RNN



$$h_t = f_x(h_{t-1}, x_t; \mathbf{W})$$

h_0 can be a vector of all zeros, or trained

$$z_t = f_z(h_t; \mathbf{W}_z)$$

Vanishing Gradients in RNN

$$h_t = \sigma(W_x x_t + W_h h_{t-1})$$

$$z_n = f(\mathbf{h}_n; \mathbf{W}_z)$$


- Training requires $\frac{\partial E_n}{\partial W_x} = \sum_{i=1}^n \frac{\partial E_n}{\partial z_n} \frac{\partial z_n}{\partial h_n} \frac{\partial h_n}{\partial h_i} \frac{\partial h_i}{\partial W_x}$
 - Everything except $\frac{\partial h_n}{\partial h_i}$ involves “nearby” variables, so focus on $\frac{\partial h_n}{\partial h_i}$

$$\frac{\partial h_n}{\partial h_i} = \frac{\partial h_n}{\partial h_{n-1}} \frac{\partial h_{n-1}}{\partial h_{n-2}} \dots \frac{\partial h_{i+1}}{\partial h_i} = \prod_{t=i}^{n-1} \frac{\partial h_{t+1}}{\partial h_t}$$

$$\frac{\partial h_{t+1}}{\partial h_t} = \frac{\partial}{\partial h_t} f_x(\mathbf{h}_t, \mathbf{x}_{t+1}; \mathbf{W}_x)$$

$$= \sigma'(W_x x_{t+1} + W_h h_t) W_h$$

$$\begin{aligned}\frac{\partial h_n}{\partial h_i} &= \prod_{t=i}^{n-1} \sigma'(W_x x_{t+1} + W_h h_t) W_h \\ &= W_h^{n-1} \prod_{t=i}^{n-1} \sigma'(W_x x_{t+1} + W_h h_t)\end{aligned}$$



- Gradient blows up if largest eigenvalue of W_h is larger than 1
- Gradient vanishes otherwise

Gradients

- Basic RNN is not very effective
- Need many time steps / complex model for challenging tasks
- Gradients in learning are a problem
 - Too large: can be handled with **gradient clipping** (truncate gradient magnitude)
 - Too small: can be handled with network structures / **gating functions** (LSTM, GRU)

Outline


Recurrent Neural Networks

Long Short-Term Memory

Temporal Convolutional Networks


Examples

Long Short-Term Memory



- Hochreiter and Schmidhuber, Neural Computation 1997
 - (Figure from Donohue et al. CVPR 2015)
- **Gating functions** $g(\cdot), f(\cdot), o(\cdot)$ reduce vanishing gradients

Long Short-Term Memory



$$i_t = \sigma(W_{xi}x_t + W_{hi}h_{t-1} + b_i) \quad (1)$$

$$f_t = \sigma(W_{xf}x_t + W_{hf}h_{t-1} + b_f) \quad (2)$$

$$o_t = \sigma(W_{xo}x_t + W_{ho}h_{t-1} + b_o) \quad (3)$$

$$g_t = \tanh(W_{sc}x_t + W_{hc}h_{t-1} + b_c) \quad (4)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot g_t \quad (5)$$

$$h_t = o_t \odot \tanh c_t \quad (6)$$

Long Short-Term Memory

Let's consider $\frac{dc_t}{dc_{t-1}}$

- Full derivative of error function w.r.t. weights will be a product of these
- Note $c_t = c_t(f_t, c_{t-1}, i_t, g_t)$, where f_t, i_t, g_t also depend on c_{t-1}

$$\begin{aligned} i_t &= \sigma(W_{xi}x_t + W_{hi}h_{t-1} + b_i) \\ f_t &= \sigma(W_{xf}x_t + W_{hf}h_{t-1} + b_f) \\ o_t &= \sigma(W_{xo}x_t + W_{ho}h_{t-1} + b_o) \\ g_t &= \tanh(W_{sc}x_t + W_{hc}h_{t-1} + b_c) \\ c_t &= f_t \odot c_{t-1} + i_t \odot g_t \\ h_t &= o_t \odot \tanh c_t \end{aligned}$$

$$\begin{aligned} \frac{dc_t}{dc_{t-1}} &= \frac{\partial c_t}{\partial f_t} \frac{\partial f_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial c_{t-1}} + \frac{\partial c_t}{\partial c_{t-1}} + \frac{\partial c_t}{\partial i_t} \frac{\partial i_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial c_{t-1}} + \frac{\partial c_t}{\partial g_t} \frac{\partial g_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial c_{t-1}} \\ &= c_{t-1} \sigma'(\dots) W_{hf} o_{t-1} \tanh'(c_{t-1}) + f_t + g_t \sigma'(\dots) W_{hi} \tanh'(c_{t-1}) + i_t \tanh'(\dots) W_{hc} \tanh'(c_{t-1}) \end{aligned}$$

Discussion:

- When repeatedly multiplying four different terms added together, there is a smaller chance of vanishing compared to a single term
- f_t can be chosen to be larger or smaller, depending on whether the gradients should propagate backwards to before stage t
- f_t is learned, so neural network learns to control gradient propagation
- Vanishing gradients is *alleviated*, not solved
- Gradients may still explode as well

Outline


Recurrent Neural Networks

Long Short-Term Memory

Temporal Convolutional Networks

Examples

Convolutions to Aggregate over Time




- Control history by d (dilation, holes in the filter) and k (width of the filter)
- Causal convolution, only use elements from the past
- Bai, Kolter, Koltun arXiv 2018

Residual (skip) Connections

$$\hat{\mathbf{z}}^{(i)} = (\hat{z}_1^{(i)}, \dots, \hat{z}_T^{(i)})$$

Residual block (k, d)




Residual block ($k=3, d=1$)

$$\hat{z}_{T-1}^{(1)} \quad \hat{z}_T^{(1)}$$

Convolutional Filter

Identity Map (or 1×1 Conv)



- Include residual connections to allow long-range modeling and gradient flow

Outline


Recurrent Neural Networks

Long Short-Term Memory

Temporal Convolutional Networks


Examples

Example: Image Captioning




- Karpathy and Fei-Fei, CVPR 2015

Example: Video Description



- S. Venugopalan, M. Rohrbach, J. Donahue, R. Mooney, T. Darrell, K. Saenko, ICCV 2015

Example: Machine Translation



- Wu et al., *Google’s Neural Machine Translation System: Bridging the Gap between Human and Machine Translation*, arXiv 2016

Conclusion

- **Readings:**

- [http://www.deeplearningbook.org/
contents/rnn.html](http://www.deeplearningbook.org/contents/rnn.html)
- [https://medium.com/datadriveninvestor/how-
do-lstm-networks-solve-the-problem-of-
vanishing-gradients-a6784971a577](https://medium.com/datadriveninvestor/how-lstm-networks-solve-the-problem-of-vanishing-gradients-a6784971a577)
- [https://weberna.github.io/blog/2017/11/15/L
STM-Vanishing-Gradients.html](https://weberna.github.io/blog/2017/11/15/LSTM-Vanishing-Gradients.html)

- Recurrent neural networks, can model sequential inputs/outputs

- Input includes state (output) from previous time
- Different structures:
 - RNN with multiple inputs/outputs
 - Gated recurrent unit (GRU)
 - Long short-term memory (LSTM)
- Error gradients back-propagated across entire sequence