

# Generative Models

CMPT 726

Mo Chen

SFU Computing Science

23 Nov., 2020

# Outline

Generative Models

Autoencoders

Variational Auto Encoders

Generative Adversarial Networks

# Outline

**Generative Models**

Autoencoders

Variational Auto Encoders

Generative Adversarial Networks

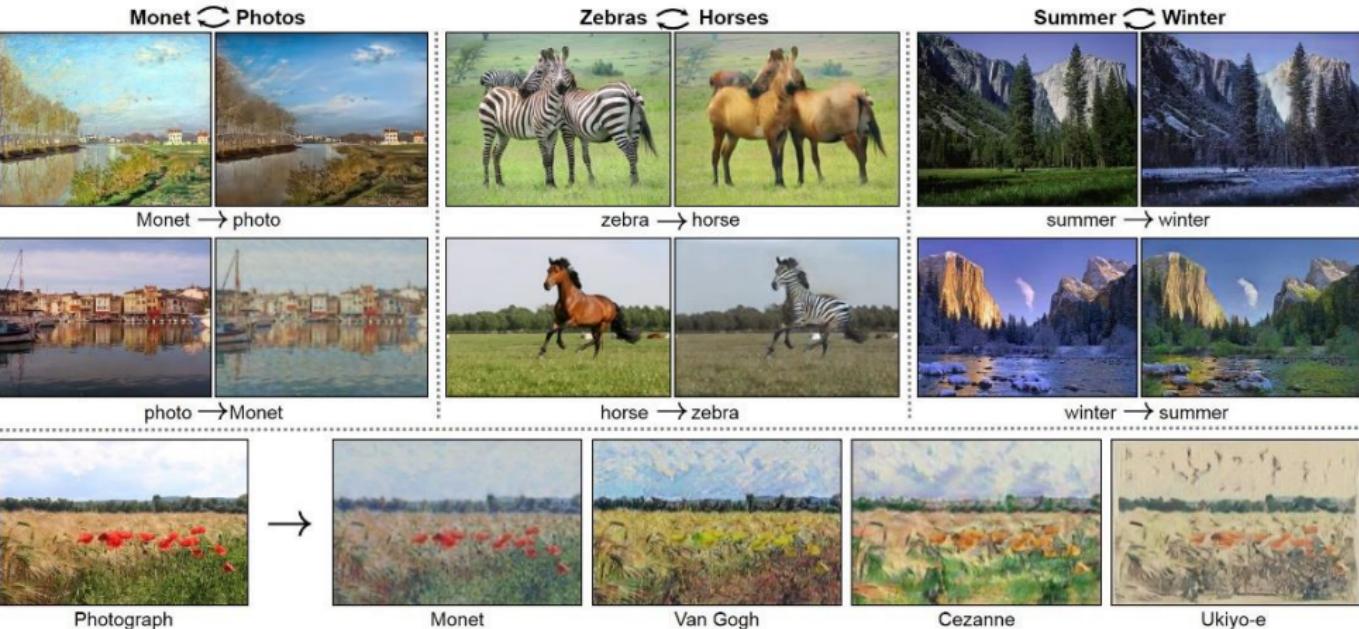
# Generative Models

- Start with training data, with **unknown** distribution  $p(x)$ 
  - No labels!
- Generate new samples from a similar distribution  $\hat{p}(x)$ 
  - $\hat{p}(x)$  is learned from data

# Generative Models

<https://arxiv.org/abs/1406.5298>

# Generative Models



<https://junyanz.github.io/CycleGAN/>

# Outline

Generative Models

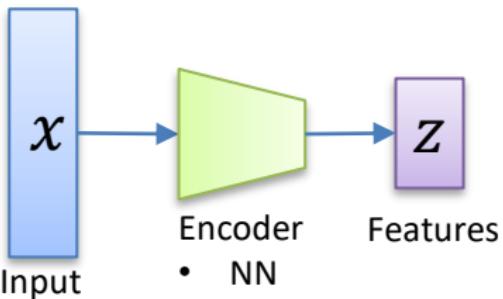
Autoencoders

Variational Auto Encoders

Generative Adversarial Networks

# Autoencoders

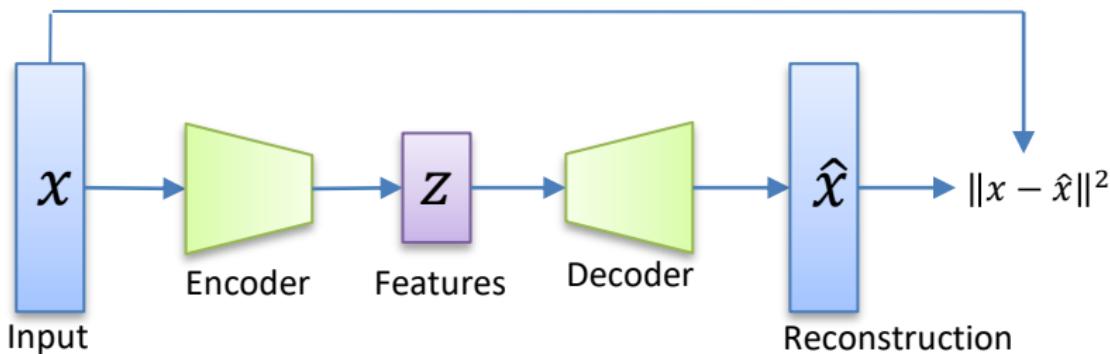
- Unsupervised learning method
  - Learns lower-dimensional features from unlabeled data
  - Labels can be expensive!



- Features  $z$ 
  - Lower-dimensional
  - Can be useful for other tasks
  - How to learn?

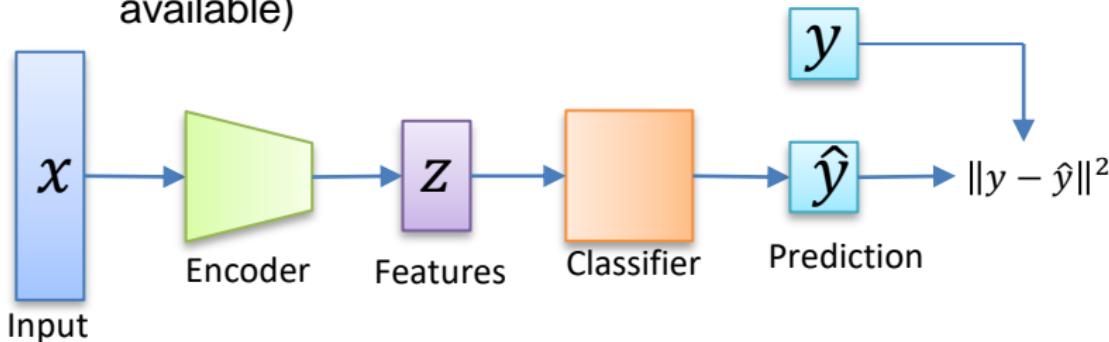
# Autoencoders

- Unsupervised learning method
  - Learn features that allow reconstruction of original data
  - Reconstruction (decoding) done by a decoder
    - Decoder: another neural network
  - Loss function:  $\|x - \hat{x}\|^2$



## After Training

- Replace decoder with another network (e.g. classifier)
  - Supervised learning can be more efficient when features are “pre-learned”
  - Example: train autoencoder on ImageNet dataset, and then train only classifier on bird classification (less data available)



# Outline

Generative Models

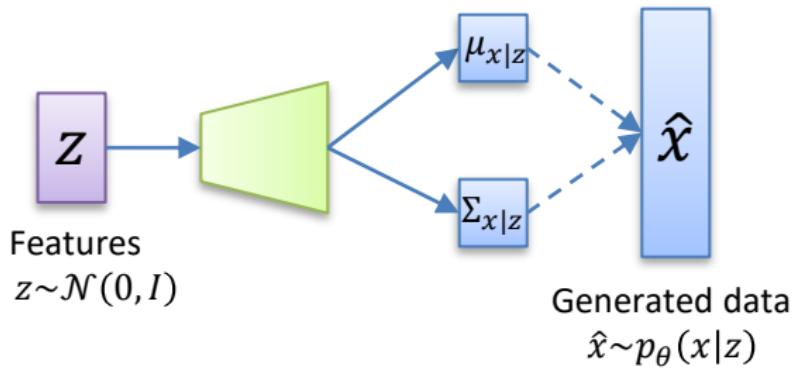
Autoencoders

**Variational Auto Encoders**

Generative Adversarial Networks

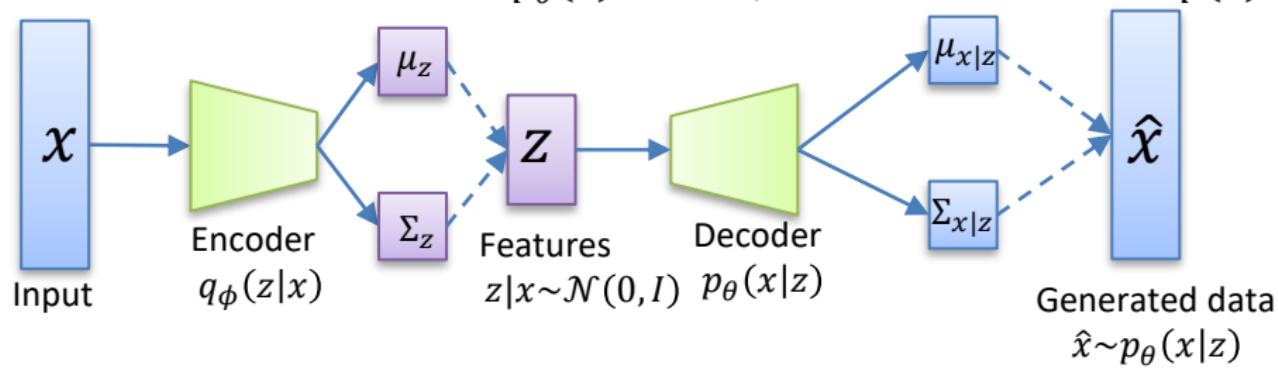
# Variational Autoencoders

- Make autoencoders probabilistic to allow us to generate a variety of new data
- New data generated from features  $z$ 
  - Feature variations:  $z \sim p_\theta(z) = \mathcal{N}(0, I)$  for simplicity
  - Synthetic data variations:  $\hat{x} \sim p_\theta(x|z)$ , represented by a neural network



# Variational Autoencoders: Training

- We still need to learn features that allow reconstruction of original data
  - This time, since there is variation in  $z$  and  $x$ , we cannot simply minimize  $\|x - \hat{x}\|^2$
- Instead, maximize the likelihood  $p_\theta(\hat{x})$  that a training image will be generated
  - fit a distribution  $p_\theta(\hat{x})$  to data, which has distribution  $p(x)$



# ELBO Loss

- How to maximize  $p_\theta(x^{(i)})$ ?
  - Data sample  $x^{(i)}$  is given; we are finding  $\theta$

$$\begin{aligned}
 \log p_\theta(x^{(i)}) &= \mathbb{E}_{z \sim q_\phi(z|x^{(i)})} [\log p_\theta(x^{(i)})] \quad (\mathbb{E}_b[a] = a \text{ if } a \text{ does not depend on } b) \\
 &= \mathbb{E}_{z \sim q_\phi(z|x^{(i)})} \left[ \log \frac{p_\theta(x^{(i)}|z)p_\theta(z)}{p_\theta(z|x^{(i)})} \right] \\
 &= \mathbb{E}_{z \sim q_\phi(z|x^{(i)})} \left[ \log \frac{p_\theta(x^{(i)}|z)p_\theta(z)}{p_\theta(z|x^{(i)})} \frac{q_\phi(z|x^{(i)})}{q_\phi(z|x^{(i)})} \right] \quad (\text{Bayes' rule: } p(a|b) = \frac{p(b|a)p(a)}{p(b)}) \\
 &= \mathbb{E}_{z \sim q_\phi(z|x^{(i)})} [\log p_\theta(x^{(i)}|z)] - \mathbb{E}_{z \sim q_\phi(z|x^{(i)})} \left[ \log \frac{q_\phi(z|x^{(i)})}{p_\theta(z)} \right] + \mathbb{E}_{z \sim q_\phi(z|x^{(i)})} \left[ \log \frac{q_\phi(z|x^{(i)})}{p_\theta(z|x^{(i)})} \right] \\
 &= \mathbb{E}_{z \sim q_\phi(z|x^{(i)})} [\log p_\theta(x^{(i)}|z)] - D_{\text{KL}}(q_\phi(z|x^{(i)}) \| p_\theta(z)) + D_{\text{KL}}(q_\phi(z|x^{(i)}) \| p_\theta(z|x^{(i)}))
 \end{aligned}$$

- $D_{\text{KL}}$  is the Kullback-Leibler divergence
  - a measure of closeness of probability distributions
  - Always  $\geq 0$

# ELBO Loss

- $D_{\text{KL}}$  is the Kullback-Leibler divergence
  - a measure of closeness of probability distributions
  - Always  $\geq 0$

$$\log p_{\theta}(x^{(i)}) = \mathbb{E}_{z \sim q_{\phi}(z|x^{(i)})} [\log p_{\theta}(x^{(i)}|z)] - D_{\text{KL}}(q_{\phi}(z|x^{(i)})||p_{\theta}(z)) + D_{\text{KL}}(q_{\phi}(z|x^{(i)})||p_{\theta}(z|x^{(i)}))$$

- **First term:** Estimate by sampling from decoder (yet more Math tricks, see [1])
- **Second term:** Closed-form expression if we assume  $p_{\theta}(z)$  is Gaussian
- **Third term:** intractable to compute but always  $\geq 0$ , so ignore...
- Maximize first two terms, the evidence lower bound (ELBO)

[1] Kingma, Welling 2013. "Auto-Encoding Variational Bayes." <http://arxiv.org/abs/1312.6114>

# ELBO Loss

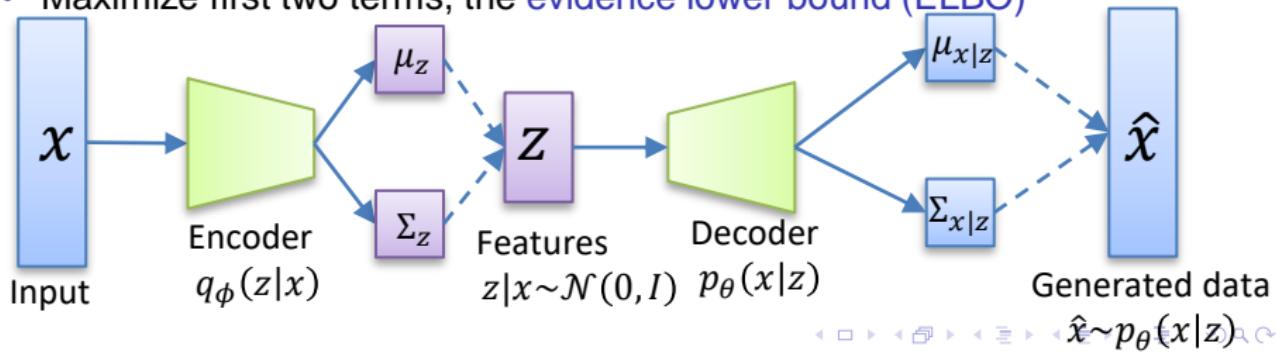
$$l(x^{(i)}, \theta, \phi) = \mathbb{E}_{z \sim q_\phi(z|x^{(i)})} [\log p_\theta(x^{(i)}|z)] - D_{\text{KL}}(q_\phi(z|x^{(i)}) \| p_\theta(z))$$

- **First term:** Estimate by sampling from decoder (yet more Math tricks, see [1])
  - Maximize data likelihood, given latent features
- **Second term:** Closed-form expression if we assume  $p_\theta(z)$  is Gaussian
  - Make encoder  $q_\phi(z|x^{(i)})$  close to  $p_\theta(z) = \mathcal{N}(0, I)$
- Maximize first two terms, the **evidence lower bound (ELBO)**

## ELBO Loss

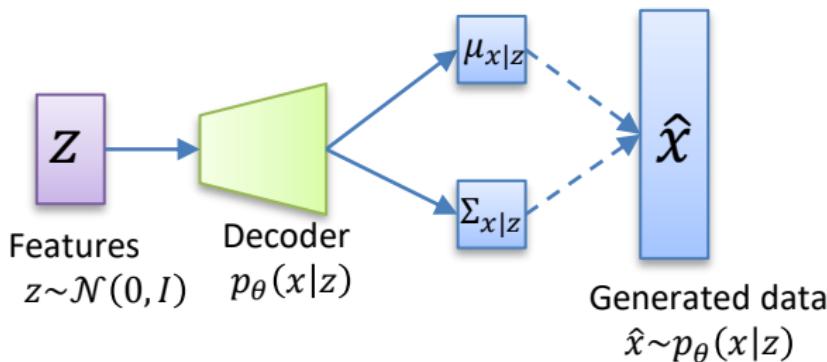
$$l(x^{(i)}, \theta, \phi) = \mathbb{E}_{z \sim q_\phi(z|x^{(i)})} [\log p_\theta(x^{(i)}|z)] - D_{\text{KL}}(q_\phi(z|x^{(i)}) \| p_\theta(z))$$

- **First term:** Estimate by sampling from decoder (yet more Math tricks, see [1])
    - Maximize data likelihood, given latent features
  - **Second term:** Closed-form expression if we assume  $p_\theta(z)$  is Gaussian
    - Make encoder  $q_\phi(z|x^{(i)})$  close to  $p_\theta(z) = \mathcal{N}(0, I)$
  - Maximize first two terms, the **evidence lower bound (ELBO)**

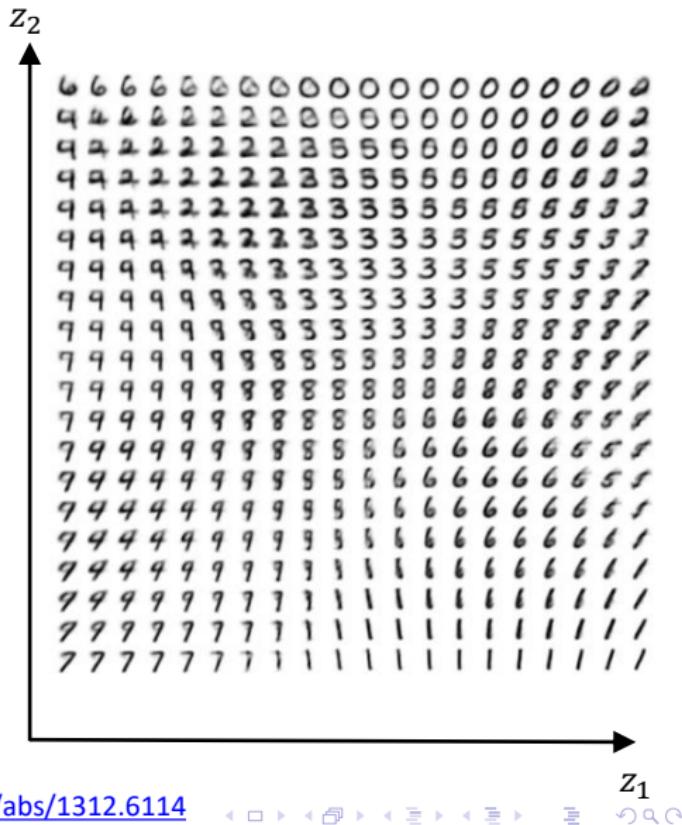
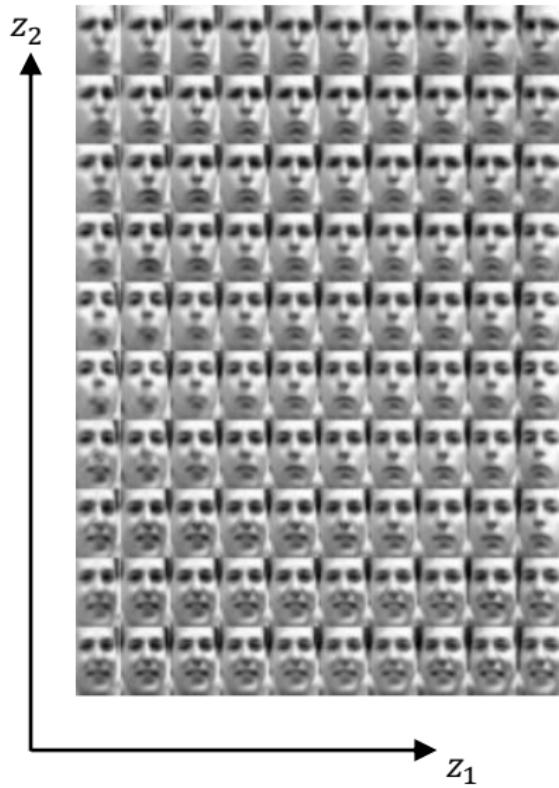


# Test Time

- Throw away encoder
- Sample  $z \sim \mathcal{N}(0, I)$
- Compute  $p_\theta(x|z)$  using decoder
- Sample from decoder to generate new data  $\hat{x}$



## Test Time



"Auto-Encoding Variational Bayes." <http://arxiv.org/abs/1312.6114>

# Outline

Generative Models

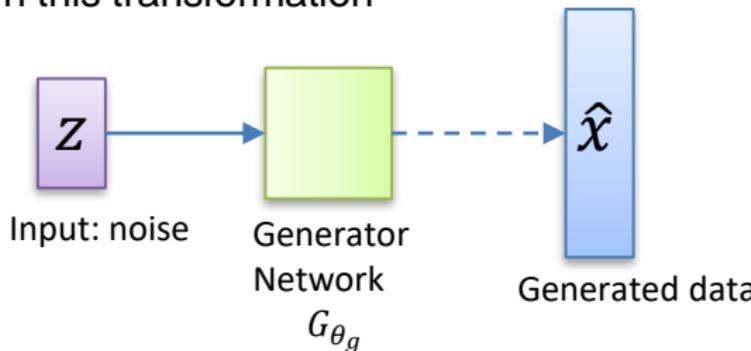
Autoencoders

Variational Auto Encoders

**Generative Adversarial Networks**

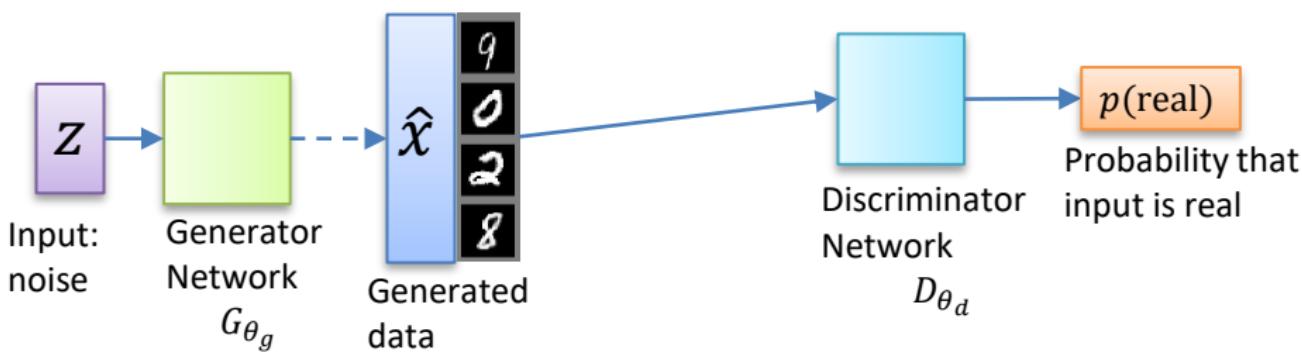
# Generative Adversarial Networks (GANs)

- No explicit representation of probability densities (mean and variance)
- Data generation can be thought of as sampling from data distribution a complex distribution
  - Sample from simple distribution (noise)
  - Transform the samples so that they become distributed according to a more complex distribution
- GANs: Learn this transformation



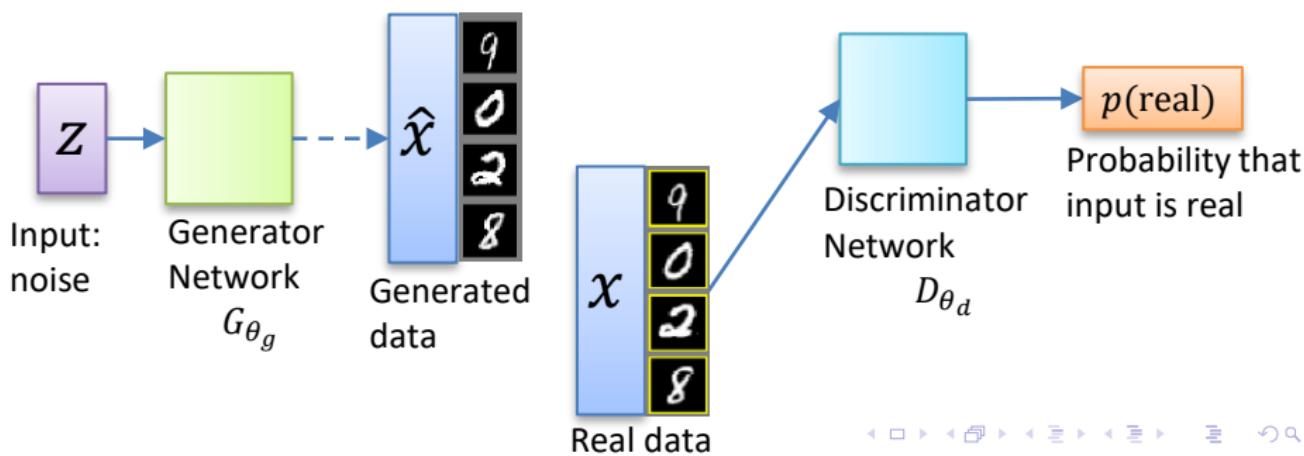
# Generative Adversarial Networks (GANs)

- Generator network needs to output realistic images, i.e. those from the same distribution as some dataset
  - But need to learn the transformation *without labels*
  - Idea: learn to fool a discriminator network
- Discriminator network: discriminate between real and fake data
  - “Labels”: real or not real



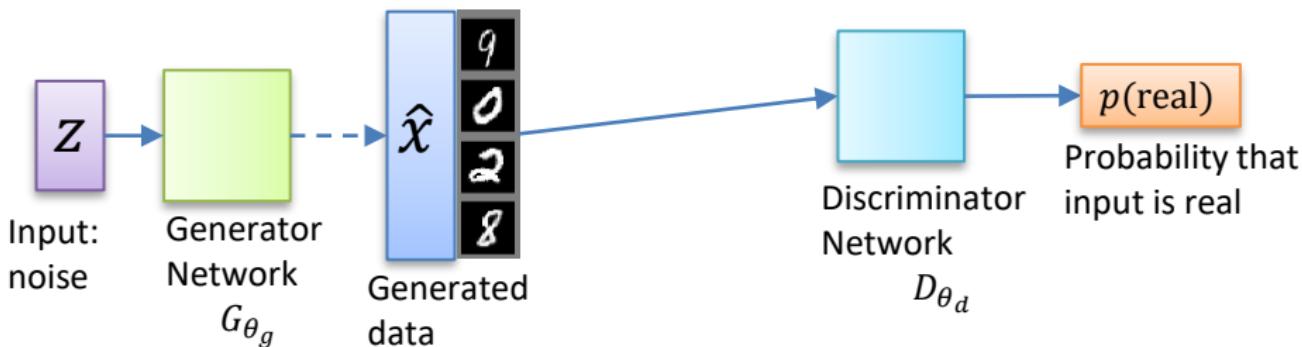
# Generative Adversarial Networks (GANs)

- Generator network needs to output realistic images, i.e. those from the same distribution as some dataset
  - But need to learn the transformation *without labels*
  - Idea: learn to fool a discriminator network
- Discriminator network: discriminate between real and fake data
  - “Labels”: real or not real



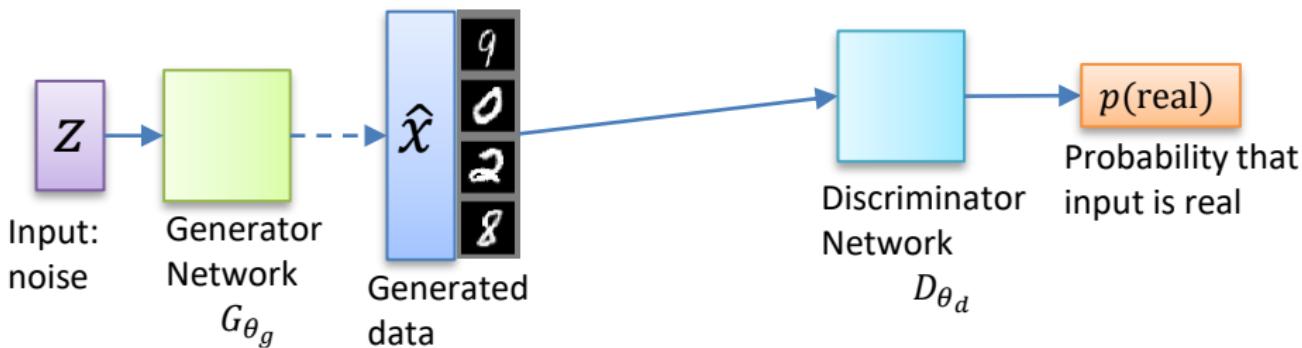
# Discriminator Loss Function

- Maximize  $l_d(x^{(i)}, z^{(i)}; \theta_d) = \log D_{\theta_d}(x^{(i)}) + \log \left( 1 - D_{\theta_d} \left( G_{\theta_g}(z^{(i)}) \right) \right)$ 
  - Fix generator parameters  $\theta_g$
- If the **input is real data**  $x^{(i)}$ , output something close to 1
- If the **input is fake data**  $\hat{x}^{(i)}$ , output something close to 0
  - Fake data is from generator:  $\hat{x}^{(i)} = G_{\theta_g}(z^{(i)})$



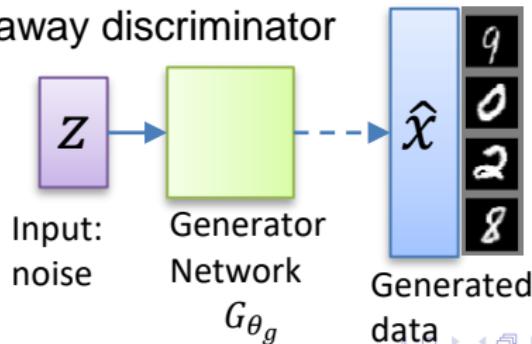
# Generator Loss Function

- Maximize  $l_g(z^{(i)}; \theta_g) = \log(D_{\theta_d}(G_{\theta_g}(z^{(i)}))$ 
  - Fix discriminator parameters  $\theta_d$
- Try to make discriminator output 1 when taking generated data as input

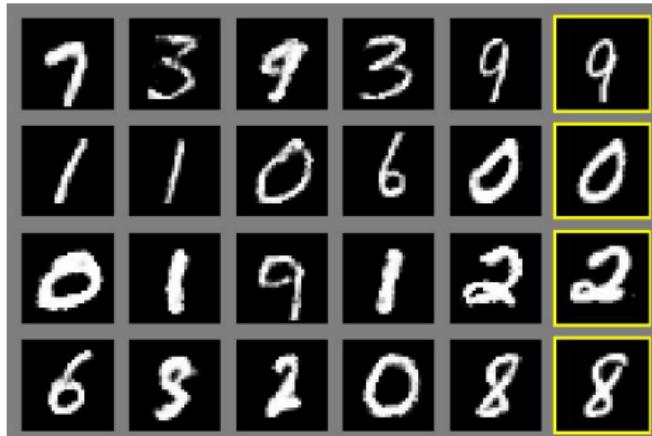


# GANs: Training and Using

- For every training iteration
  - Perform  $k$  steps of discriminator updates
    - Sample  $\{z^{(i)}\}_{i=1}^M$  (random noise),  $\{x^{(i)}\}_{i=1}^M$  (from data)
    - Take ascent step in the  $\nabla_{\theta_d} \sum_{i=1}^M l_d(x^{(i)}, z^{(i)}; \theta_d)$  direction
  - Perform a step of generator update
    - Sample  $\{z^{(i)}\}_{i=1}^M$  (random noise)
    - Take ascent step in the  $\nabla_{\theta_g} \sum_{i=1}^M l_g(z^{(i)}; \theta_g)$  direction
- After training, throw away discriminator



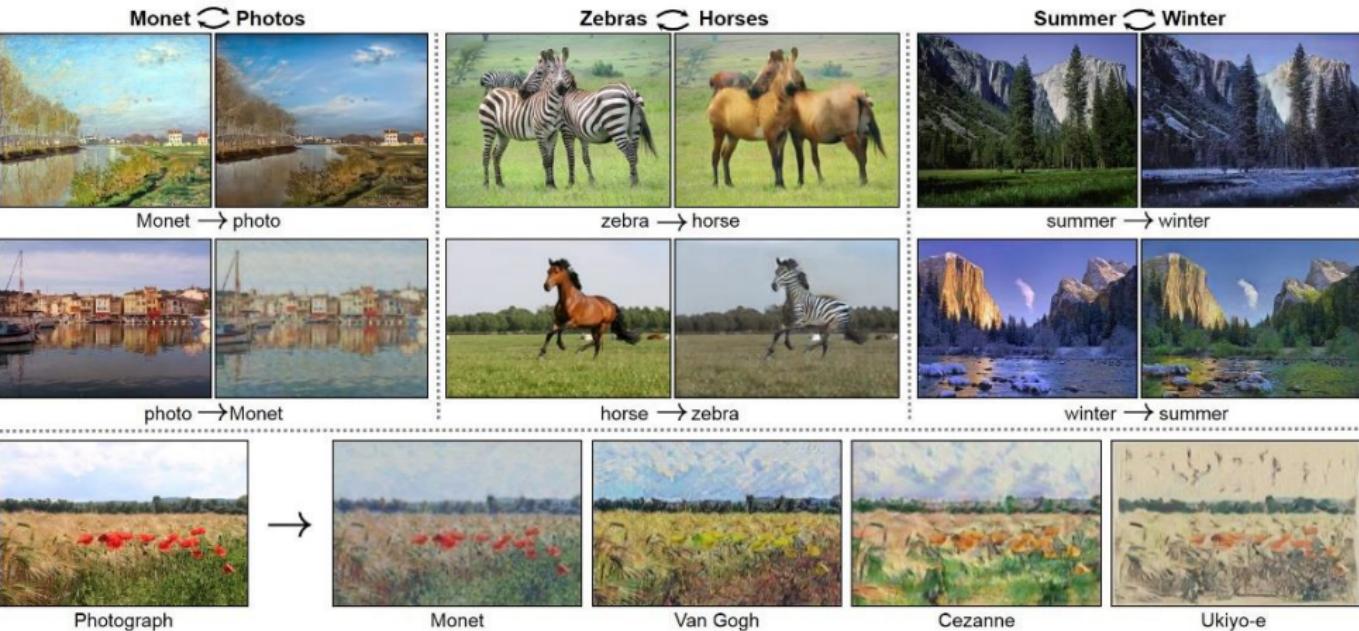
# Generative Models



Goodfellow et al., 2014. "Generative Adversarial Networks."

<http://arxiv.org/abs/1406.2661>

# Generative Models



<https://junyanz.github.io/CycleGAN/>